

LAB2 SOLUTIONS

Computer Vision Foundations: Images and Filtering

EX. 1

Perform binary thresholding using four different values: 30, 60, 90, 120.
As input image, use a grayscale image.

Using Numpy

```
thrs = [30,60,90,120]
fig, ax = plt.subplots(nrows = 1, ncols = len(thrs), figsize = (3*len(thrs),
3*len(thrs)))
for i, thr in enumerate(thrs):
    img = np.where(lena_gs <= int(thr), 0, 255 )
    ax[i].imshow(img, cmap=plt.cm.gray)
    ax[i].axis('off')
    ax[i].set_title(f'thrs = {thr}')
plt.show()
```

OpenCV2 example

cv2.threshold(src, thresh, maxval, type[, dst]) → retval, dst
It returns the used threshold and the thresholded image.

```
thrs = [30,60,90,120]
fig, ax = plt.subplots(nrows = 1, ncols = len(thrs), figsize = (3*len(thrs),
3*len(thrs)))
for i, thr in enumerate(thrs):
    _, img = cv2.threshold(lena_gs, int(thr), 255, cv2.THRESH_BINARY)
    ax[i].imshow(img, cmap=plt.cm.gray)
    ax[i].axis('off')
    ax[i].set_title(f'thrs = {thr}')
plt.show()
```

EX. 2

```
def box_filter(kernel_size):
    ...
```

Create a $k \times k$ box filter.

Input parameters

kernel_size: int

Size of the squared kernel.

Output parameters

```

-----
Filter: np.array [k x k]
Box filter

To check
-----
kernel size must be an odd number
'''

    assert kernel_size % 2 != 0

    kernel = np.ones([kernel_size, kernel_size]) / kernel_size**2

    return kernel

def conv_2d(image, kernel):
    '''
    Apply a k x k filter to the input image.
    Input parameters
    -----
    image: np.array
    Input image.
    kernel: np.array [k x k]
    Filter to apply

    Output parameters
    -----
    Image: np.array

    To check
    -----
    Grayscale image is required
    '''

    assert len(image.shape) == 2

    kernel_size = kernel.shape[0]

    image = np.asarray(image, dtype=np.float32)

    output = np.zeros_like(image)

    kernel = np.flipud(np.fliplr(kernel))

    padding_size = (kernel_size - 1)//2
    img_padded = np.pad(image, (padding_size, padding_size))

    h,w = image.shape
    for row in range(h):

```

```

    for col in range(w):
        output[row,col] = (img_padded[row:row+kernel_size,
col:col+kernel_size]*kernel).sum()

    return output

```

EX. 3

```

def gaussian_2d(size=11, sigma=1):
    """
    2-D Gaussian Filter

    Input Parameters
    -----
    size: int
    Size of the squared kernel
    sigma: float
    Standard deviation

    Output Parameter
    -----
    filter: np.array
    2-D Gaussian Filter
    """

    # Check input parameters
    assert size % 2 != 0, "Kernel size must be an odd number."

    ax = np.linspace(-(size - 1) / 2., (size - 1) / 2., size)
    xx, yy = np.meshgrid(ax, ax)

    kernel = 1/(np.sqrt(2 * np.pi * np.square(sigma))) * np.exp(-0.5 *
(np.square(xx) + np.square(yy)) / np.square(sigma))

    return kernel / np.sum(kernel)

```

EX. 4

```

def median_filter(image, kernel_size, padding=True):
    """
    This function applies the median filter to the input image.

    Input Parameters
    -----
    image: np.array

```

Input grayscale image

kernel_size: int

Dimension of a squared kernel.

Padding: bool

If True, the input image is padded with zeros in order to have as output an image with the same spatial size of the input image.

If False, no padding is applied and the output image will be smaller than the input one.

Output Parameters

image: np.array

Filtered image

'''

Check inputs

assert len(image.shape) == 2

assert kernel_size % 2 != 0

image = np.asarray(image, dtype=np.float32)

h, w = image.shape

if padding:

output = np.zeros_like(image)

padding_size = (kernel_size - 1) // 2

img_padded = np.pad(image, (padding_size, padding_size))

for row in range(h):

for col in range(w):

output[row, col] = np.median(img_padded[row:row+kernel_size, col:col+kernel_size])

else:

out_h = h - kernel_size + 1

out_w = w - kernel_size + 1

output = np.zeros([out_h, out_w])

for row in range(out_h):

for col in range(out_w):

output[row, col] = np.median(image[row:row+kernel_size, col:col+kernel_size])

return output

EX. 5

```
lena_blur = cv2.GaussianBlur(lena_gs, (21,21), 10)
lena_details = cv2.addWeighted(lena_gs,1,lena_blur,-1,0)

fig, ax = plt.subplots(nrows = 2, ncols = 3, figsize = (15,12))
ax[0,0].imshow(lena_gs, cmap=plt.cm.gray)
ax[0,0].axis('off')
ax[0,0].set_title('Original')

ax[0,1].imshow(lena_blur, cmap=plt.cm.gray)
ax[0,1].axis('off')
ax[0,1].set_title('Blur')

ax[0,2].imshow(lena_details, cmap=plt.cm.gray)
ax[0,2].axis('off')
ax[0,2].set_title('Details')

for i,k in enumerate([1,2,3]):
    lena_sharp = cv2.addWeighted(lena_gs,1,lena_details,k,0)
    ax[1,i].imshow(lena_sharp, cmap=plt.cm.gray)
    ax[1,i].axis('off')
    ax[1,i].set_title(f'k={k}')
```