

Computer Vision & Cognitive Systems

SCQ5109806 - LM CS,DS,CYB,PD



Foundations: image filtering, edges

Prof. Lamberto Ballan

Summary

- Image sampling and quantization
- Image histograms
- Images as functions
- Filters (linear systems)
- Convolution and (cross-)correlation

Recap: linear filters

- We define a filter (system) as a unit that converts an input function $f[n, m]$ into an output (response) function $g[n, m]$, where (n, m) are the independent variables
- In the case of images, (n, m) represents the **spatial position** in the image

$$f[n, m] \rightarrow \boxed{\text{System } \mathcal{S}} \rightarrow g[n, m]$$

Recap: moving average filter

- 2D moving average over a 3×3 window of neighborhood (this is also referred as to *box filter*)

$$g[n, m] = \frac{1}{9} \sum_{k=n-1}^{n+1} \sum_{l=m-1}^{m+1} f[k, l]$$

$$= \frac{1}{9} \sum_{k=-1}^1 \sum_{l=-1}^1 f[n - k, m - l]$$

h

1	1	1
1	1	1
1	1	1

Image Kernel



$$(f * h)[n, m] = \sum_{k, l} f[k, l] h[n-k, m-l]$$

Convolution

Recap: moving average filter

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

	0	10	20	30	30				

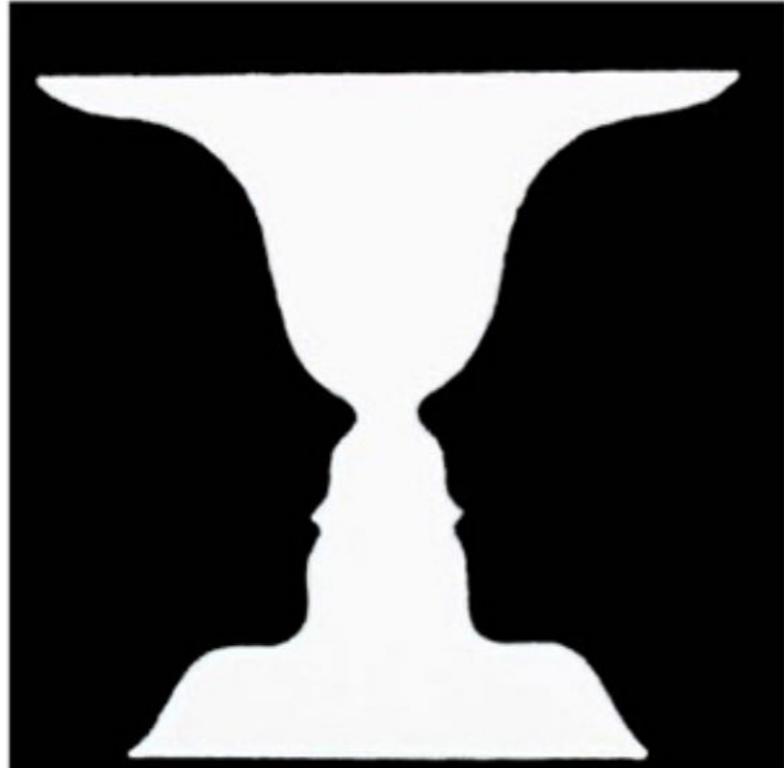
$$(f * h)[n, m] = \sum_{k, l} f[k, l] h[n-k, m-l]$$

Convolution

Image filtering recap

- Compute a function of the local neighborhood at each pixel in the image
 - ▶ The function is specified by a “filter” saying how to combine values from neighbors
- Applications of image filtering:
 - ▶ extract information (edges, corners, blobs, ...)
 - ▶ detect patterns (template matching)
 - ▶ de-noising, super-resolution, in-painting

What's an edge?



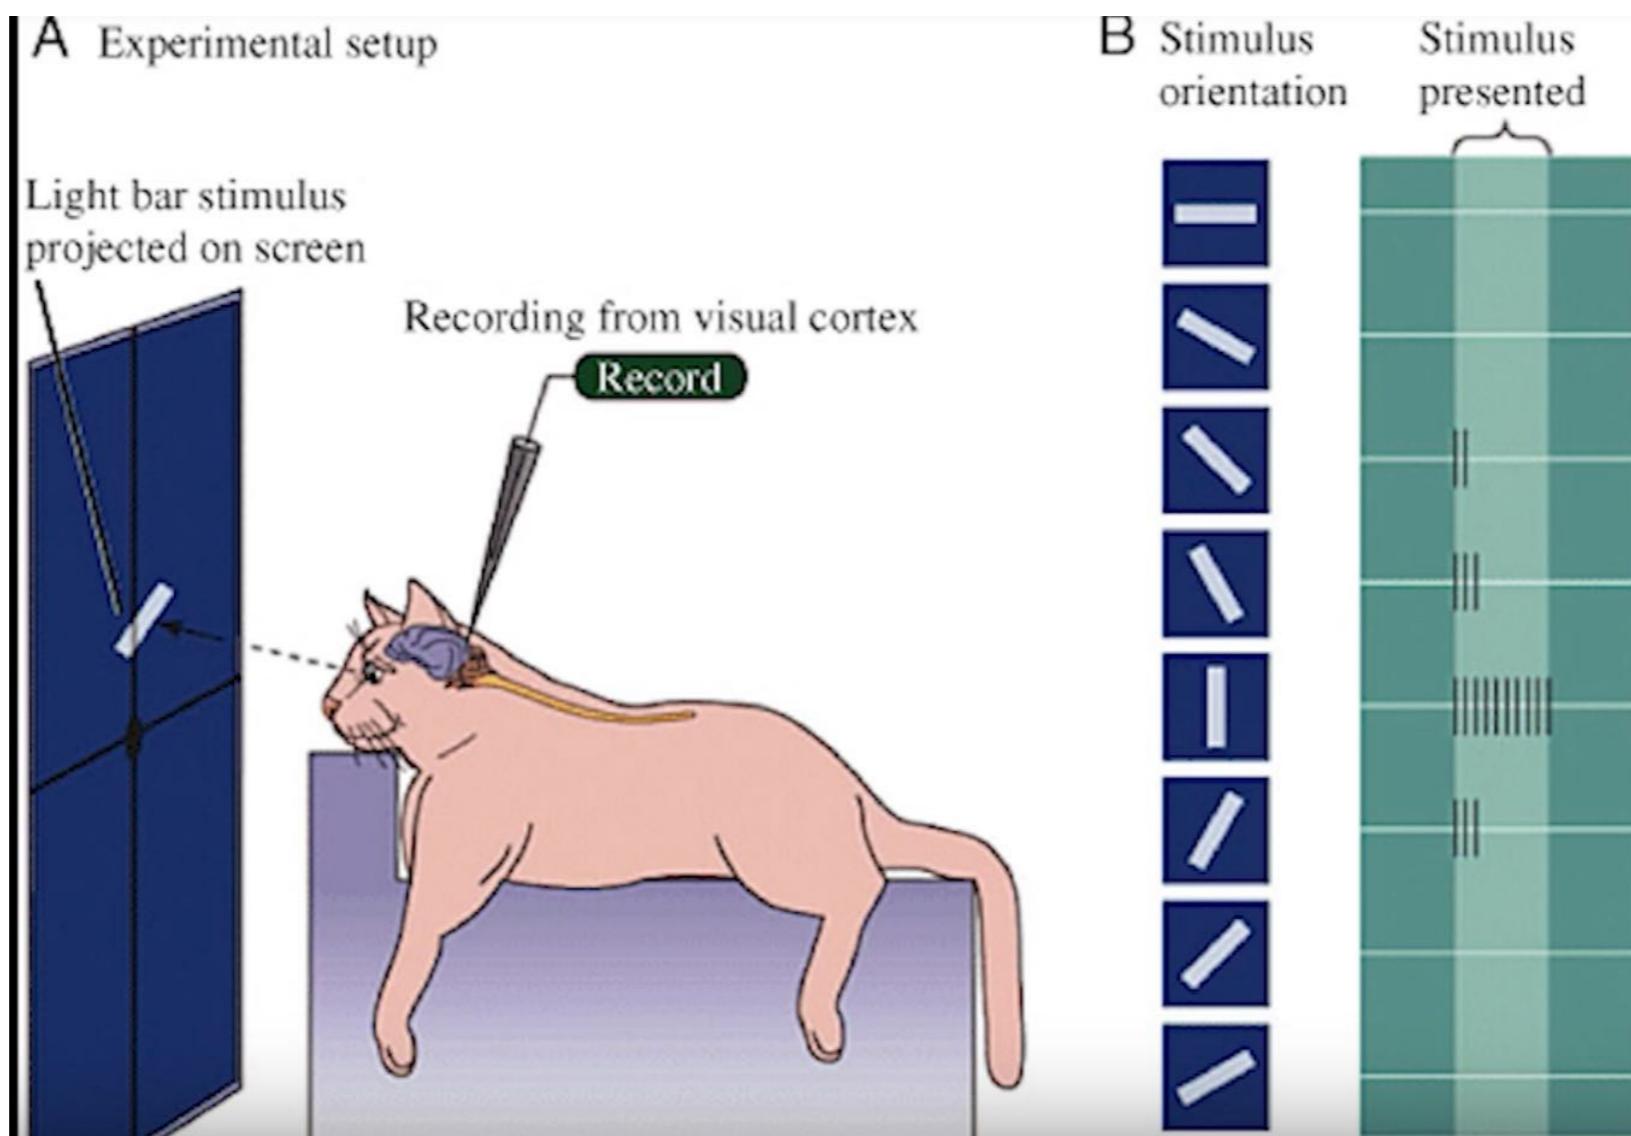
Cave painting (Chauvet, France, ~30,000 BC)



Shen Zhou's "Poet on a mountain top" (China, ~1,500 AD)

What's an edge?

- We know edges are special from human (mammalian) vision studies

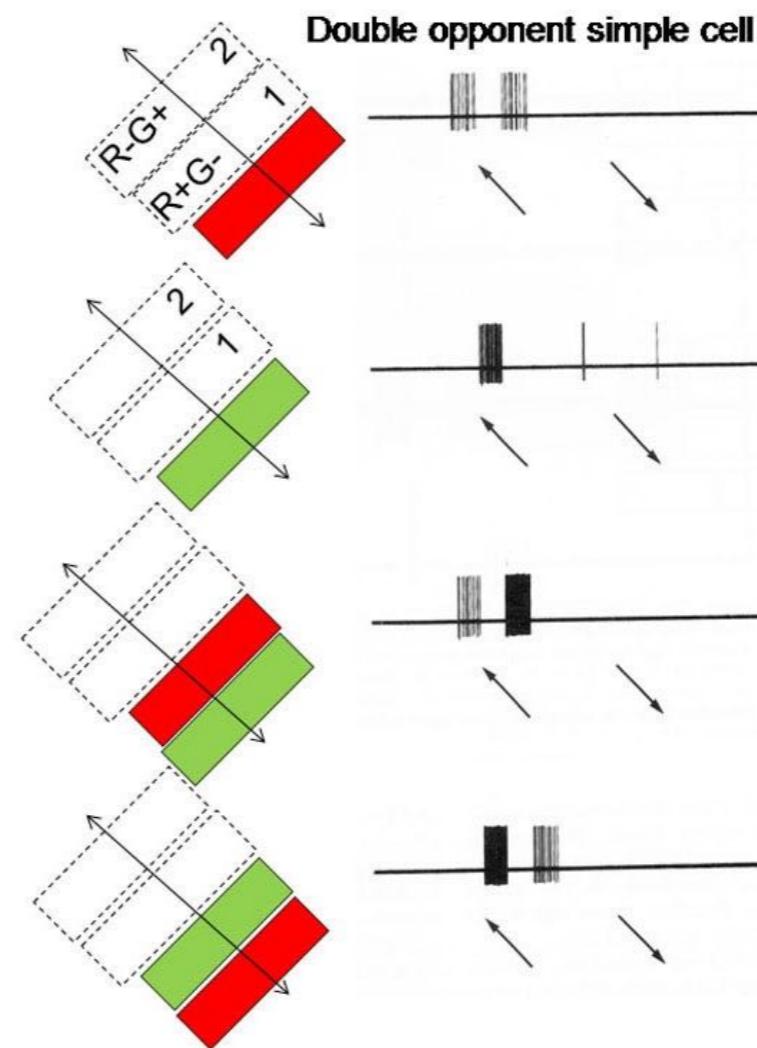


Hubel and Wiesel, 1960s

What's an edge?

- We know edges are special from human (mammalian) vision studies

Hubel and Wiesel, 1960s



Edge detection

- **Goal:** identify sudden changes (discontinuities) in an image
 - ▶ Intuitively, most semantic and shape information from the image can be encoded in the edges
 - ▶ More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)

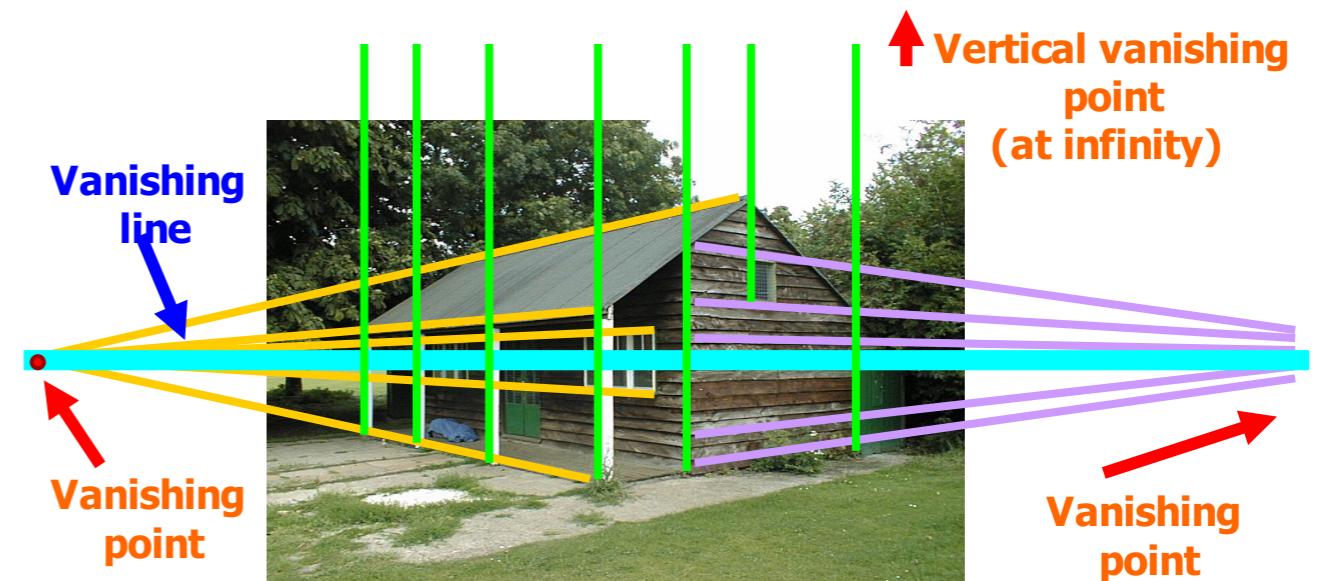


Why do we care about edges?

- Extract information, recognize objects

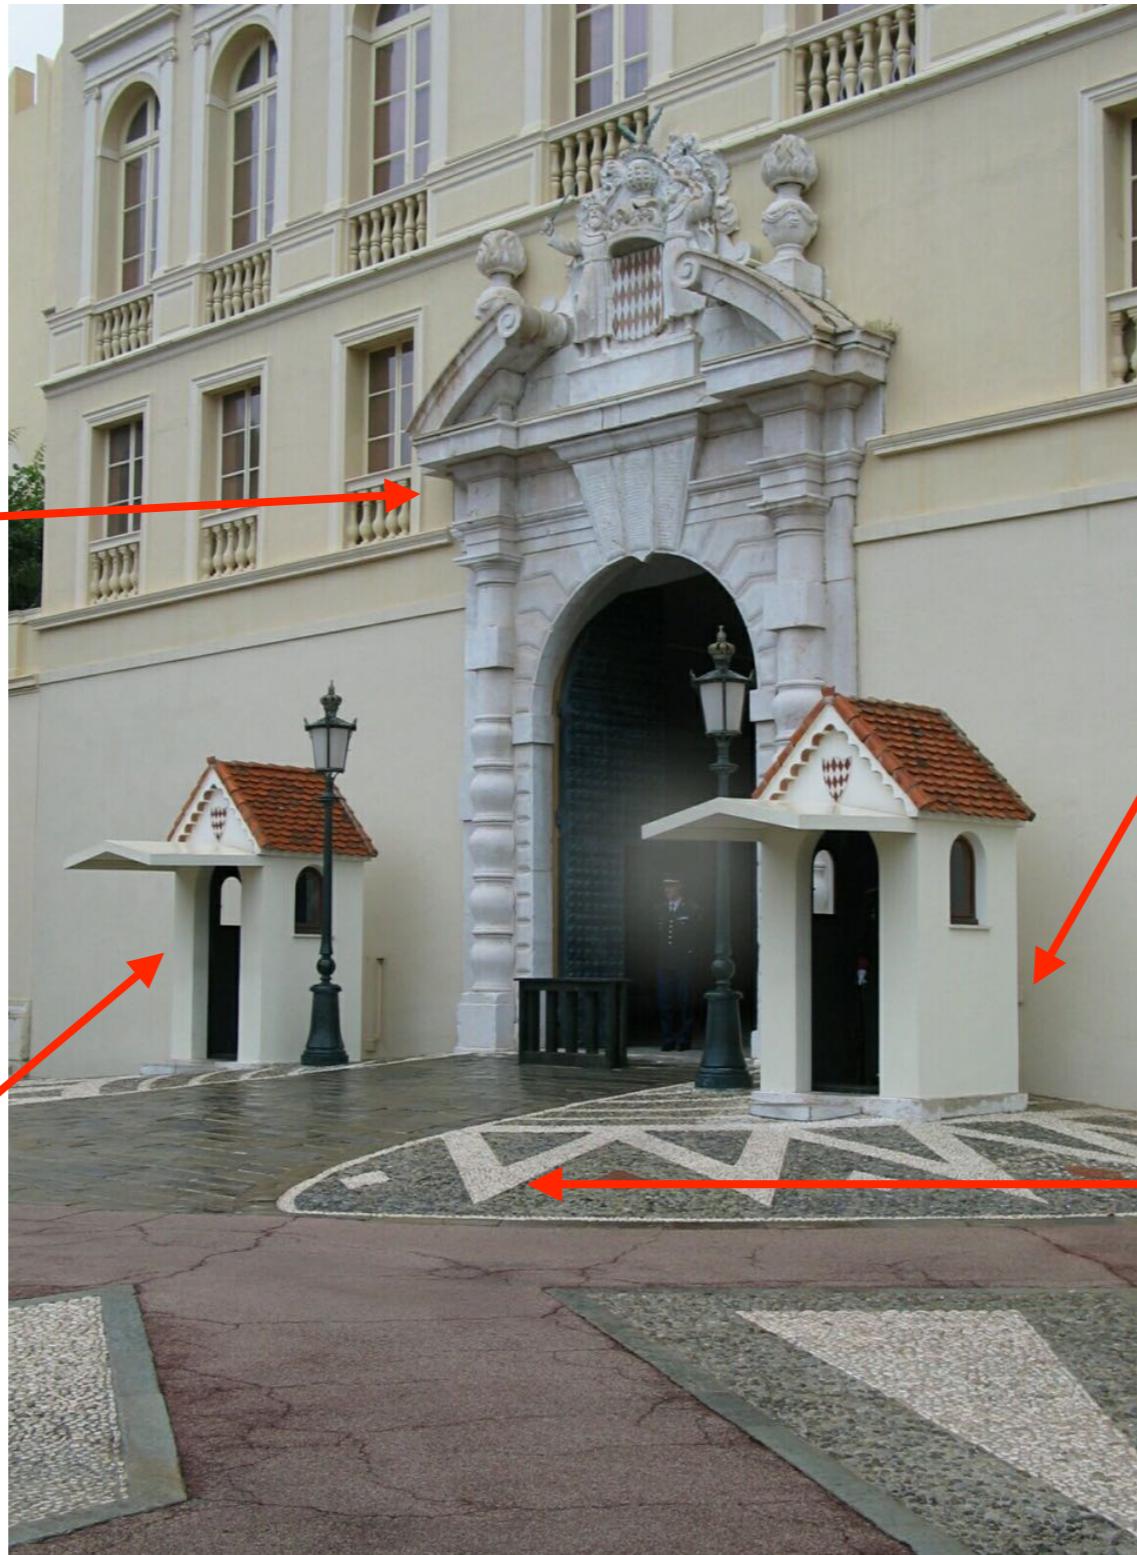


- Recover geometry and viewpoint

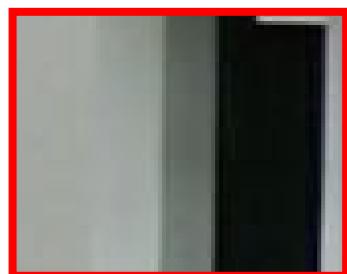


What causes an edge?

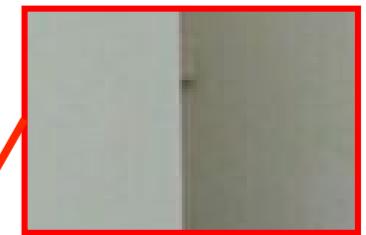
Illumination discontinuity:
cast shadows



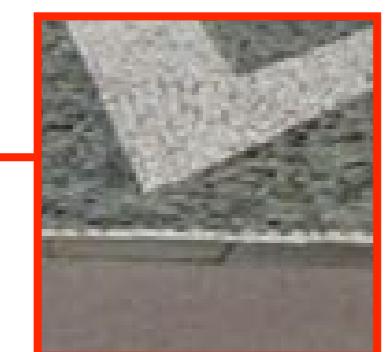
Change in surface
orientation: shape



Depth discontinuity:
object boundary

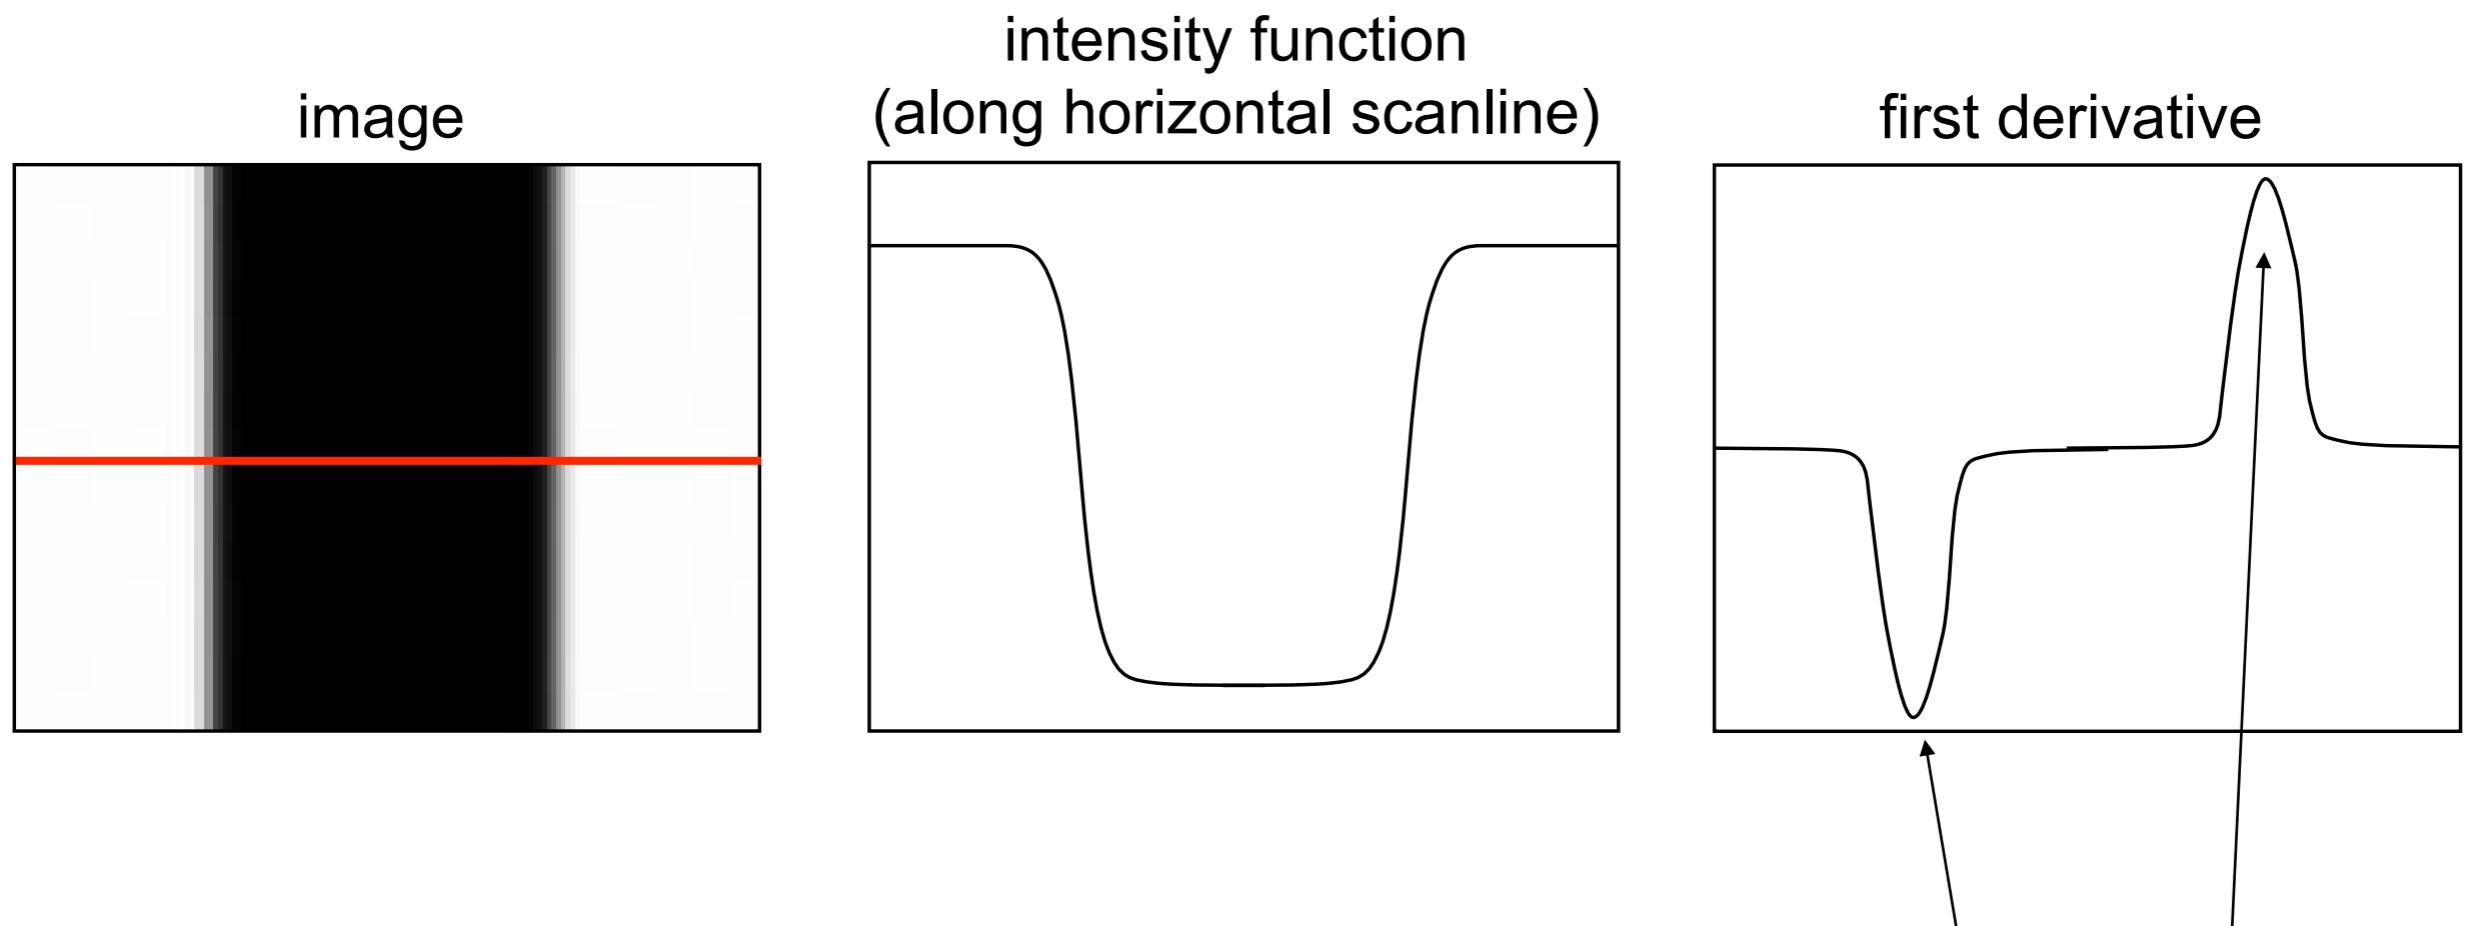


Surface color
discontinuity



Edges and image gradients

- An edge is a place of rapid change in the image intensity function



This could be the sketch of an edge detection algorithm

edges correspond to extrema of derivative

Derivatives: a quick recap

- Derivative in 1D:

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x) = f_x \quad \longrightarrow \quad \boxed{\begin{array}{l} A \text{ simple example:} \\ \\ y = x^2 + x^4 \\ \frac{dy}{dx} = 2x + 4x^3 \end{array}}$$

- Discrete derivative in 1D:

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x)$$

$$\frac{df}{dx} = \frac{f(x) - f(x - 1)}{1} = f'(x)$$

$$\frac{df}{dx} = f(x) - f(x - 1) = f'(x)$$

we can approximate using finite differences

Derivatives: a quick recap

- Types of Discrete derivative in 1D:

- Backward: $\frac{df}{dx} = f(x) - f(x-1) = f'(x)$ [0 1 -1] *Filter*
- Forward: $\frac{df}{dx} = f(x) - f(x+1) = f'(x)$ [-1 1 0]
- Central: $\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$ [1 0 -1]

Derivatives: a quick recap

- 1D discrete derivative example:

$$f(x) = \begin{matrix} -1 & 1 & 0 \\ 10 & 15 & 10 \end{matrix} \quad 10 \ 25 \ 20 \ 20 \ 20$$
$$f'(x) = \begin{matrix} 0 & 5 & -5 & 0 & 15 & -5 & 0 & 0 \end{matrix}$$


- What's the derivative filter used in this example?

Backward

Forward

Central

Filter [0 1 -1]

Derivatives: a quick recap

- 1D discrete derivative example:

$$f(x) = \begin{matrix} 0 \\ 10 & 15 & 10 \end{matrix} \quad \begin{matrix} 10 & 25 & 20 & 20 & 20 \end{matrix}$$
$$f'(x) = \begin{matrix} 0 & 5 & 0 & -15 & 5 & 0 & 0 & 0 \end{matrix}$$

↓

- What's the derivative filter used in this example?

Backward

Forward

Central

Filter [-1 1 0]

2D discrete derivative: image gradient



Given function

$$f(x, y)$$

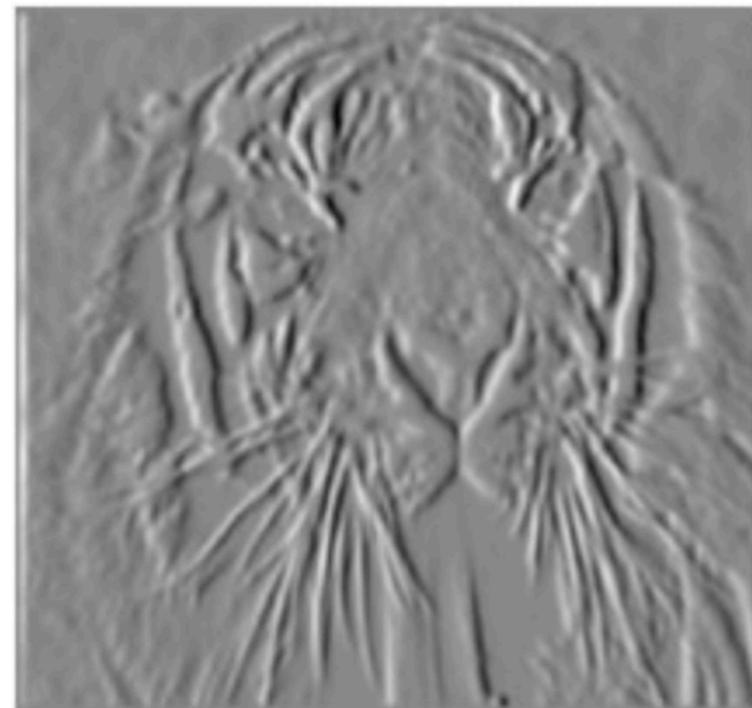
Gradient vector

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} \\ \frac{\partial f(x, y)}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$$

Which one is the gradient in the x-direction? How about y-direction?

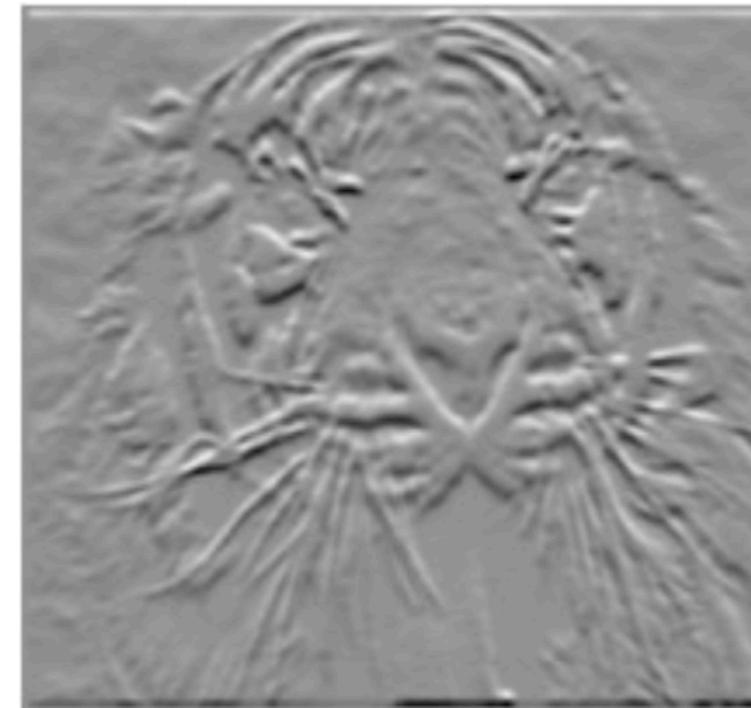
$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

x-direction



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

y-direction



2D discrete derivate example

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \boxed{\text{Filter}} = ?$$
$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

What happens when we apply this filter?

2D discrete derivate example

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \text{Filter} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Filter

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix} =$$

2D discrete derivate example

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \boxed{\text{Filter}} = ?$$

Filter

$$\frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

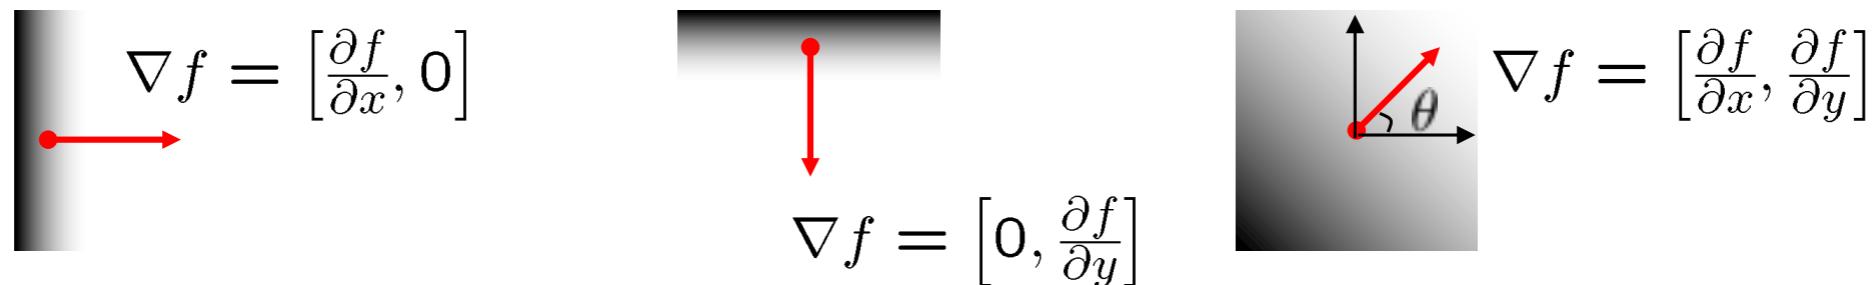
What happens when we apply this filter?

2D discrete derivate example

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \text{Filter} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Image gradient

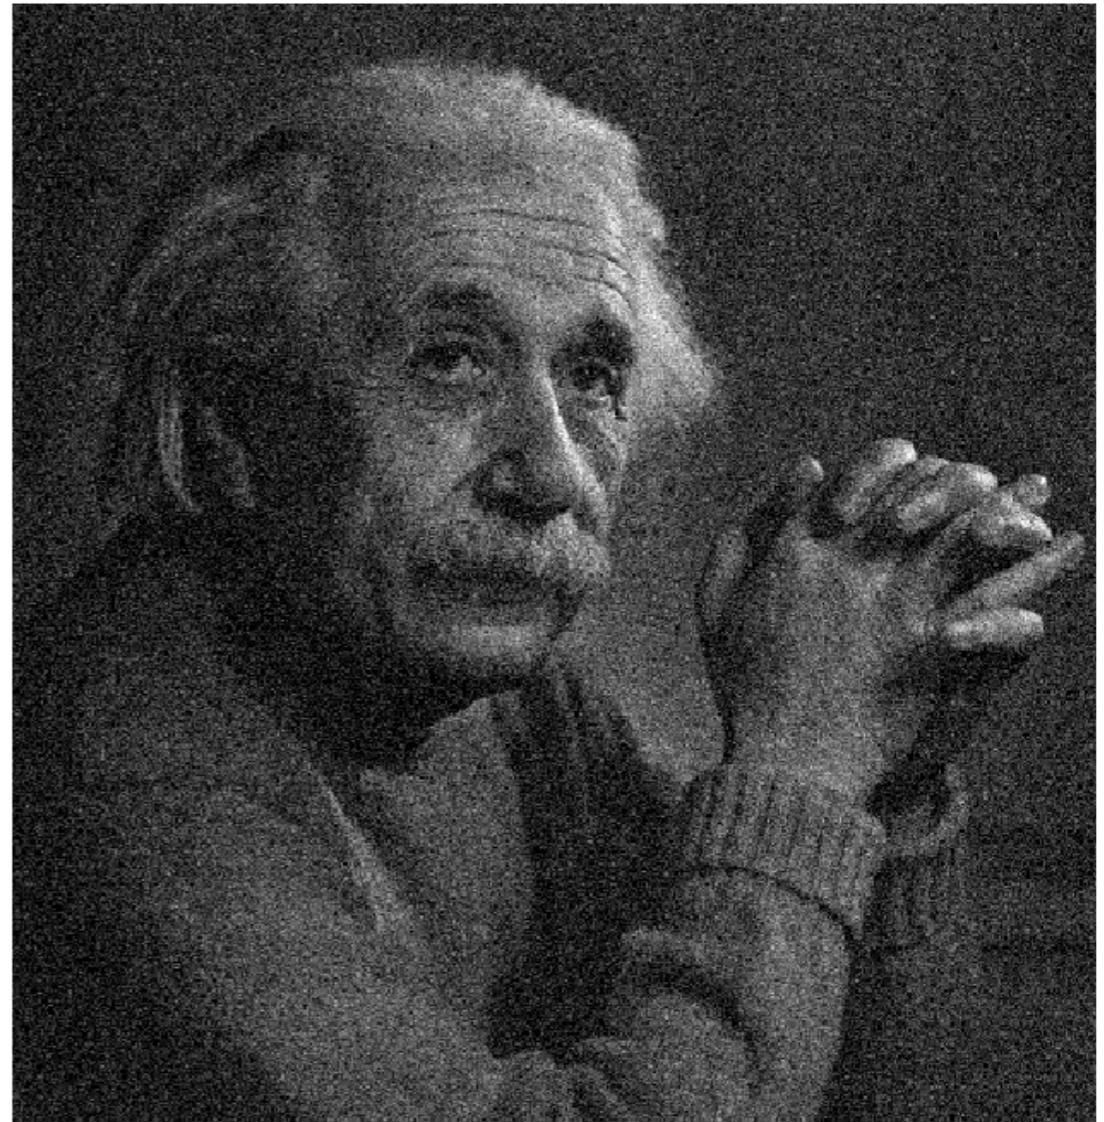
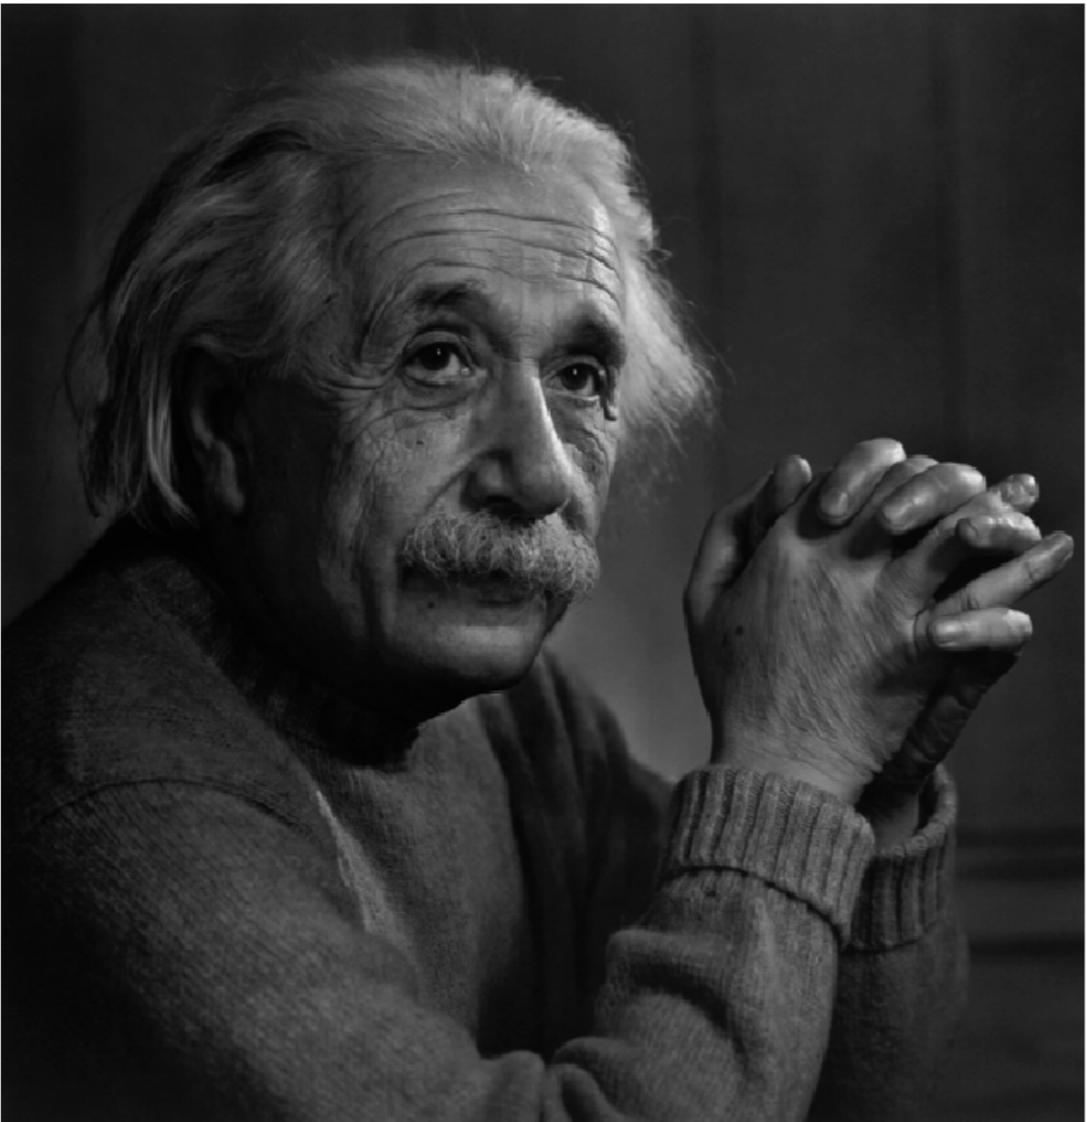
- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$



- The gradient vector points in the direction of most rapid increase in intensity
- The gradient direction is given by: $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$
- The *edge strength* is given by the gradient magnitude:

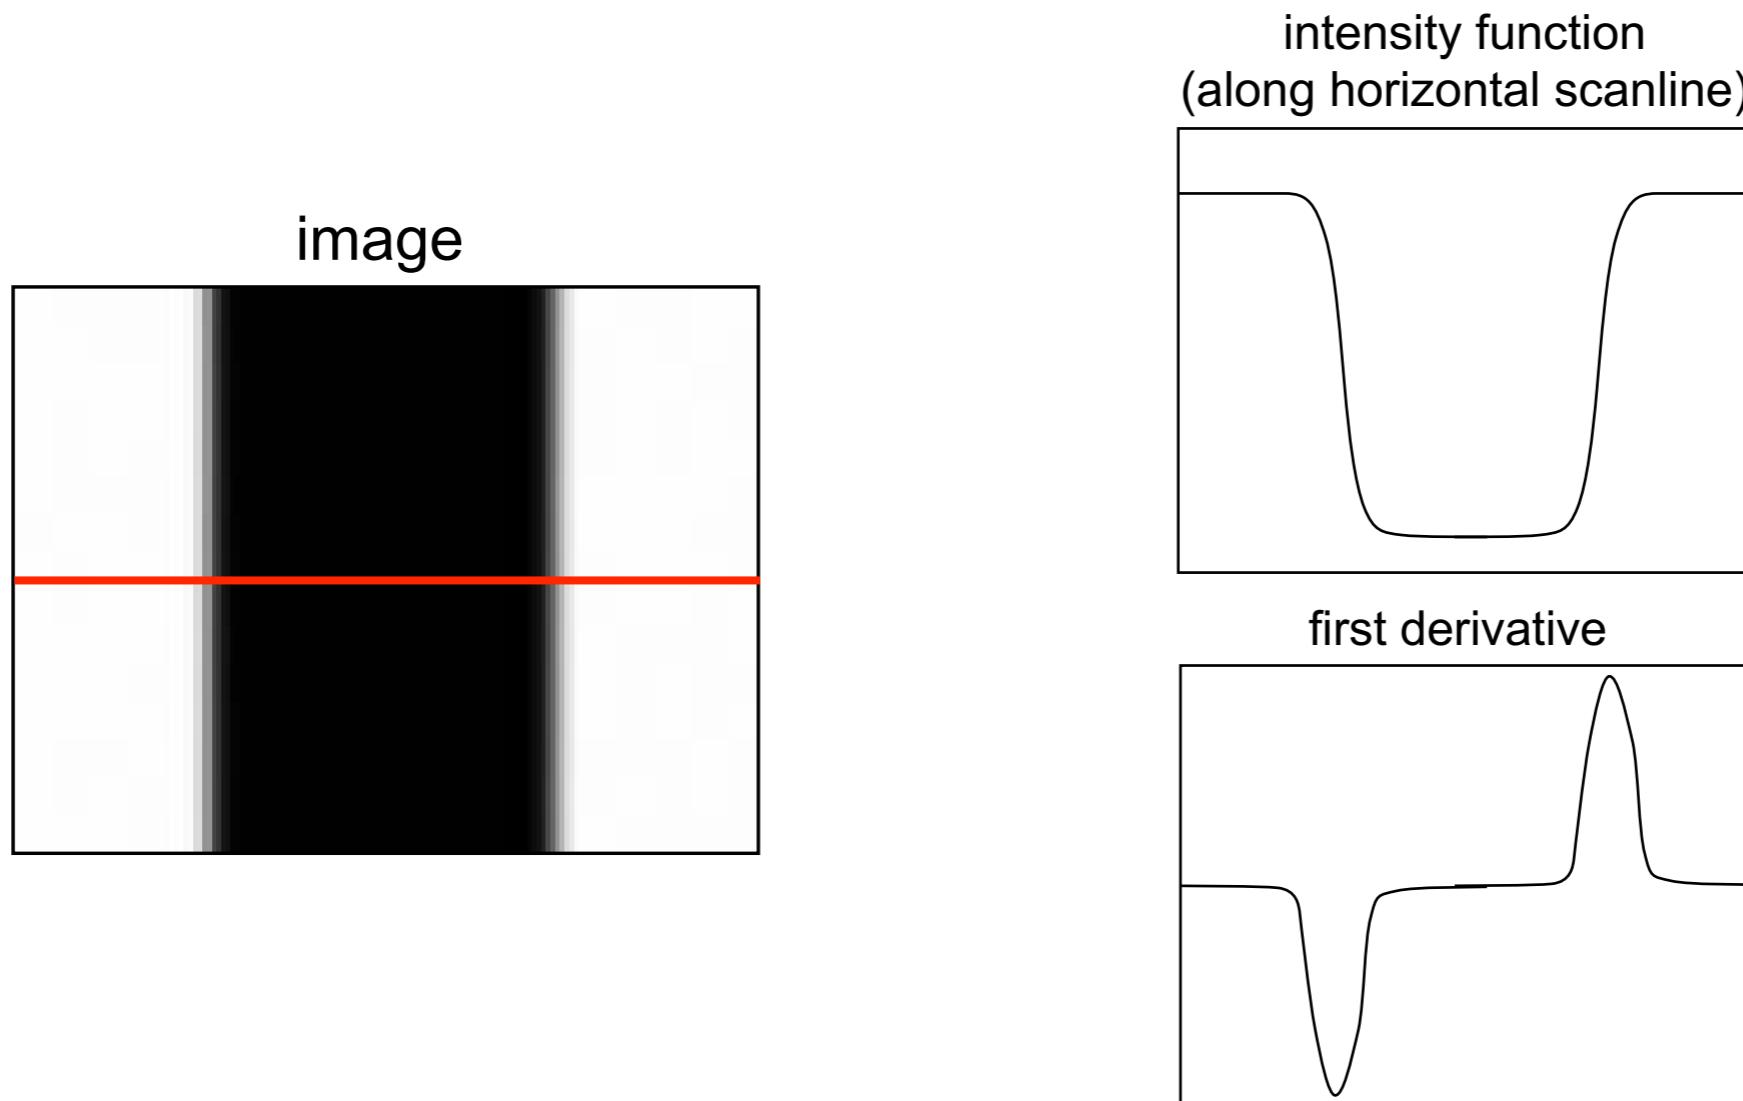
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Effects of noise



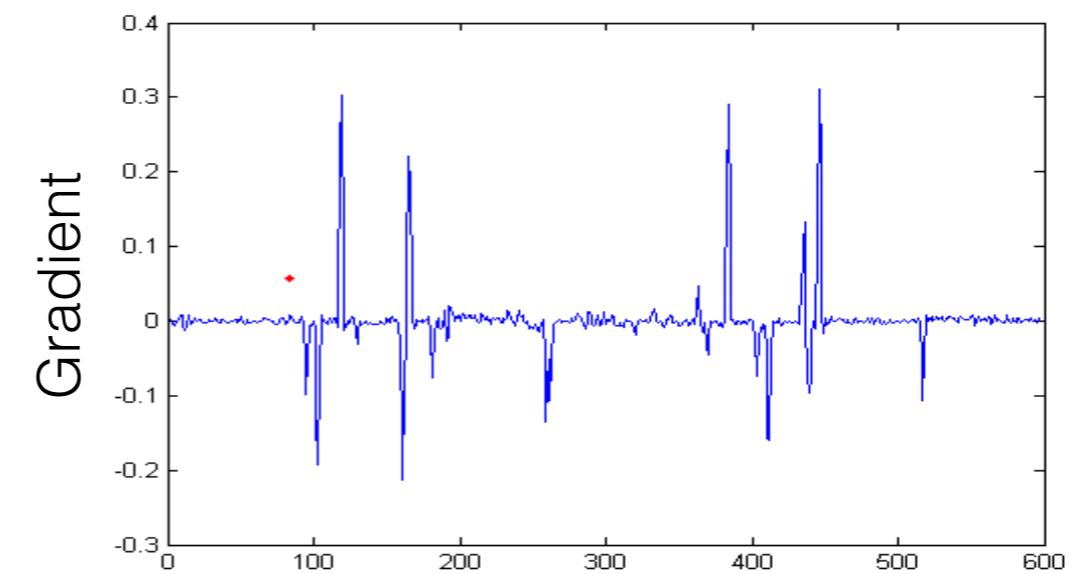
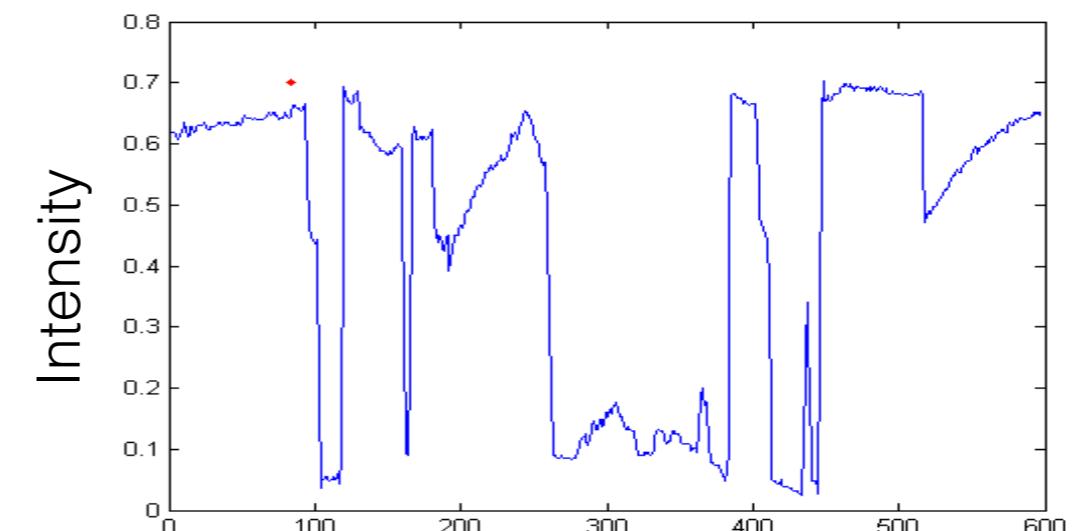
Intensity profile

- Consider a single row or column of the image
 - ▶ Plot intensity as a function of position gives a signal



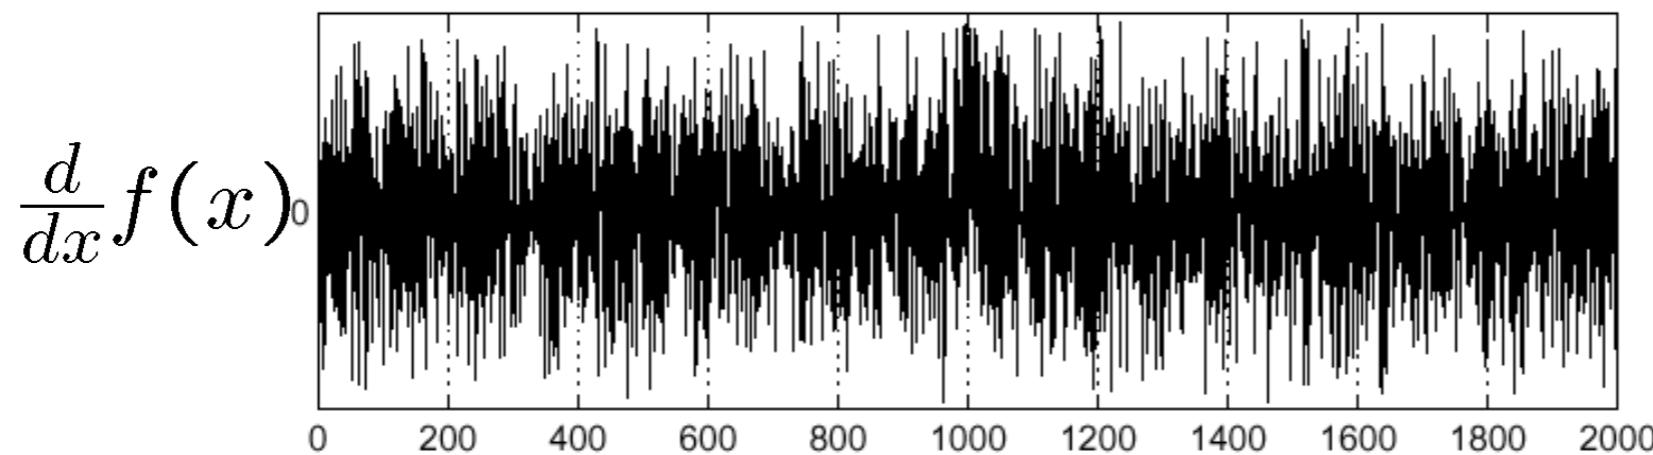
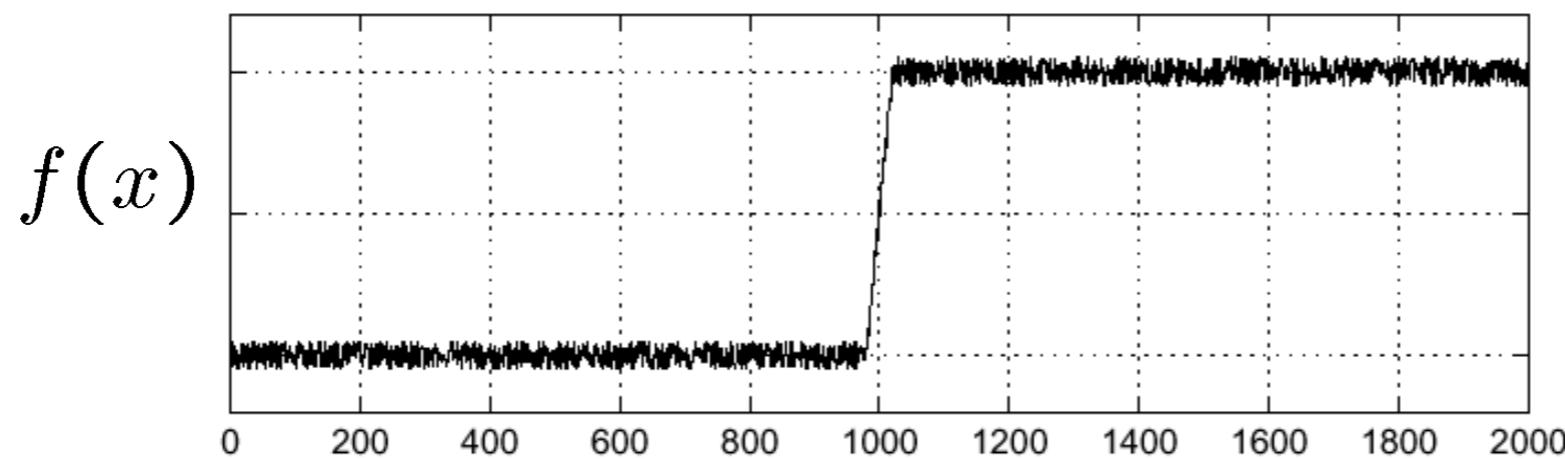
Intensity profile

- Consider a single row or column of the image
 - ▶ Plot intensity as a function of position gives a signal



Effects of noise

- Consider a single row or column of the image
 - ▶ Plotting intensity as a function of position gives a signal

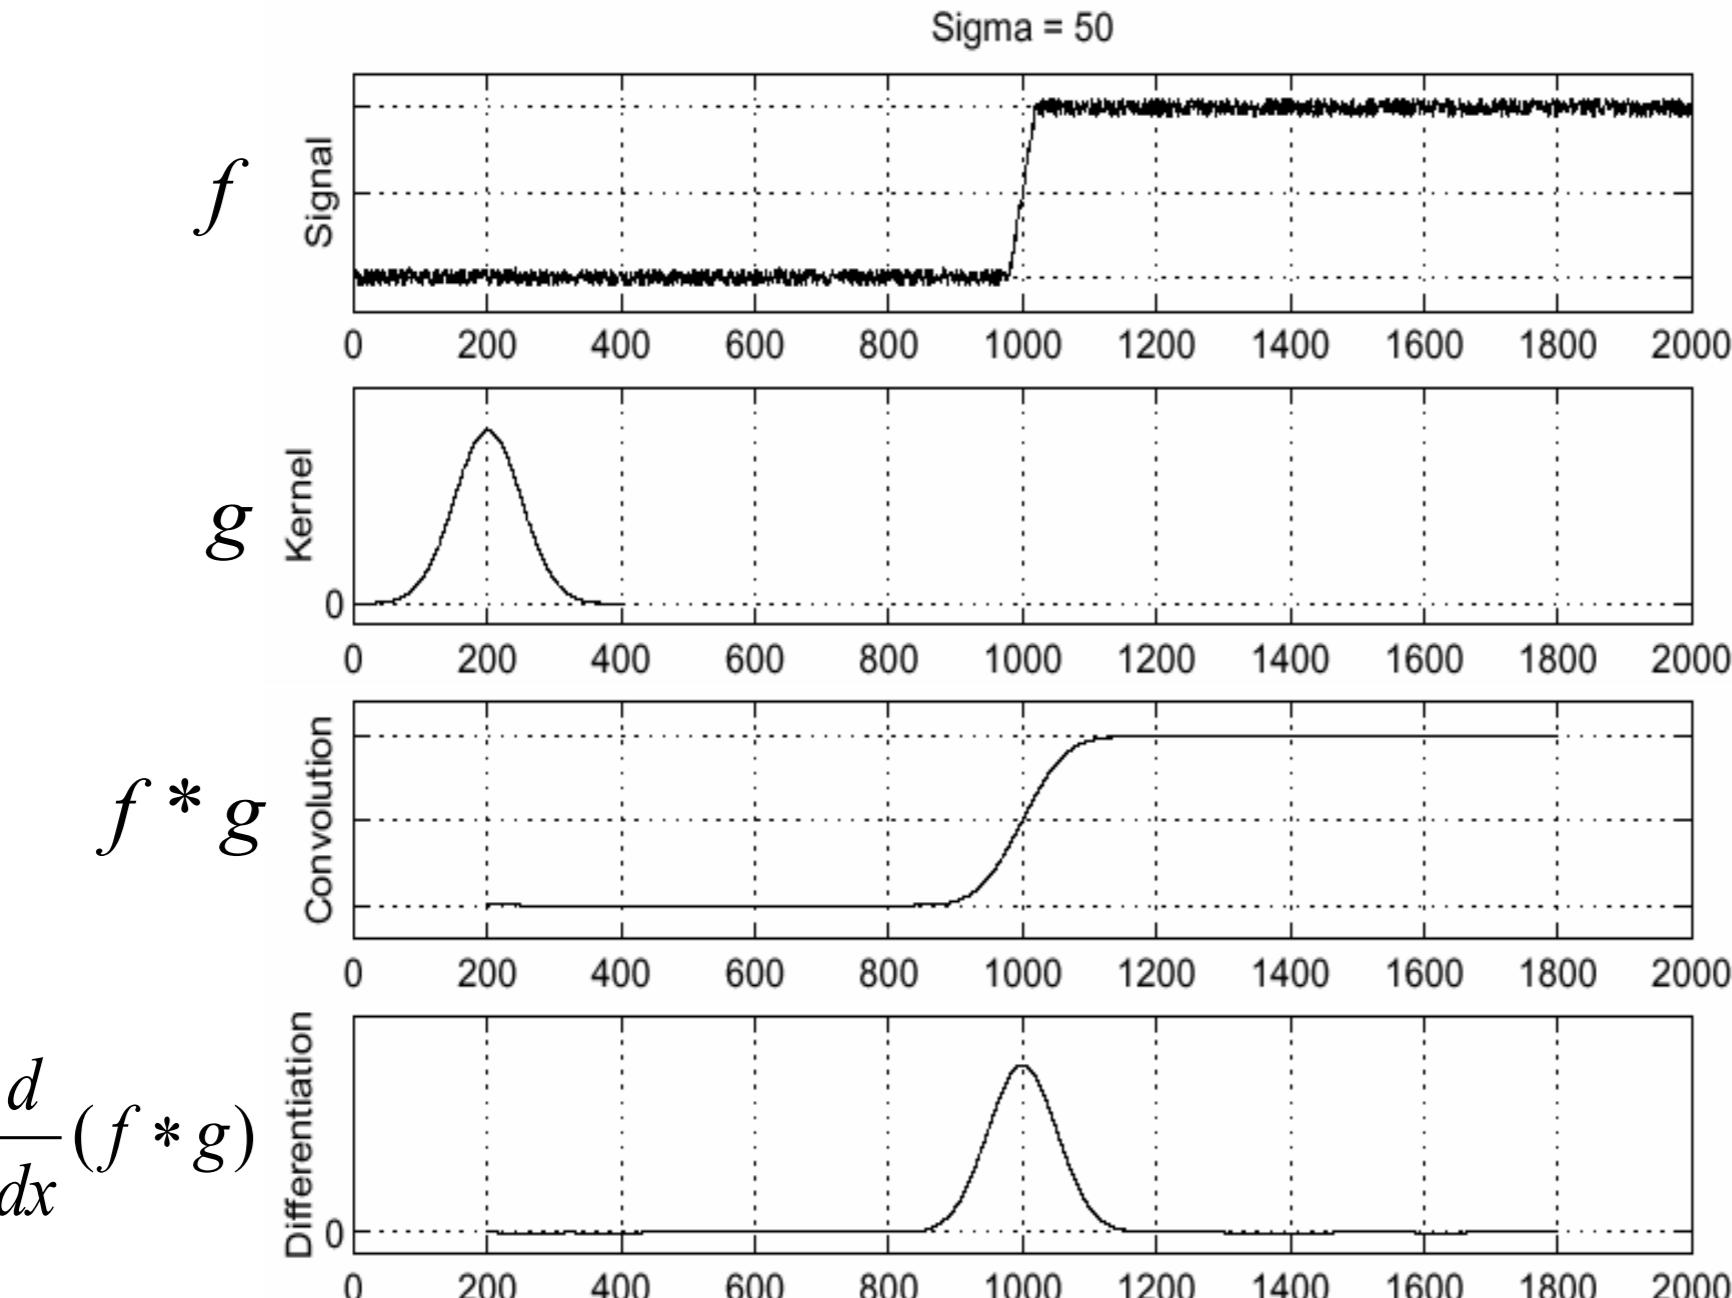


Where is the edge?

Effects of noise

- Finite difference filters respond strongly to noise
 - ▶ Image noise results in pixels that look very different from their neighbors
 - ▶ The larger the noise, the stronger the response
- What is to be done?
 - ▶ Smoothing the image should help, by forcing pixels different to their neighbors (=noise pixels?) to look more like neighbors

Solution: smooth first



Where is the
edge?

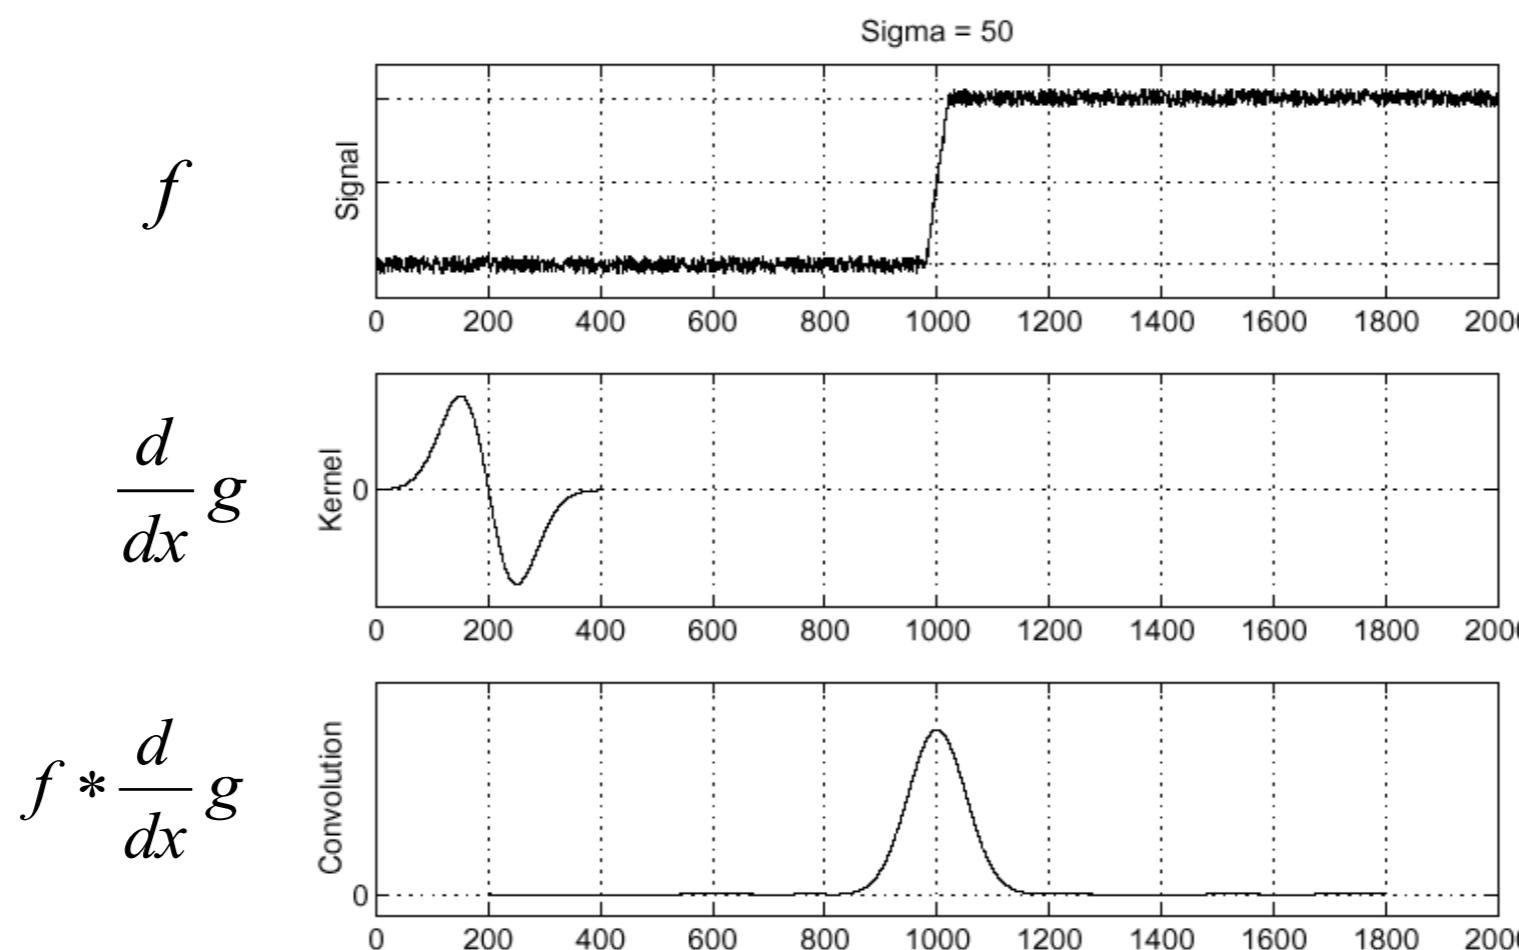
Look for peaks in $\frac{d}{dx}(f * g)$

Derivative theorem of convolution

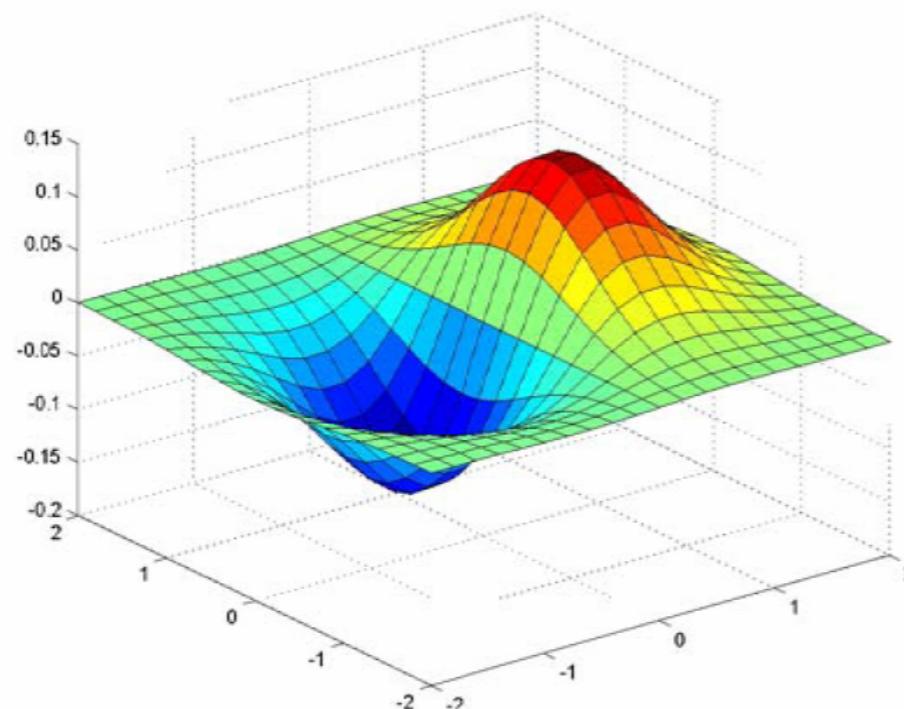
- Differential property of convolution:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

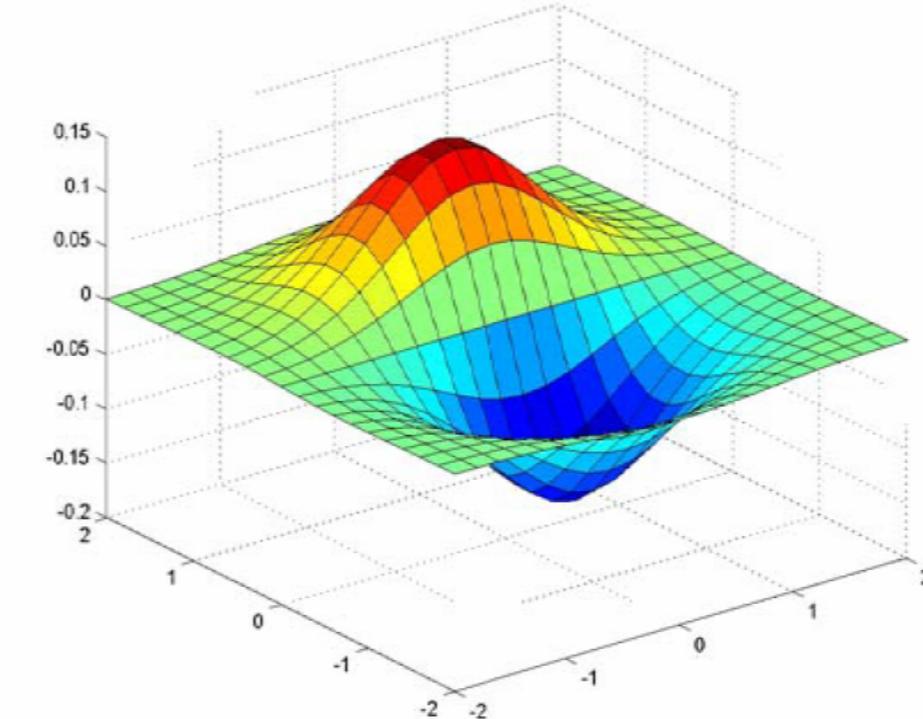
- ▶ This is a useful property and save us one operation:



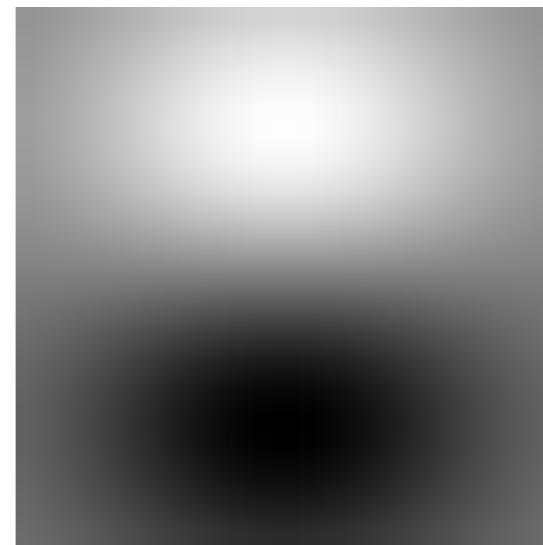
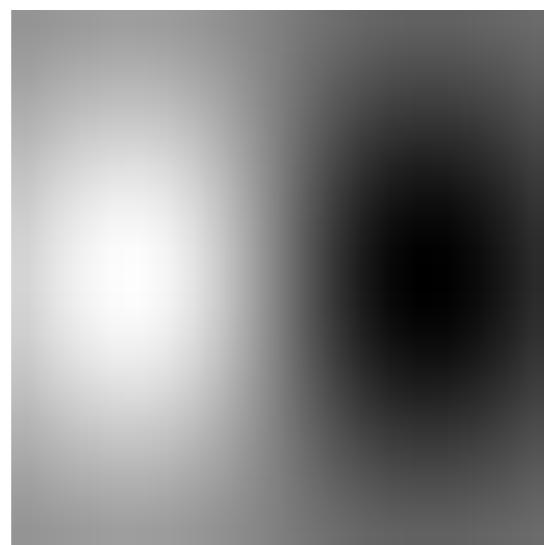
Derivative of Gaussian filter



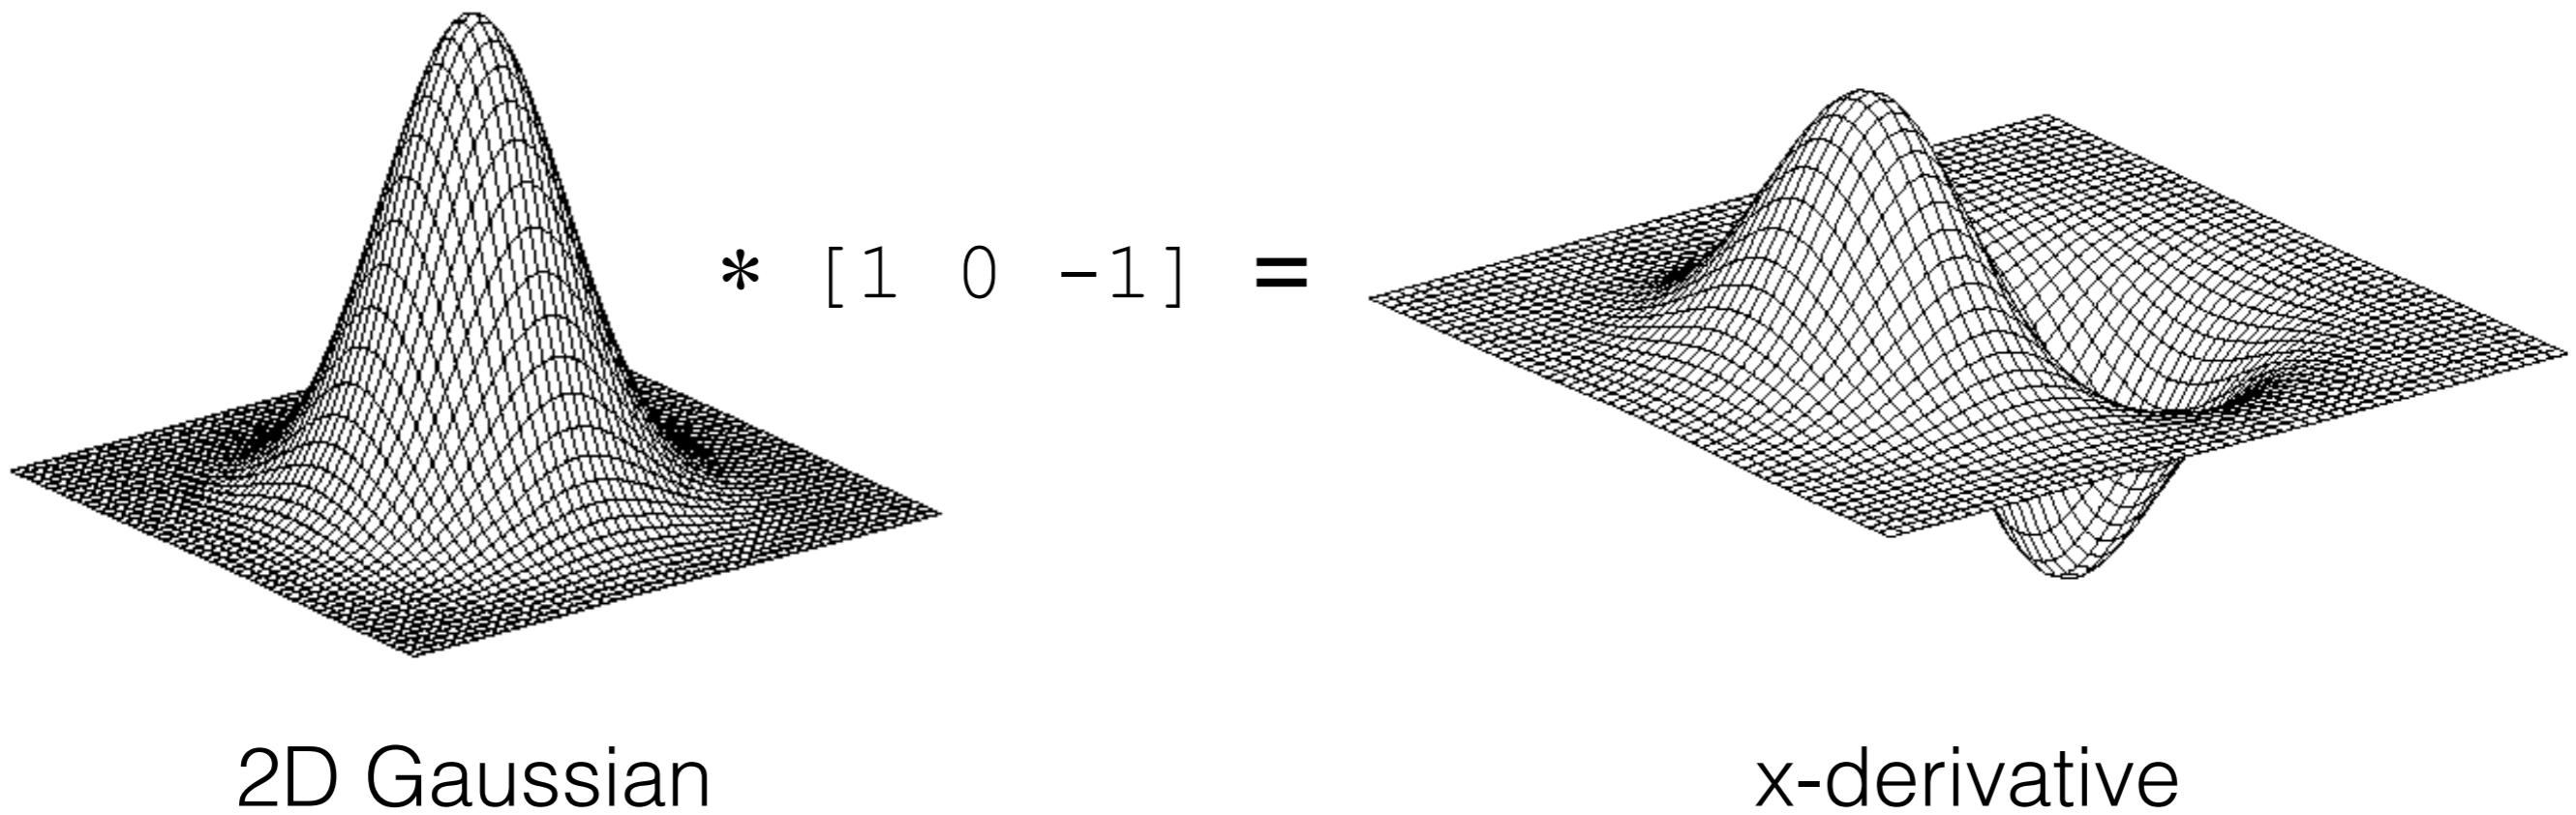
x-direction



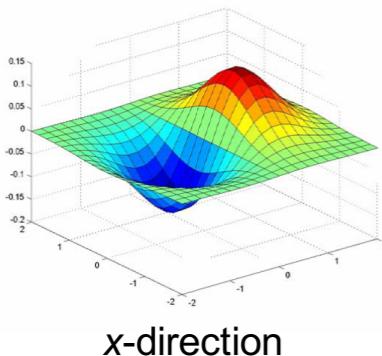
y-direction



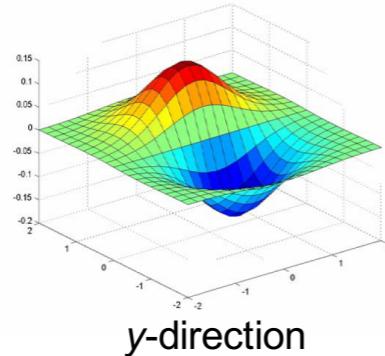
Derivative of Gaussian filter



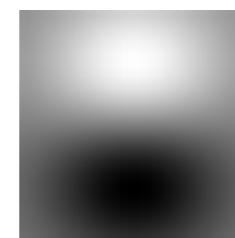
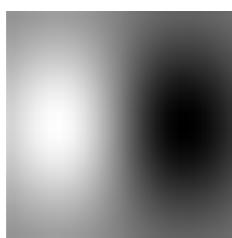
Derivative of Gaussian filter



x-direction



y-direction

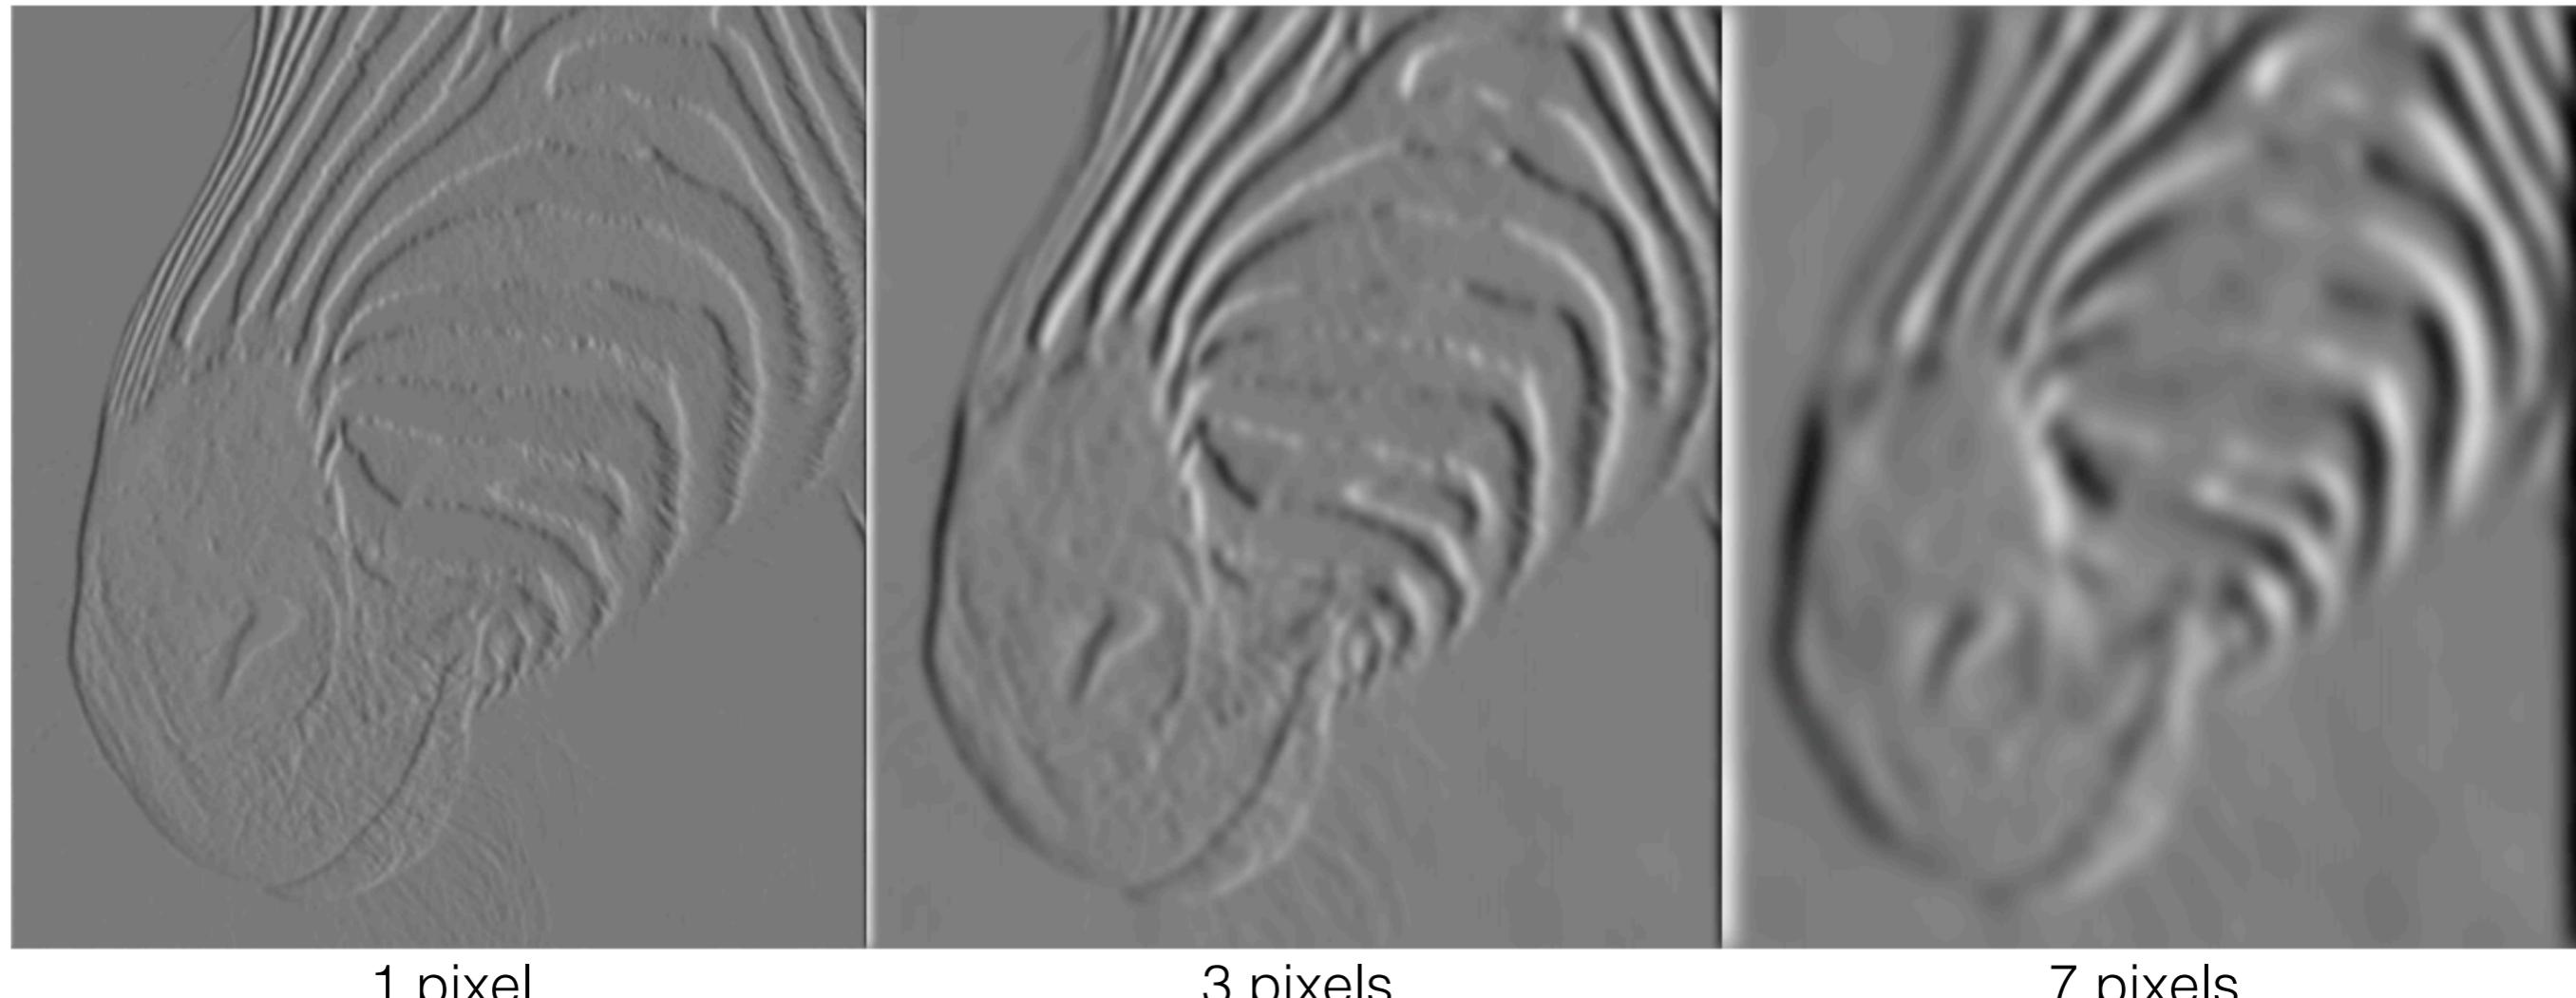


Note: separability of the Gaussian Filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

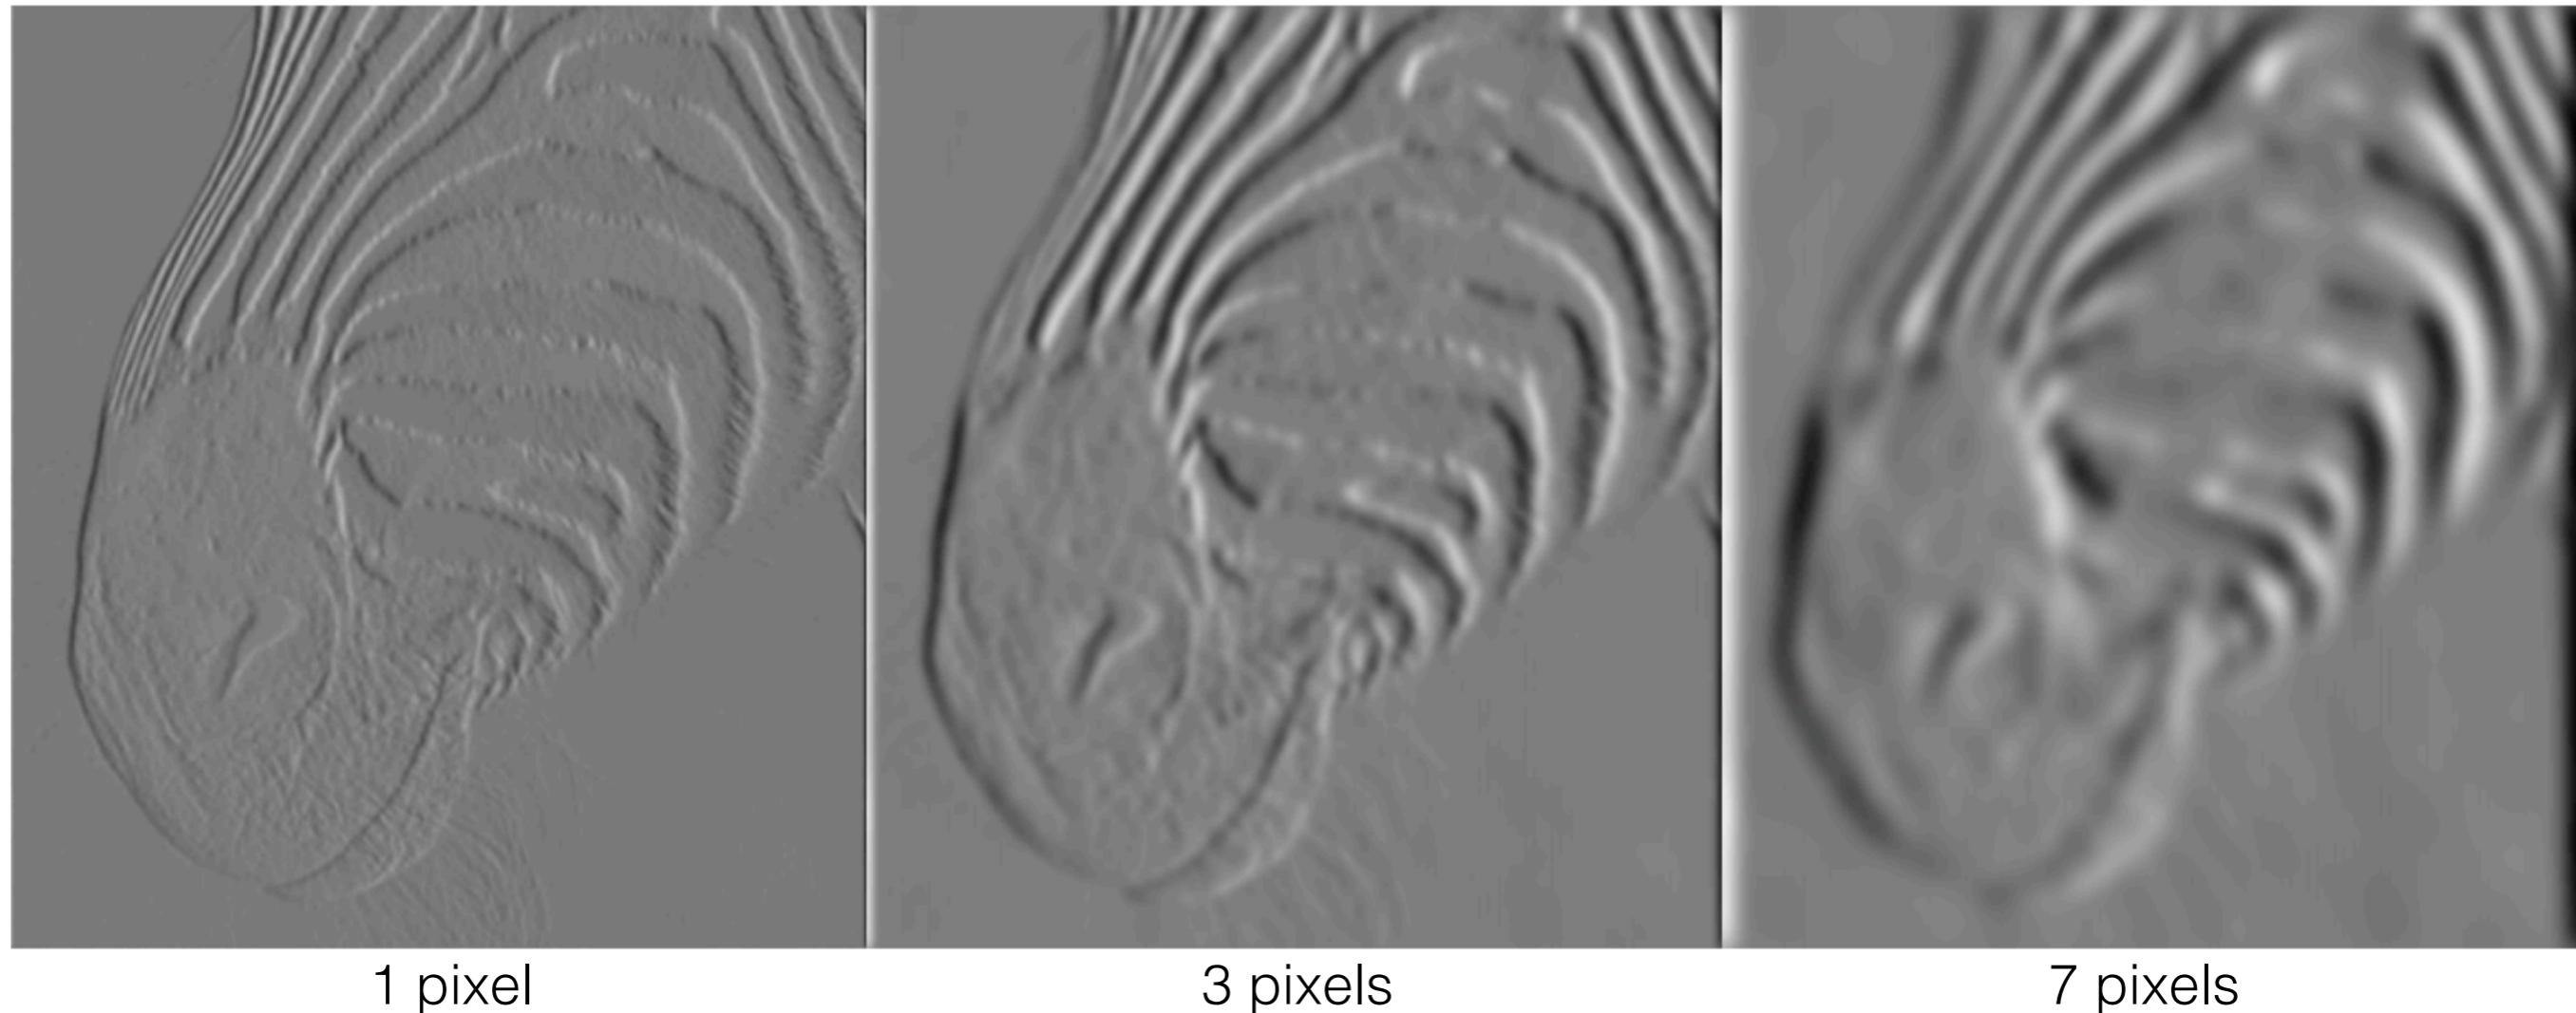
- ▶ The 2D Gaussian can be expressed as the product of 2 functions (one a function of x and the other a function of y)
- ▶ In this case the two functions are the (identical) 1D Gaussian

Scale of Gaussian derivative filter



- ▶ Note: σ is the *scale (width/spread)* of the Gaussian kernel
- ▶ Smoothed derivative remove noise, but blurs edges

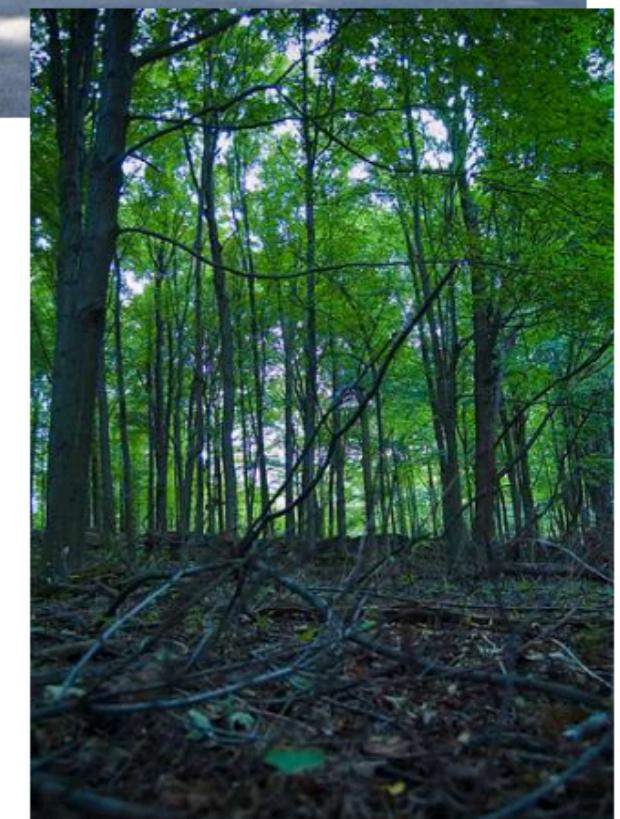
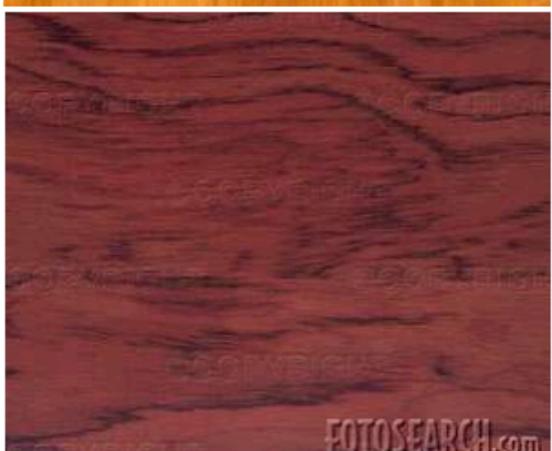
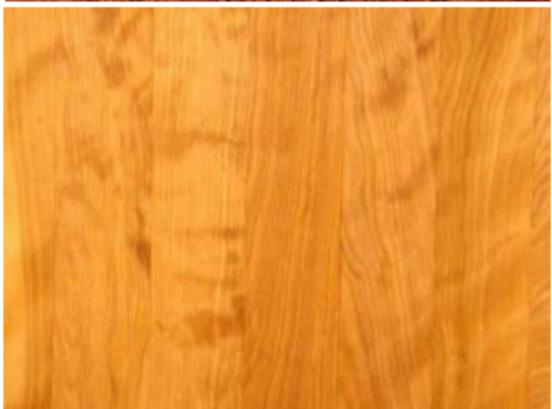
Scale of Gaussian derivative filter



- The apparent structures differ depending on scale σ
 - Larger values: larger scale edges detected
 - Smaller values: finer features detected

So, what scale to choose?

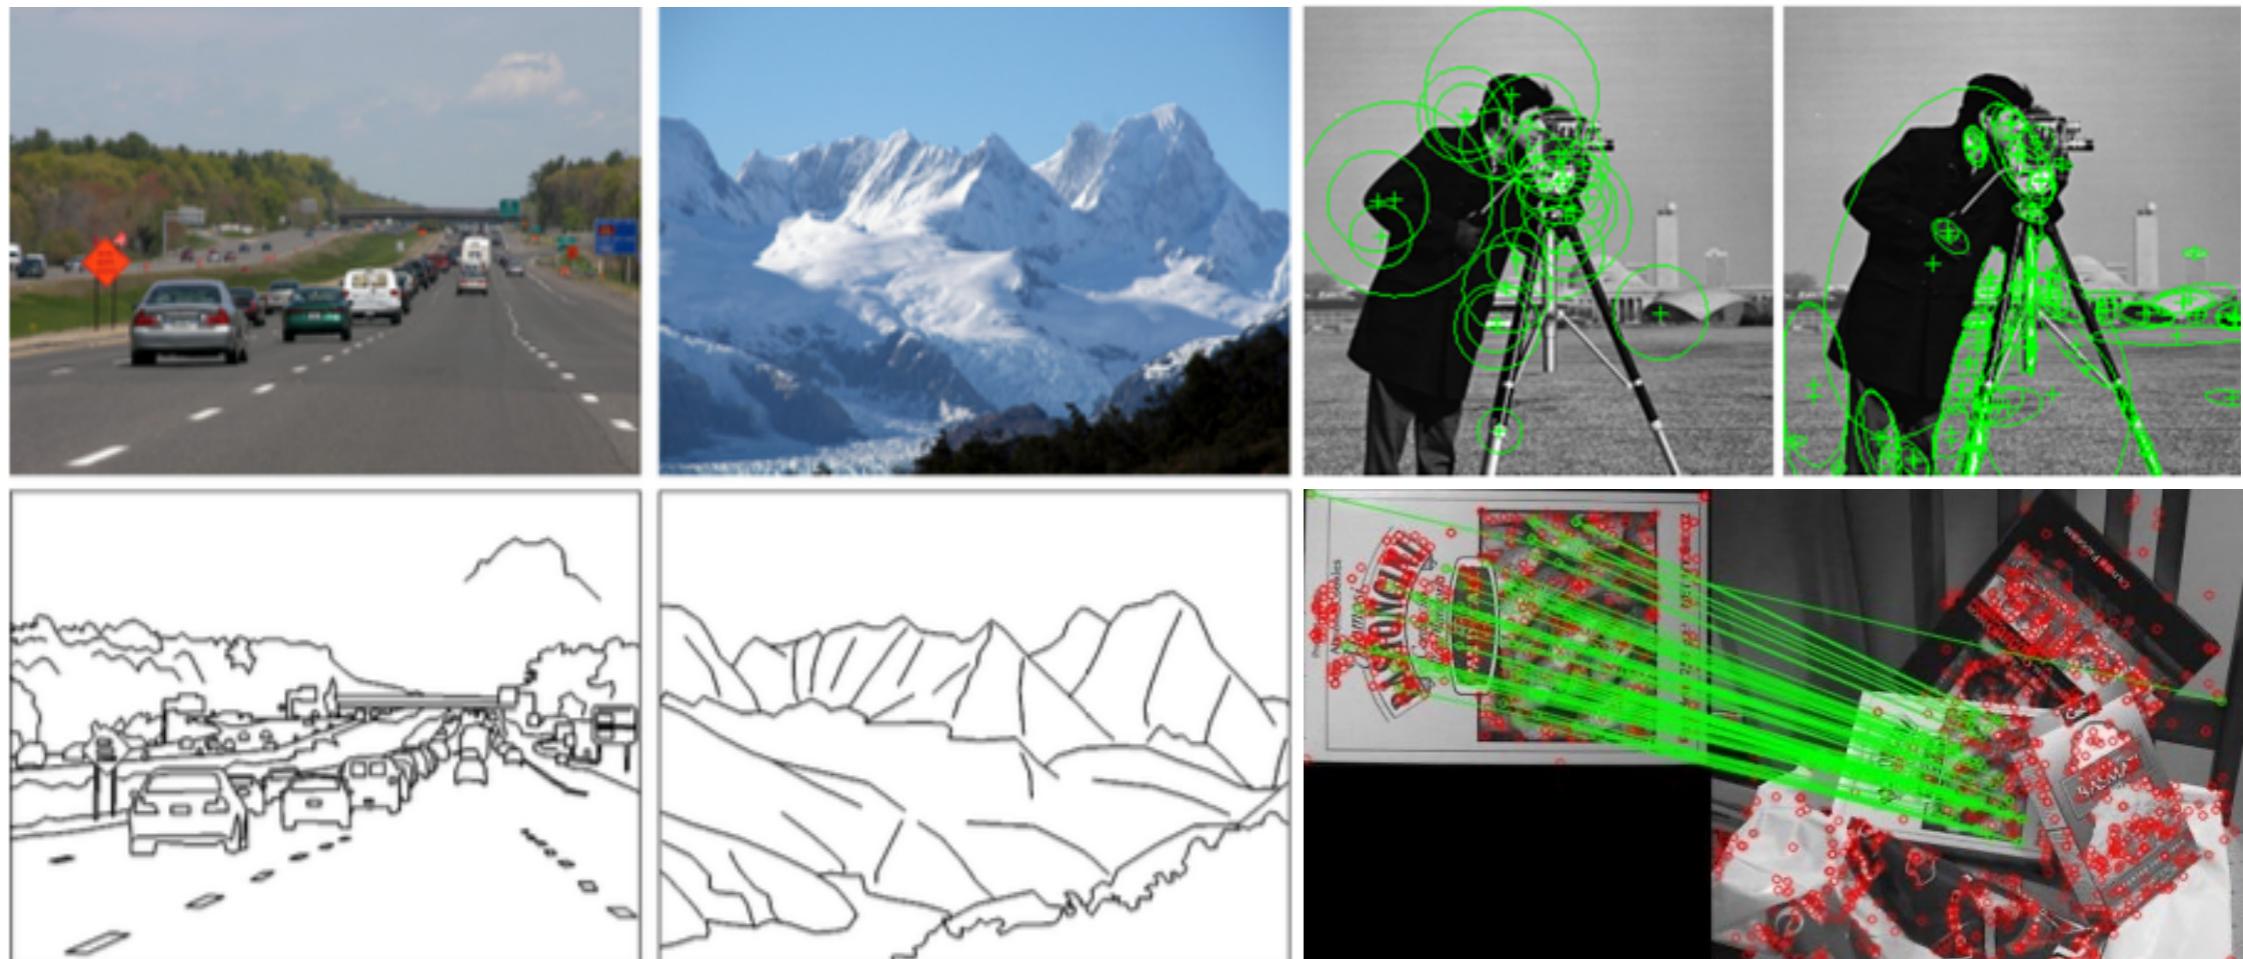
- It depends what we are looking for...



Coming up

- **Next lectures:**

- ▶ Lab on image filtering
- ▶ Image gradients, edges, feature detectors/descriptors



Contact

- **Office:** Torre Archimede, room 6CD3
- **Office hours** (ricevimento): Friday 9:00-11:00

✉ lamberto.ballan@unipd.it
🏡 <http://www.lambertoballan.net>
🏡 <http://vimp.math.unipd.it>
('@) [@ lambertoballan.bsky.social](https://lambertoballan.bsky.social)