


Computer Vision & Cognitive Systems

SCQ5109806 - LM CS,DS,CYB,PD




Intro to visual recognition, BoW and image classification

Prof. Lamberto Ballan

What we learned until now

(a brief summary)

- Edge detection
- Image gradients
- Local visual features
- DoG and SIFT descriptors



What we learned until now

(ongoing...)

- We are working on a first example of instance-based object recognition (logo recognition)
- Proposed pipeline:
 - Detect keypoints (e.g. DoG / SIFT detector)
 - Build keypoint descriptors (e.g. SIFT features)
 - Match keypoint features
 - Recognition: threshold on #matches



Recap: Visual Recognition Contest

(ongoing...)




- ▶ Deadline: November 12, 2025 - 11:00pm
- ▶ You can participate individually or as a group (2 persons)
- ▶ In your entries (max 3 submissions) as username put full name(s) and add a short description of what you have done in the note field
- ▶ Top-ranked groups shall present their approach (~2 min)




- ▶ **Query images:** 4 logos
- ▶ **Dataset** (testset): 110 images
- ▶ **Evaluation metrics:** precision, recall, mean avg precision (mAP)

What we will learn today?

- Introduction to object recognition
- Image classification
- Bag of (visual) words
- Spatial pyramids



Vision as a source of semantic information



Object recognition


- From specific object instances to classes
- **Image classification:** a core task in computer vision

Image



Assume given set of discrete labels (classes)

{cat, dog, plane, car, person, bird, sky, ...}




“bird”

Computer vision meets ML

- Data-driven approach:
 - Collect a dataset of images and labels
 - Use Machine Learning to train a classifier
 - Evaluate the classifier on new images


Example: Caltech-101 http://www.vision.caltech.edu/Image_Datasets/Caltech101/



Datasets in computer vision research




BSD (2001)



Caltech 101 (2004), Caltech 256 (2006)



PASCAL (2007-12)




LabelMe (2007)



ImageNet (2009)



SUN (2010)



Places (2014)



MS COCO (2014)



Visual Genome (2016)

2001

now

Image classification pipeline

- How to implement it?
- We need two major “ingredients”:
 - A way to *describe* images (global image representation)
 - A procedure to *compare* different images and *learn* a statistical model of a specific class

Image




img_feat: $[c_1, \dots, c_N]$



“bird”


Nearest Neighbor classifier

- Which is the simplest approach?
 - Nearest Neighbor! This is the intuition:



Nearest Neighbor classifier

- Which is the simplest approach?
 - Nearest Neighbor! This is the intuition:




Nearest Neighbor classifier

- Compute distances (Euclidean):

$$d = \text{sqrt}((x_q - x_p)^2 + (y_q - y_p)^2)$$

- ▲ Class 1
- Class 2




x_q	y_q
10	8

x_q	y_q	d	x_q	y_q	d
9	2	6.1	16	11	6.7
5	3	7.1	14	12	5.7
14	3	6.4	18	12	8.9
8	4	4.5	18	14	10
3	5	7.6	20	14	11.7
6	5	5.0	11	15	7.1
13	5	4.2	15	15	8.6
12	6	2.8	6	16	8.9
6	7	4.1	15	17	10.3
8	7	2.2	13	18	10.4
15	7	5.1	10	19	11
5	10	5.4	20	19	14.9

Nearest Neighbor classifier

- Compute distances (Euclidean):

$$d = \text{sqrt}((x_q - x_p)^2 + (y_q - y_p)^2)$$




x_q	y_q
9	11

x_q	y_q	d	x_q	y_q	d
9	2	9.0	16	11	7.0
5	3	8.9	14	12	5.1
14	3	9.4	18	12	9.1
8	4	7.1	18	14	9.5
3	5	8.5	20	14	11.4
6	5	6.7	11	15	4.5
13	5	7.2	15	15	7.2
12	6	5.8	6	16	5.8
6	7	5.0	15	17	8.5
8	7	4.1	13	18	8.1
15	7	7.2	10	19	8.1
5	10	4.1	20	19	13.6


Nearest Neighbor classifier

- Properties: NN is sensitive to outliers




k-Nearest Neighbors classifier

- A “generalization”: from NN to k-NN




k-Nearest Neighbors classifier

- Parameter k has a very strong effect
 - Small k : sensitive to outliers




$k=5$
(a.k.a. 5-NN)

majority voting:



k-Nearest Neighbors classifier

- Parameter k has a very strong effect
 - Large k : everything is classified as the most frequent class



k-Nearest Neighbors classifier

- k-NN: large data is good!




Image classification pipeline

- How to implement it?
- We need two major “ingredients”:
 - A way to *describe* images (global image representation)
 - A procedure to *compare* different images and *learn* a statistical model of a specific class

Image



img_feat: $[c_1, \dots, c_N]$



“bird”

Image classification

- What about images? How to represent them?
 - The simplest representation is provided by the raw pixels
- A simple image classification pipeline:
 - Resize images to the same scale/resolution
 - Represent images as a 2D or 1D array of raw pixels
 - Define a distance (e.g. L1 or L2) to compare images
 - Use k-NN for classification

Image classification

- Example dataset: CIFAR10

10 classes

50,000 training images

10,000 testing images


airplane



automobile




bird



cat




deer




dog



frog




horse



ship



truck




test images and nearest neighbors



Image classification

- Distance metric to compare images:

- L1 distance (*Manhattan*): $d_{L1}(I_1, I_2) = \sum_p |I_1^p - I_2^p|$



- Example:

test image

56	32	10	18
90	23	128	133
24	26	178	200
2	0	255	220

-

training image

10	20	24	17
8	10	89	100
12	16	178	170
4	32	233	112

=

pixel-wise absolute value differences

46	12	14	1
82	13	39	33
12	10	0	30
2	32	22	108

add
→ 456

Image classification

- Example (Python implementation):

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

“Memorize” training data

For each test image:

- Find closest train image
- Predict label of nearest train image

Image classification

- Example (Python implementation):

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Q: With N examples,
how fast are NN
training and prediction?

A: Train $O(1)$, Test $O(N)$

Q: Is this good or bad?

*A: This is bad! We want
classifiers that are fast at
prediction (slow for
training is ~ok)*

Image classification

- What can we do better?

- ▶ We can design/use a better image representation
- ▶ We can rely on better functions to compare images
- ▶ We can design/use better classifiers

Bag of (visual) words model

- Intuition: build an intermediate representation based on the number of occurrences of “*visual primitives*”

Object




Bag of “words”



Bag of words model

- Origin: bag of words model for text categorization
 - The task is to assign a textual document to one or more categories based on its content



- Approach: a document is represented as an *unordered* collection of words disregarding grammar & word order

Bag of (visual) words / Bag of features

- The very same idea can be applied for image classification (object recognition)



- Works pretty well for image/video classification and also for recognizing object instances

Sivic & Zisserman (2003, 2005), Csurka et al. (2004), Grauman & Darrell (2005), Niebles et al. (2006, 2008), Ballan et al. (2009, 2012), ...

Bag of features: pipeline

- Three main steps:
 - Feature detection (sampling) and description
 - *Take a bunch of images and extract local visual features*
 - Codebook formation and image representation
 - *Build up a dictionary, i.e. our intermediate representation*
 - *For each feature find the closest visual word in the visual dictionary and build a frequency histogram*
 - Learning and recognition

Tutorials / short courses:

“Recognizing and Learning Object Categories” - Fei-Fei, Fergus, Torralba [[web](#)]

“Hands on Advanced BoW Models for Visual Recognition” - Ballan, Seidenari [[web](#)]

Bag of features: pipeline


- Feature extraction:



- Codebook formation:



- Quantize features using visual vocabulary
- Represent images by frequencies of “visual words”:



Feature extraction (sampling)

- Keypoint detectors (e.g. DoG, Harris, MSER, ...)




Feature extraction (sampling)

- Dense sampling using a regular grid
(or other simple methods, e.g. random sampling)




Codebook formation

- Learning the visual vocabulary:




Codebook formation

- Learning the visual vocabulary:



Codebook formation

- Learning the visual vocabulary:




Clustering recap: k-means

- **Goal:** minimize sum of squared Euclidean distances between points x_i and their nearest cluster centers m_k

$$D(X, M) = \sum_k \sum_{i \in k} (x_i - m_k)^2$$

- (Lloyd's) Algorithm:
 - Randomly initialize K cluster centers
 - Iterate until convergence:
 - Assign each data point to the nearest center
 - Recompute each cluster center as the mean of all points assigned to it




From clustering to vector quantization

- Clustering is a common method for learning a visual vocabulary or codebook
 - This is an unsupervised learning process
 - Each cluster center becomes a *codeword* (or *codevector*)
 - Codebook can be learned on separate training set
 - Provided the training set is sufficiently representative, the codebook will be “universal”
- The codebook is used for quantising visual features
 - A *vector quantizer* takes a feature vector and maps it to the index of the nearest codevector in a codebook
 - codebook = visual vocabulary; codevector = visual word


Codebook formation: visual words

- Visual word example 1: what's inside a cluster




Codebook formation: visual words

- Visual word example 2: what's inside a cluster



Codebook formation: visual words

- Visual word example 3: what's inside a cluster



Visual vocabularies: issues

- How to choose vocabulary size?
 - ▶ Too small: visual words not representative of all patches
 - ▶ Too large: quantization artifacts, overfitting
- Computational efficiency
 - ▶ Vocabulary trees

Nister and Stewenius, "Scalable Recognition with a Vocabulary Tree", CVPR 2006




Image credit: D. Nister

Features quantization

- Assign each feature to the most representative visual word





Image representation

- Compute histograms of visual word frequencies



Uses of BoW representation

- Image classification: treat as feature vector for standard classifiers
 - e.g. k-Nearest Neighbors, Support Vector Machines, ...
- Large-scale image search / retrieval
 - BoW vectors have been useful in matching an image to a large database of object instances
- Cluster BoW vectors over image collection
 - i.e. discover visual themes

Weighting the (visual) words

- Just as with text, some visual words are more discriminative than others
 - The bigger fraction of the documents a word appears in, the less useful it is for matching

the, and, or


vs

bank, brain, perception

- TF-IDF: instead of computing a regular histogram distance, we weight each word by its Inverse Document Frequency
 - IDF of word j : $\log \frac{\text{number of documents}}{\text{number of documents in which } j \text{ appears}}$


What about spatial information?

- Spatial information between objects/parts is lost



Spatial pyramids

- Intuition: locally orderless representation at several levels of spatial resolution




Spatial pyramids

- Very useful and effective for representing images:
 - Pyramid is built by using multiple copies of image
 - Each level is $1/4$ of the size of previous level
 - The lowest level is of the highest resolution
 - The highest level is of the lowest resolution
- You can apply the same idea using different image partitions

S.Lazebnik, C.Schmid, J.Ponce, “Beyond bags of features: Spatial pyramid matching for recognizing natural scene categories”, CVPR 2006

Spatial pyramids

- Example: experimental results on Caltech-101



Multi-class classification results (30 training images per class)

	Weak features (16)		Strong features (200)	
Level	Single-level	Pyramid	Single-level	Pyramid
0	15.5 $\pm$ 0.9		41.2 $\pm$ 1.2	
1	31.4 $\pm$ 1.2	32.8 $\pm$ 1.3	55.9 $\pm$ 0.9	57.0 $\pm$ 0.8
2	47.2 $\pm$ 1.1	49.3 $\pm$ 1.4	63.6 $\pm$ 0.9	64.6 $\pm$ 0.8
3	52.2 $\pm$ 0.8	54.0 $\pm$ 1.1	60.3 $\pm$ 0.9	64.6 $\pm$ 0.7

Contact

- **Office:** Torre Archimede, room 6CD3
- **Office hours** (ricevimento): Friday 9:00-11:00

✉ lamberto.ballan@unipd.it

🏠 <http://www.lambertoballan.net>

🏠 <http://vimp.math.unipd.it>

@ lambertoballan.bsky.social