



Vision and Cognitive Systems

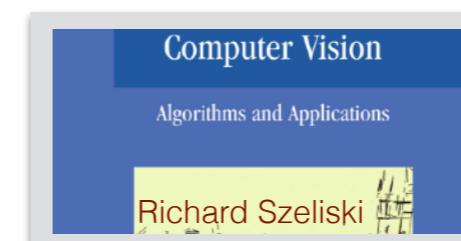
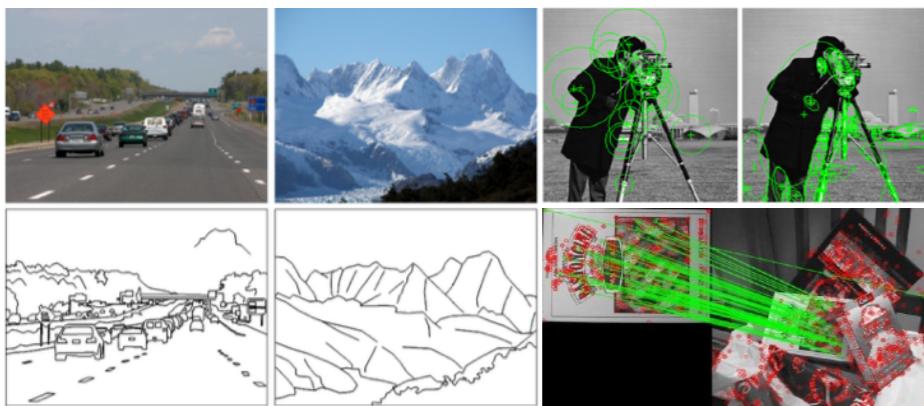
SCQ1097939 - LM CS,DS,CYB,PD

Foundations: edges, keypoints and local features

Prof. Lamberto Ballan

What we will learn this week?

- Edge detection
- Image gradients
- Local invariant features
- DoG and SIFT descriptors



Chapter 4

Image filtering recap

- Compute a function of the local neighborhood at each pixel in the image
 - ▶ The function is specified by a “filter” saying how to combine values from neighbors
- Applications of image filtering:
 - ▶ extract information (edges, corners, blobs, ...)
 - ▶ detect patterns (template matching)
 - ▶ de-noising, super-resolution, in-painting

Recap: derivate filter (Prewitt)

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

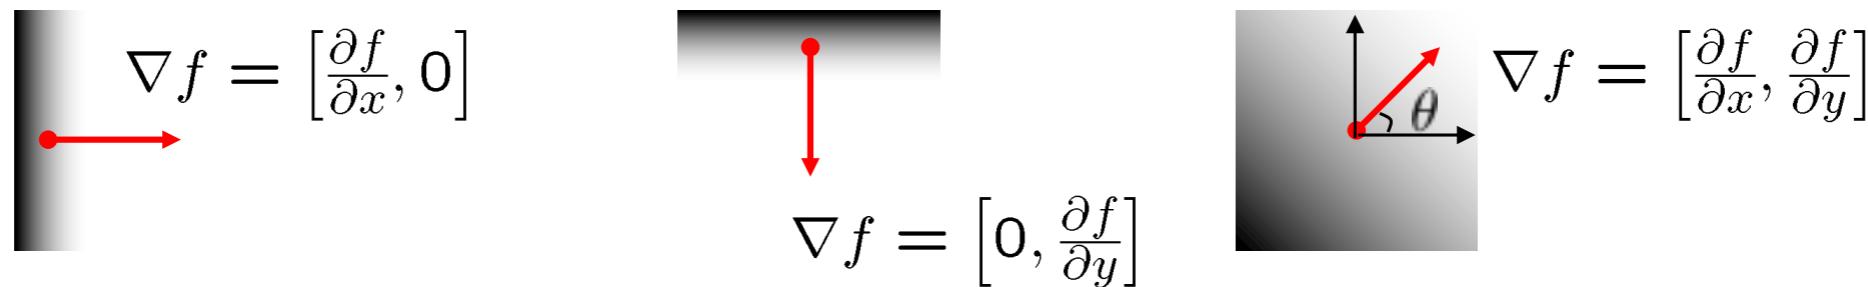
Vertical Derivatives

$$* \frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Horizontal Derivatives

Recap: image gradient

- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

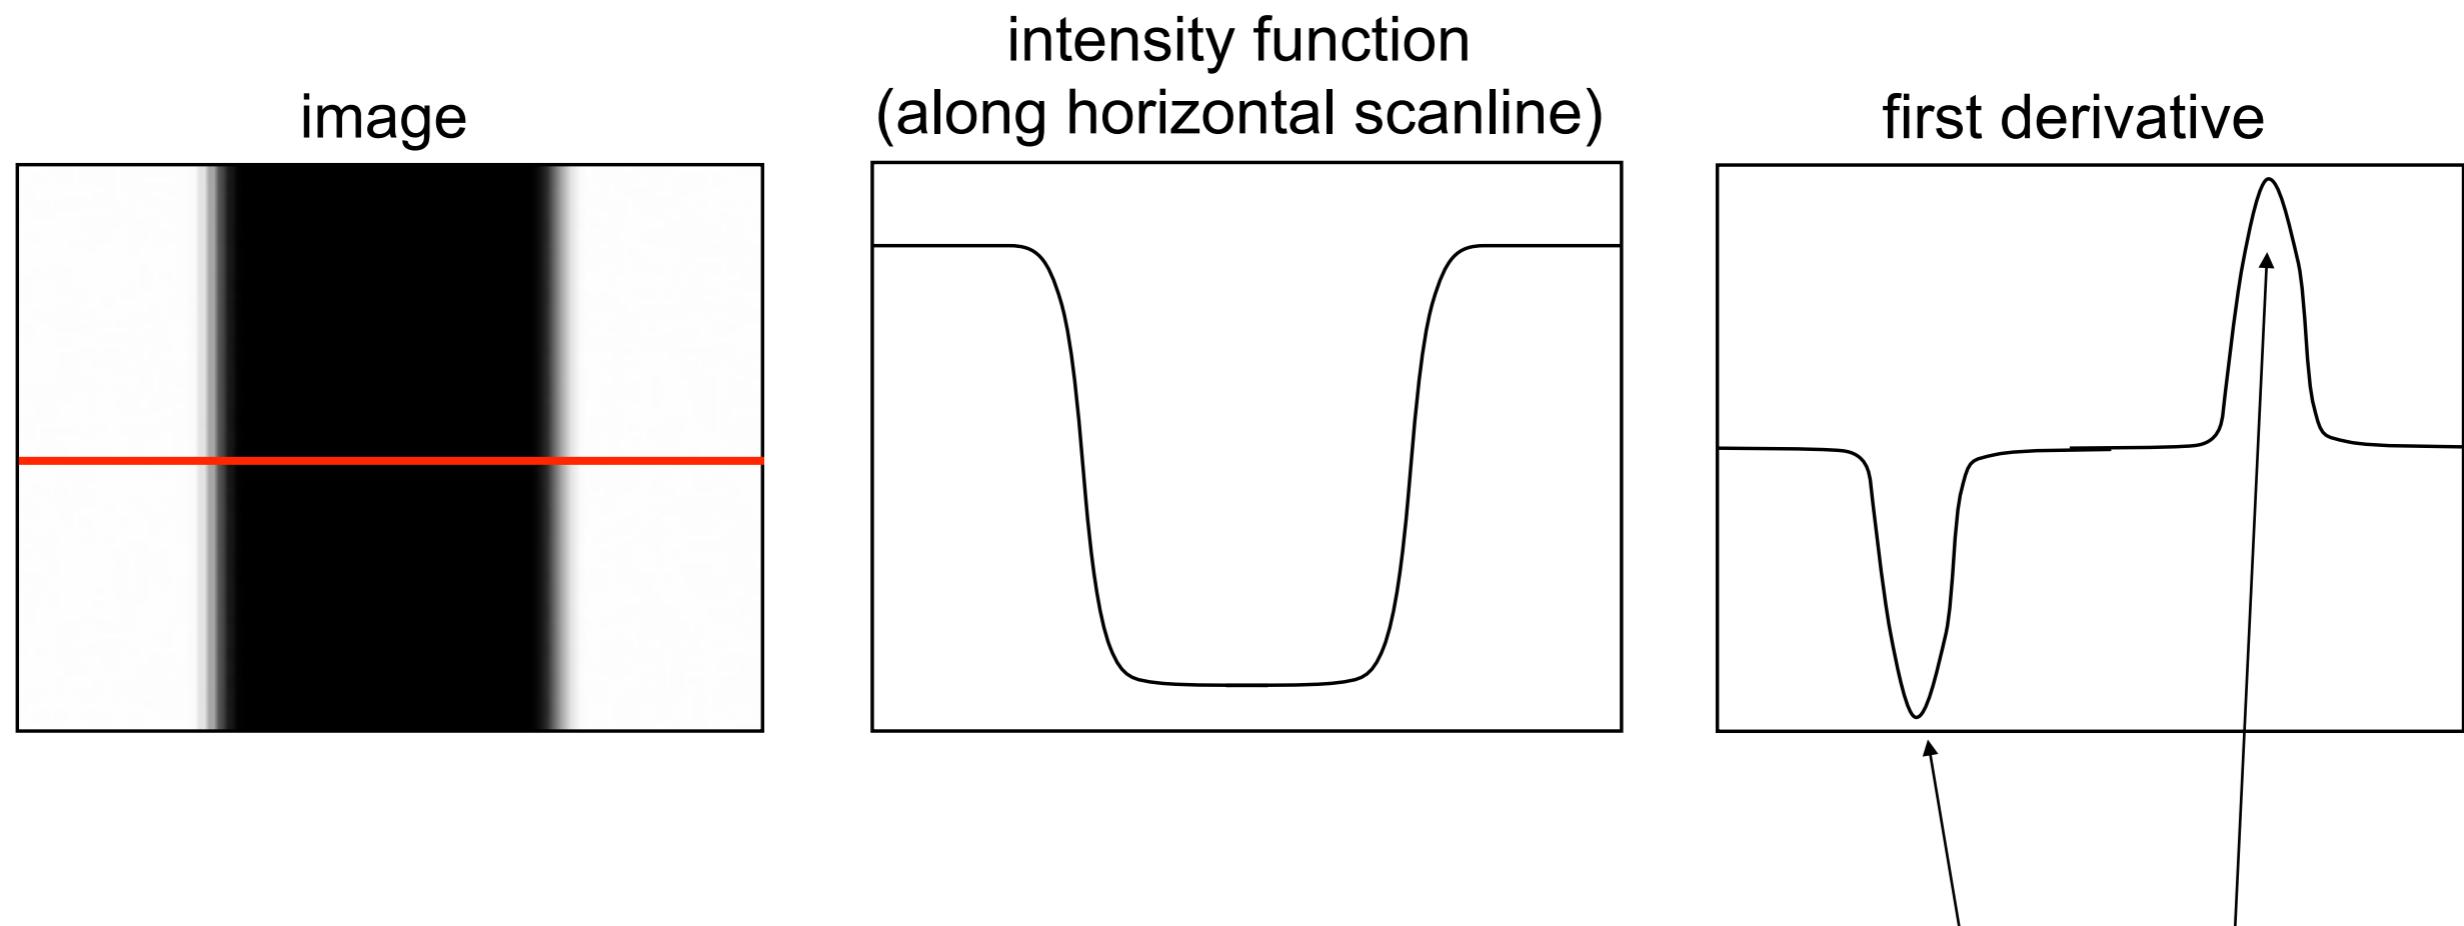


- The gradient vector points in the direction of most rapid increase in intensity
- The gradient direction is given by: $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$
- The *edge strength* is given by the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Recap: edges and image gradients

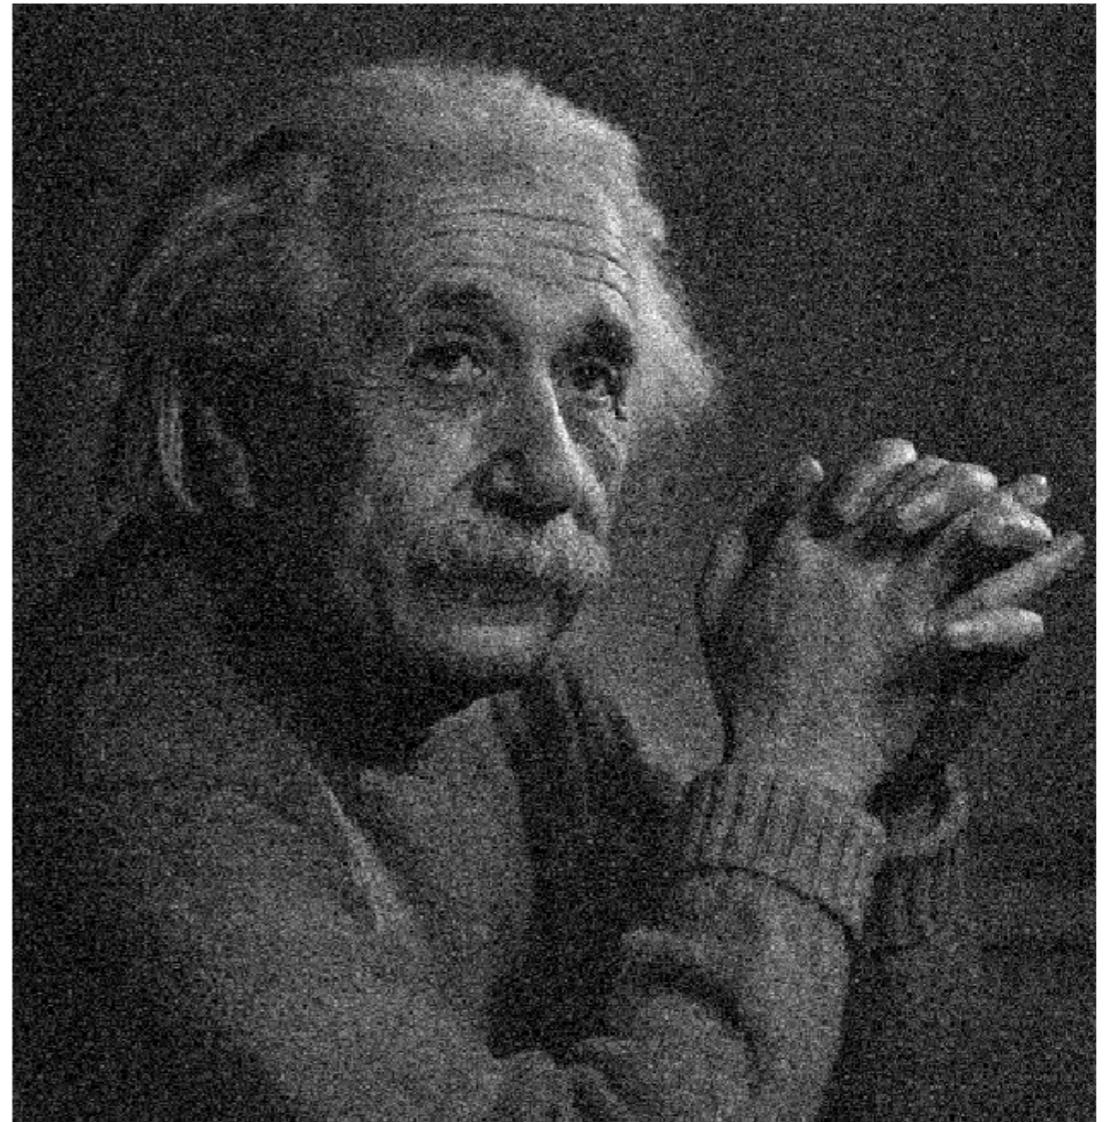
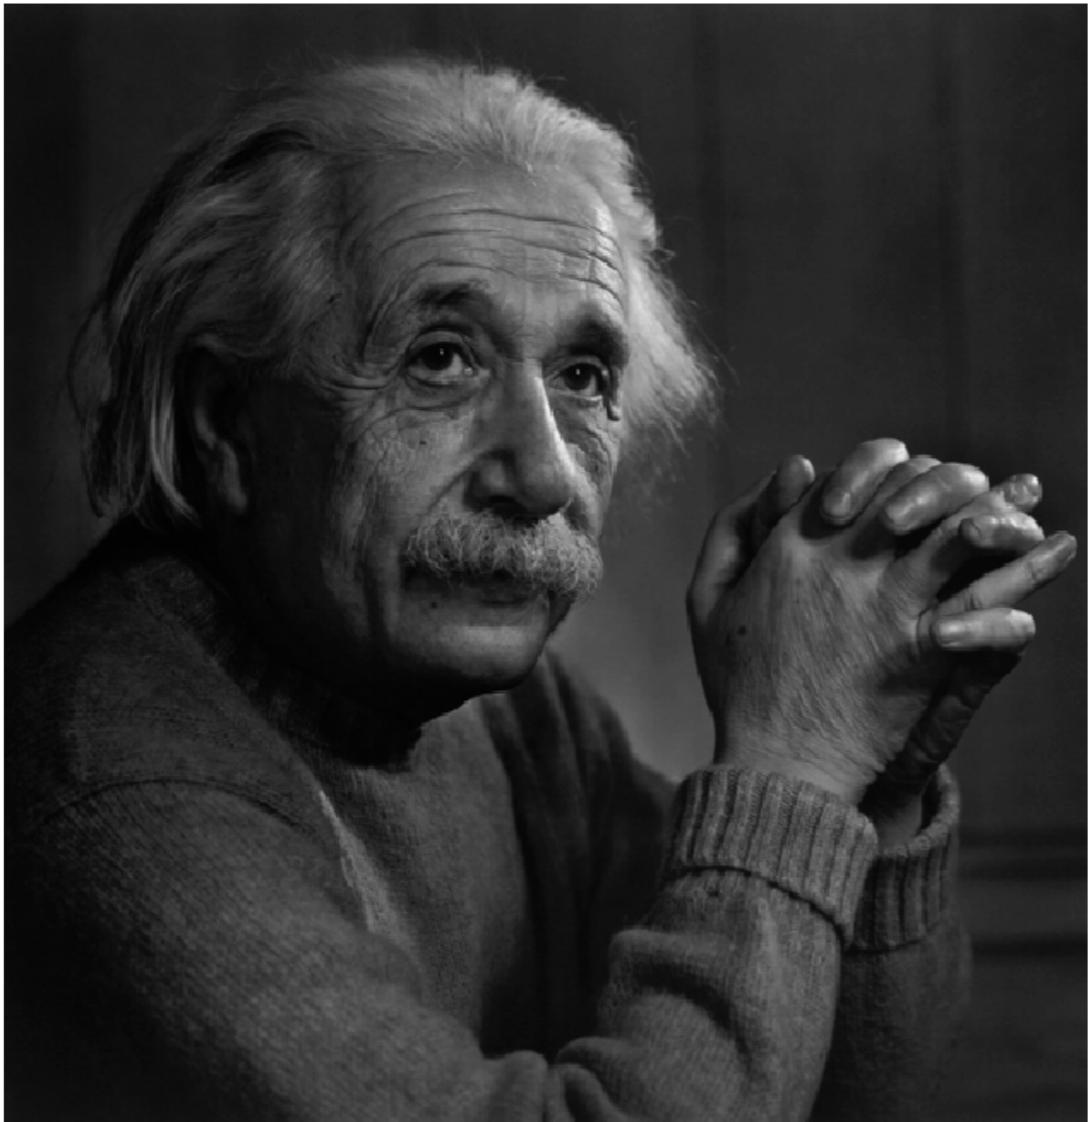
- An edge is a place of rapid change in the image intensity function



This could be the sketch of an edge detection algorithm

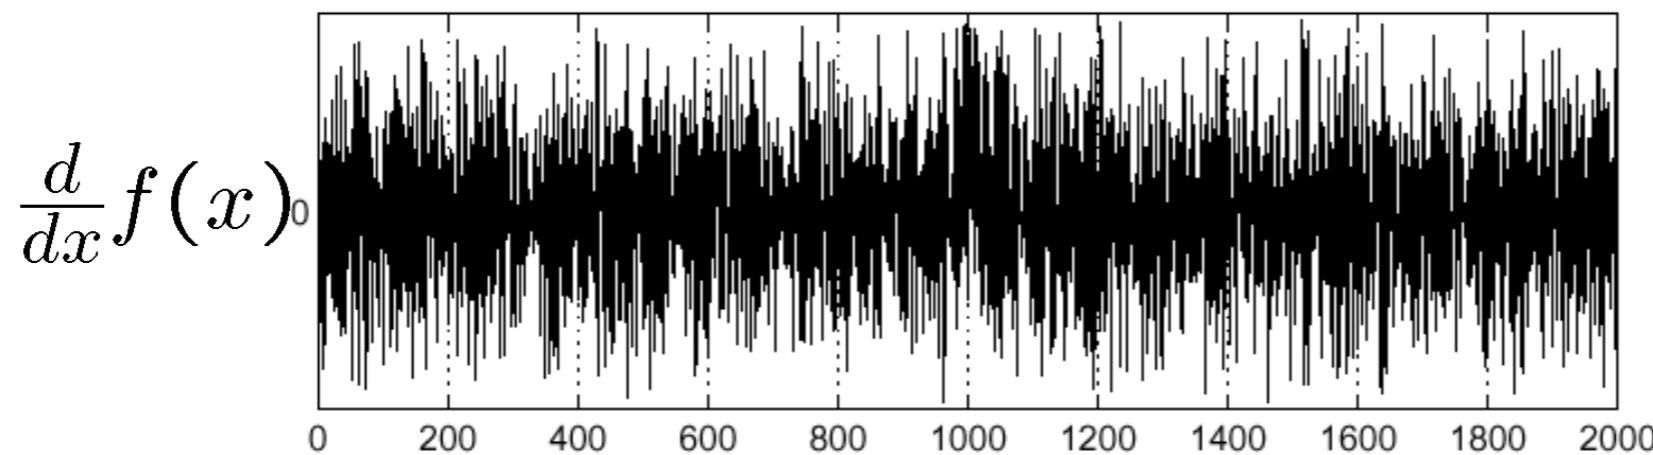
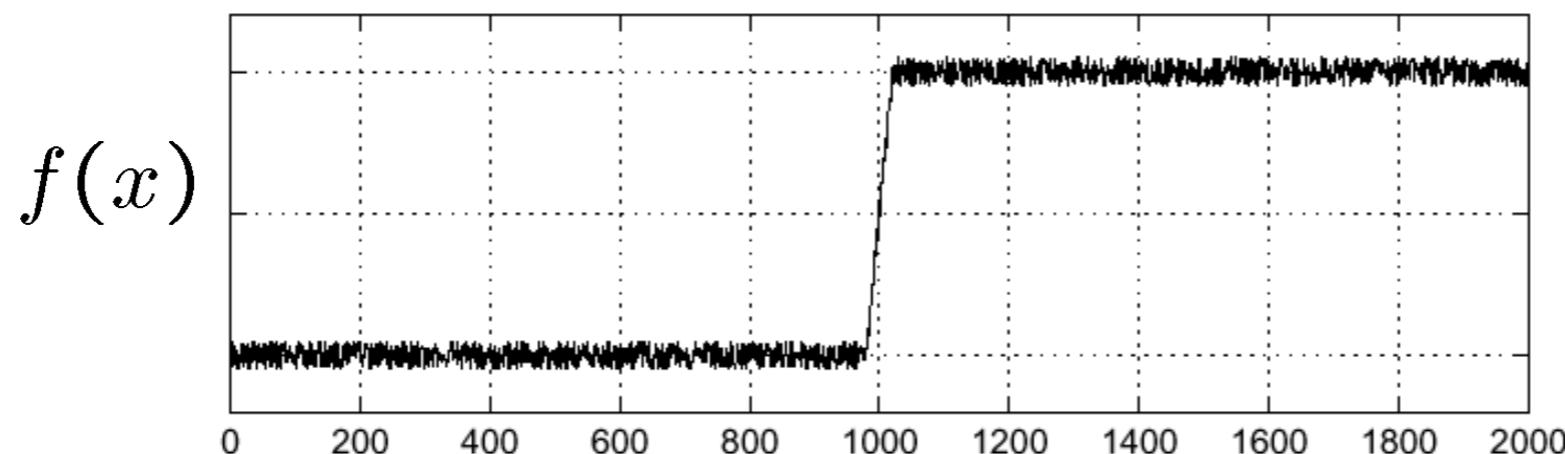
edges correspond to extrema of derivative

Effects of noise



Effects of noise

- Consider a single row or column of the image
 - ▶ Plotting intensity as a function of position gives a signal

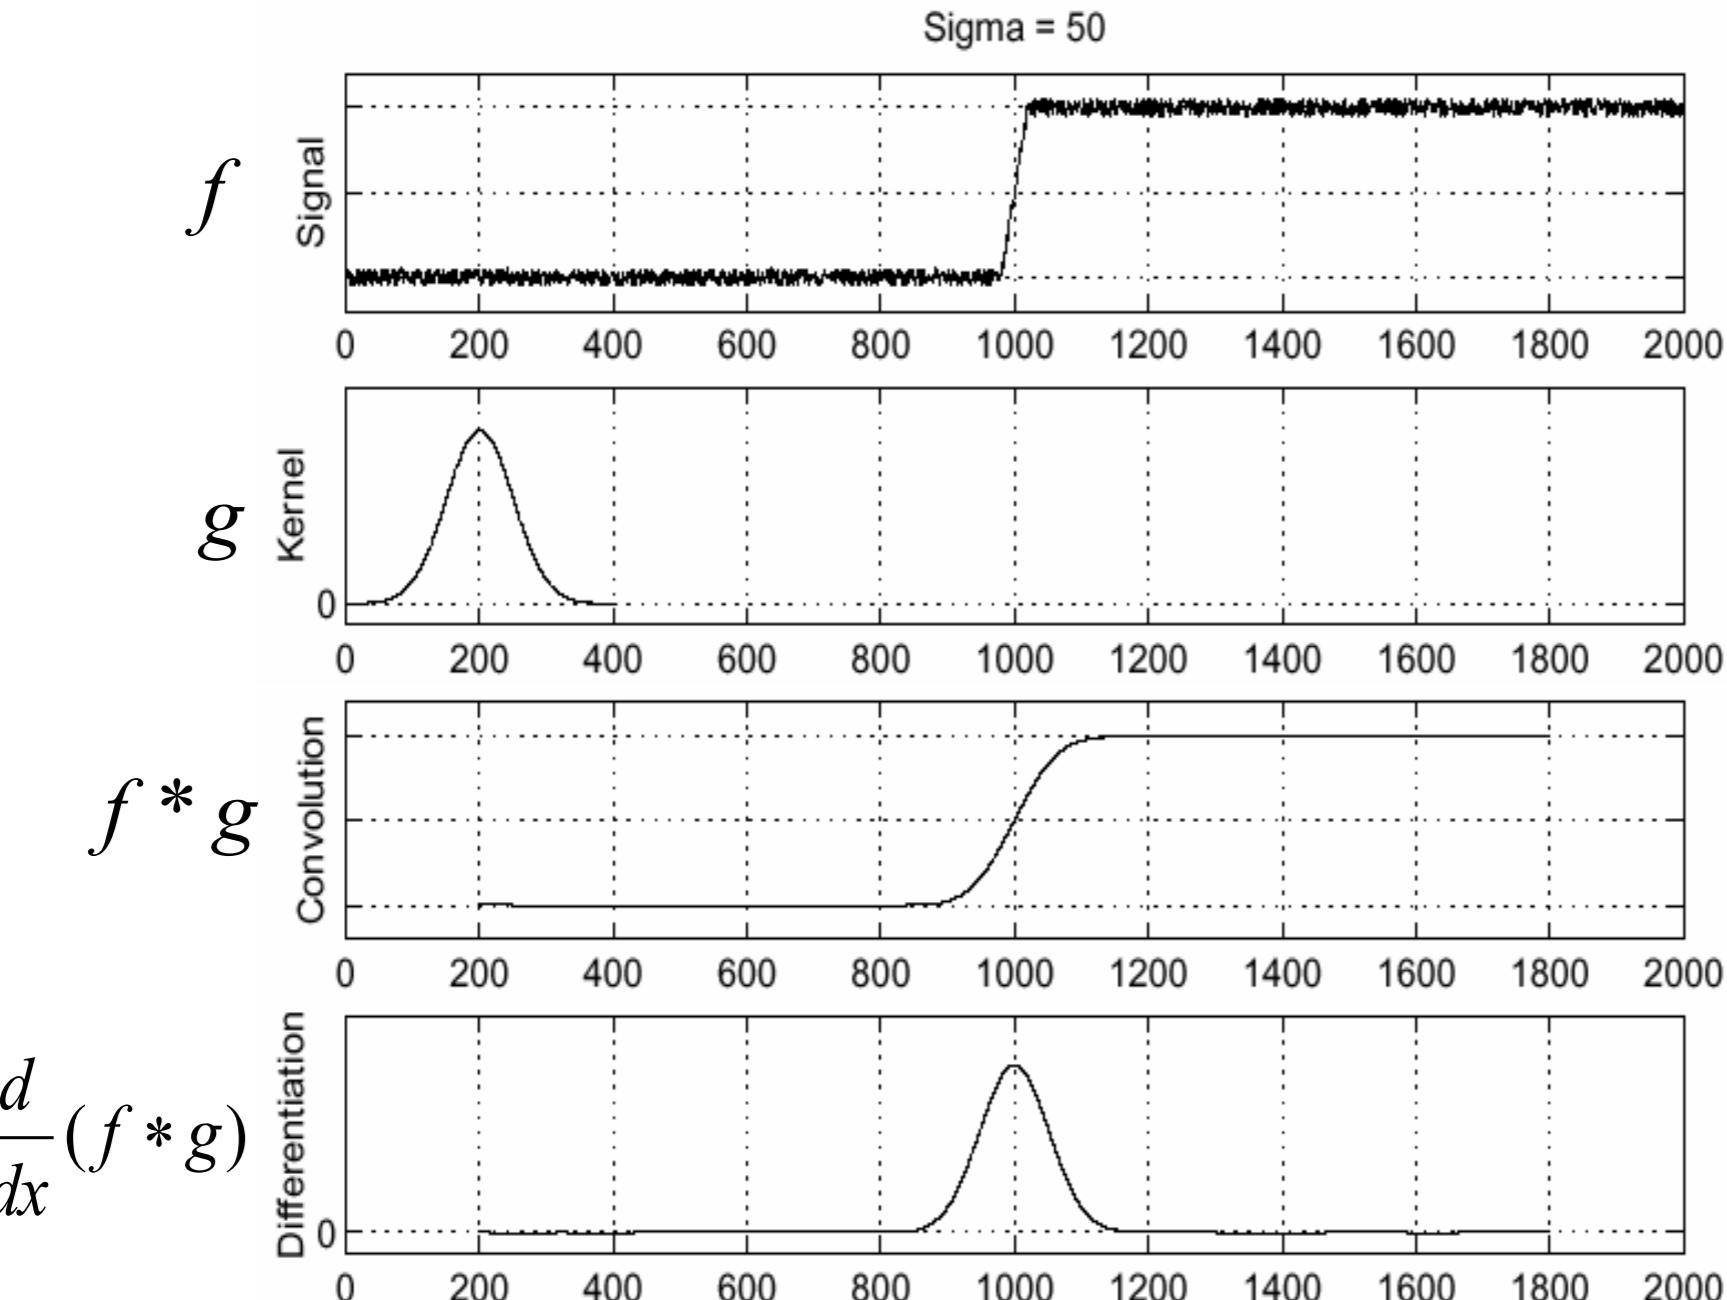


*Where is the
edge?*

Effects of noise

- Finite difference filters respond strongly to noise
 - ▶ Image noise results in pixels that look very different from their neighbors
 - ▶ The larger the noise, the stronger the response
- What is to be done?
 - ▶ Smoothing the image should help, by forcing pixels different to their neighbors (=noise pixels?) to look more like neighbors

Solution: smooth first



Where is the
edge?

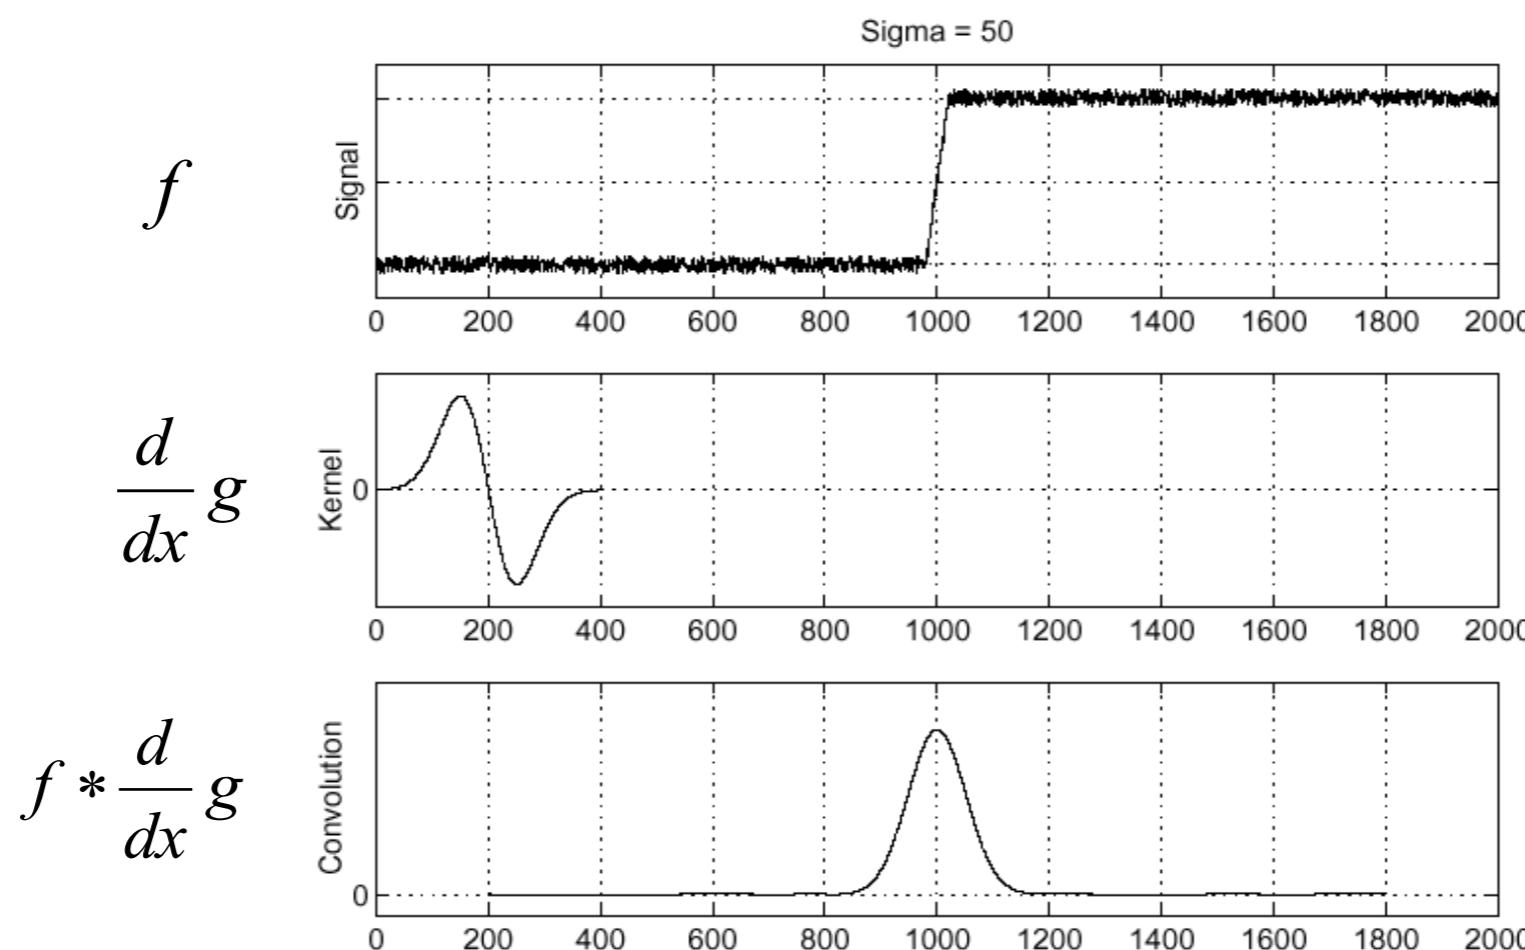
Look for peaks in $\frac{d}{dx}(f * g)$

Derivative theorem of convolution

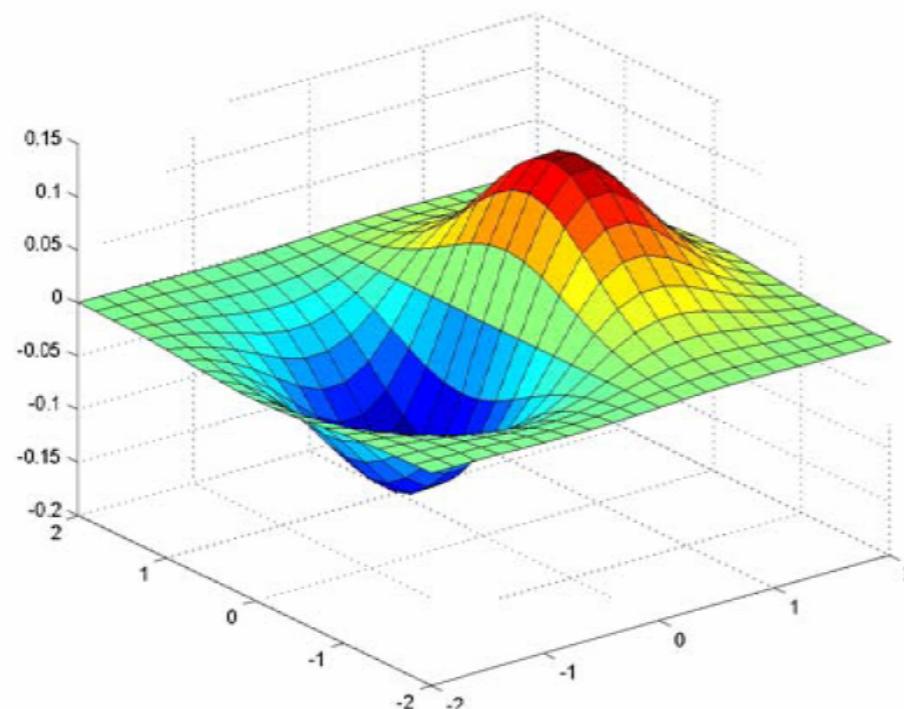
- Differential property of convolution:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

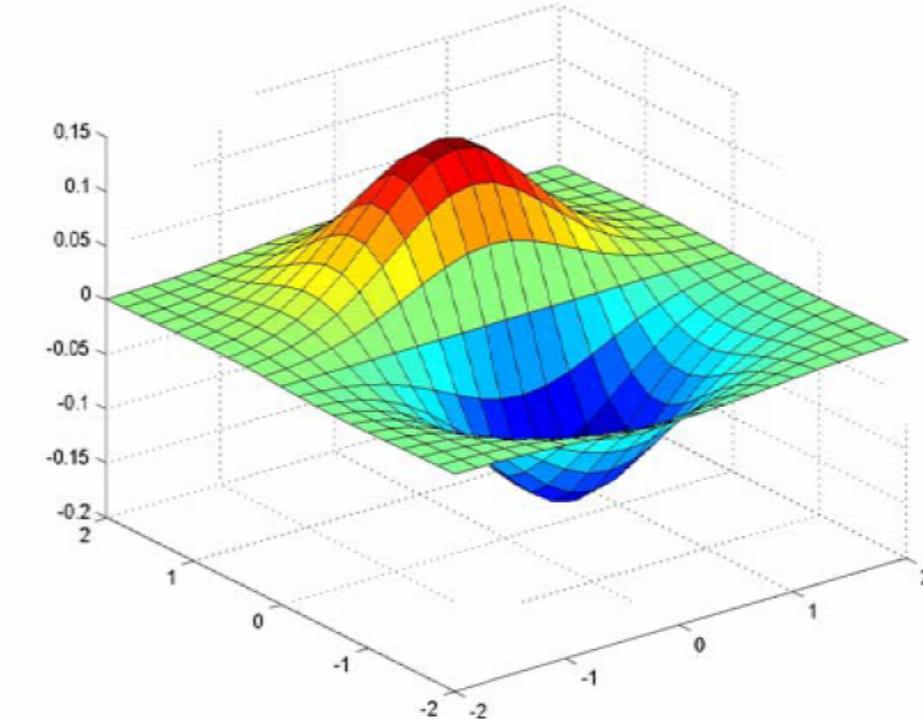
- ▶ This is a useful property and save us one operation:



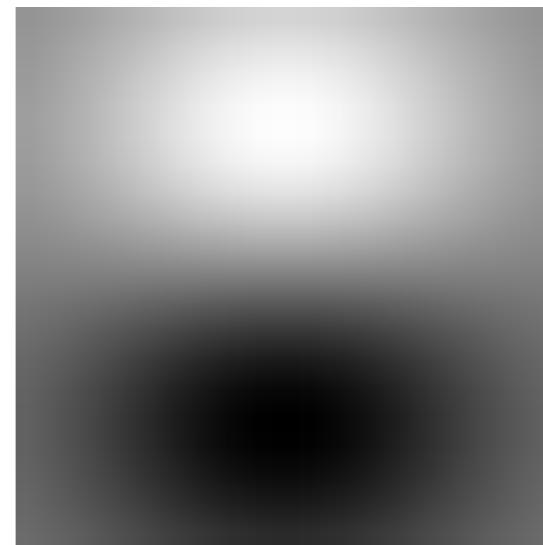
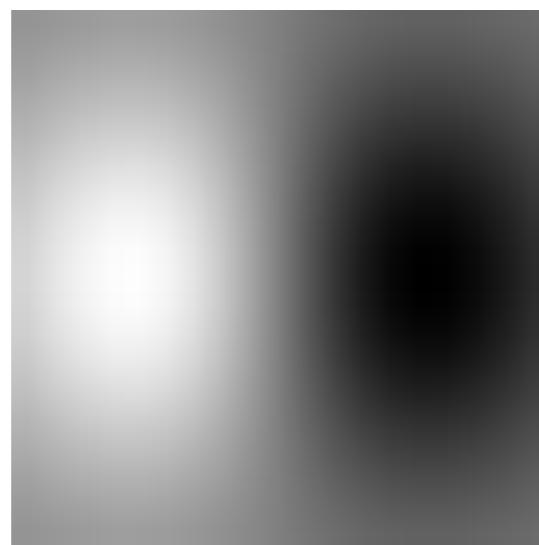
Derivative of Gaussian filter



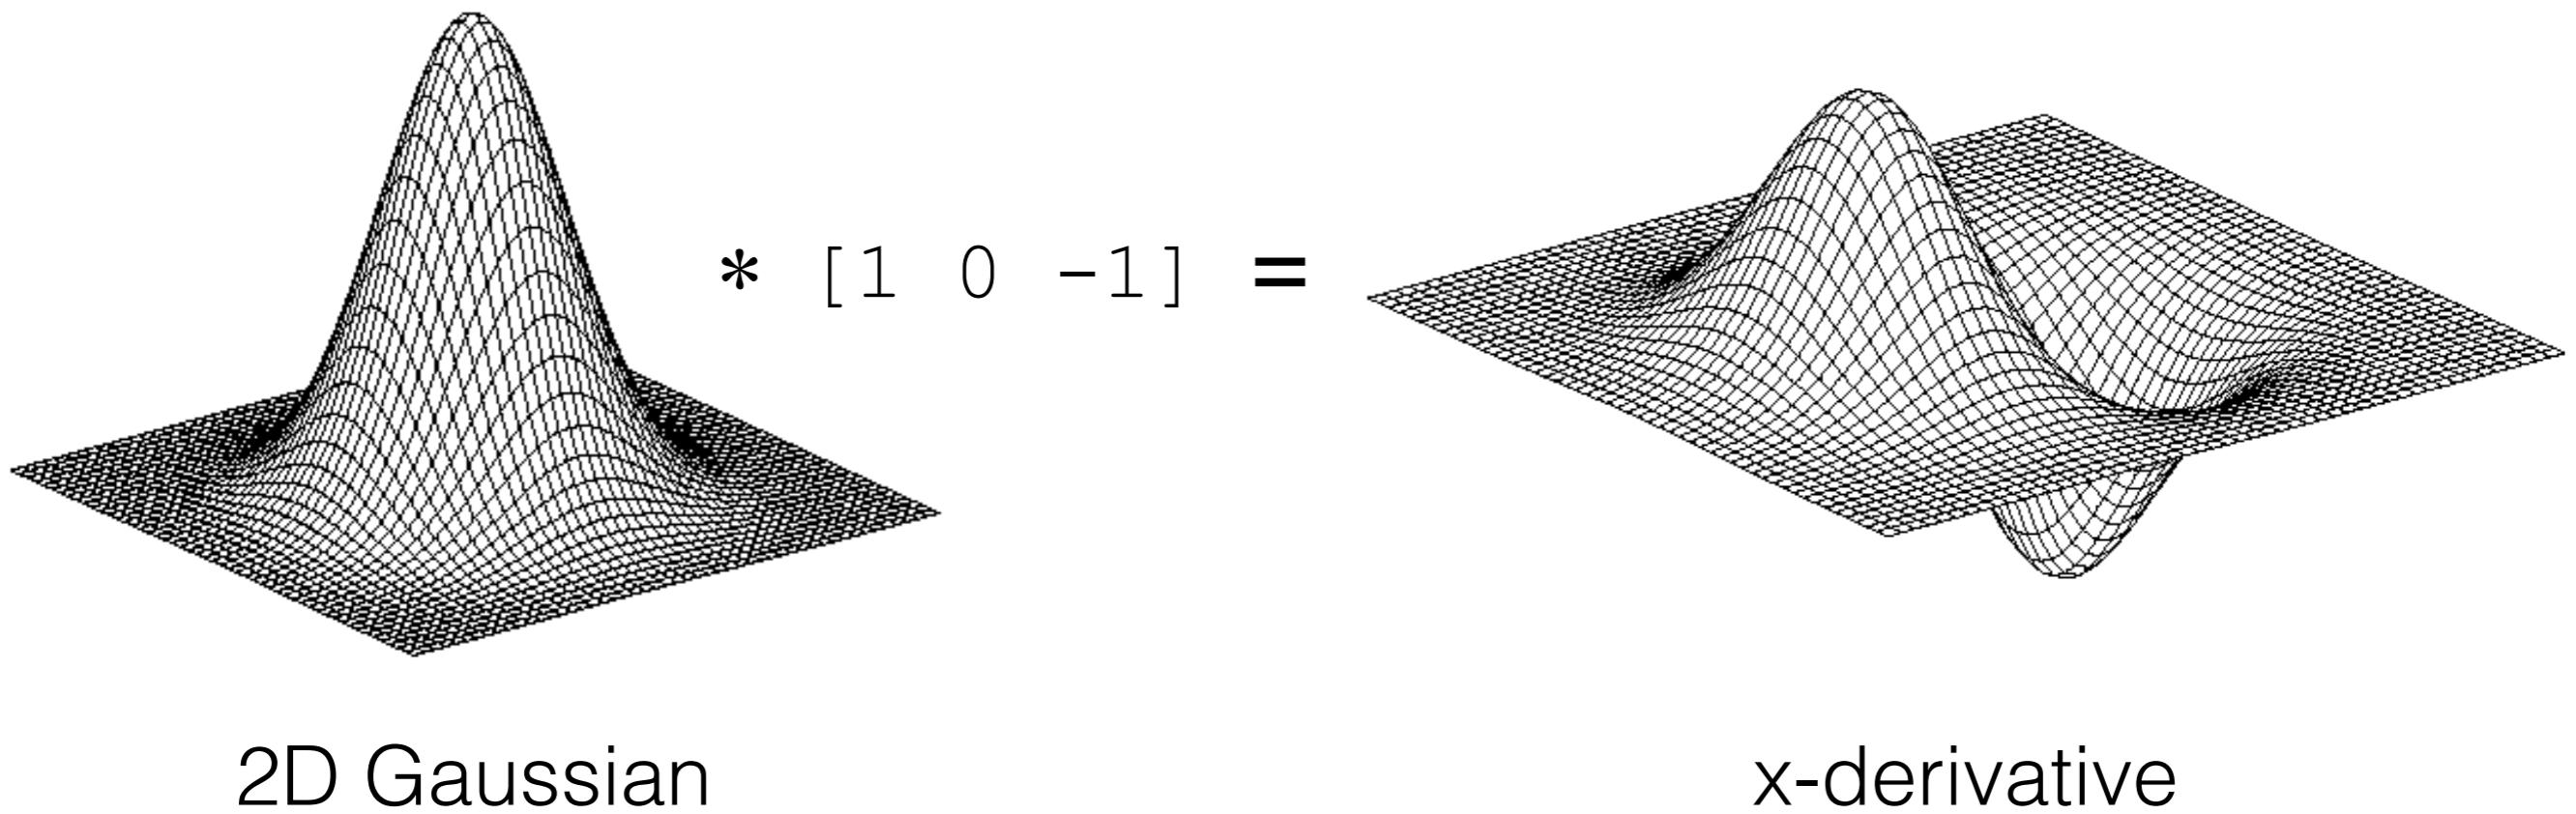
x-direction



y-direction



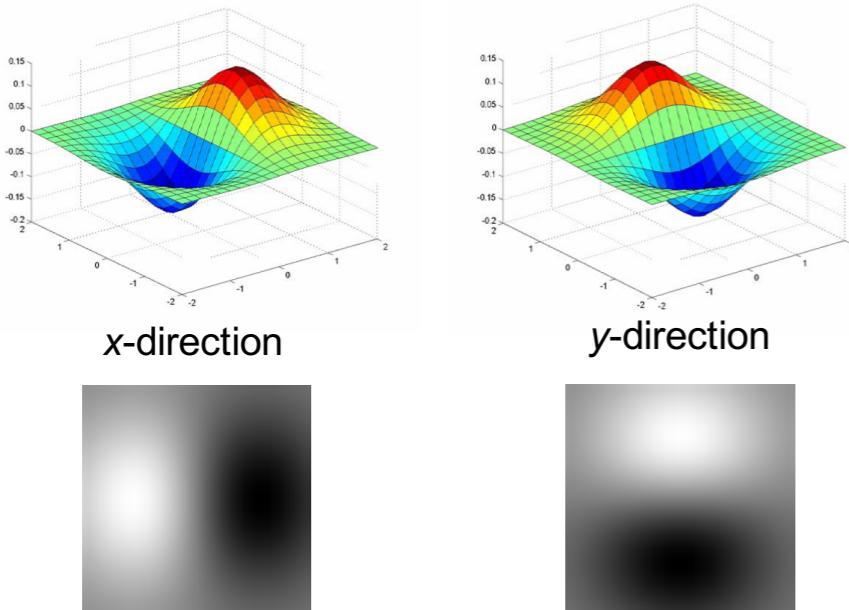
Derivative of Gaussian filter



2D Gaussian

x-derivative

Derivative of Gaussian filter

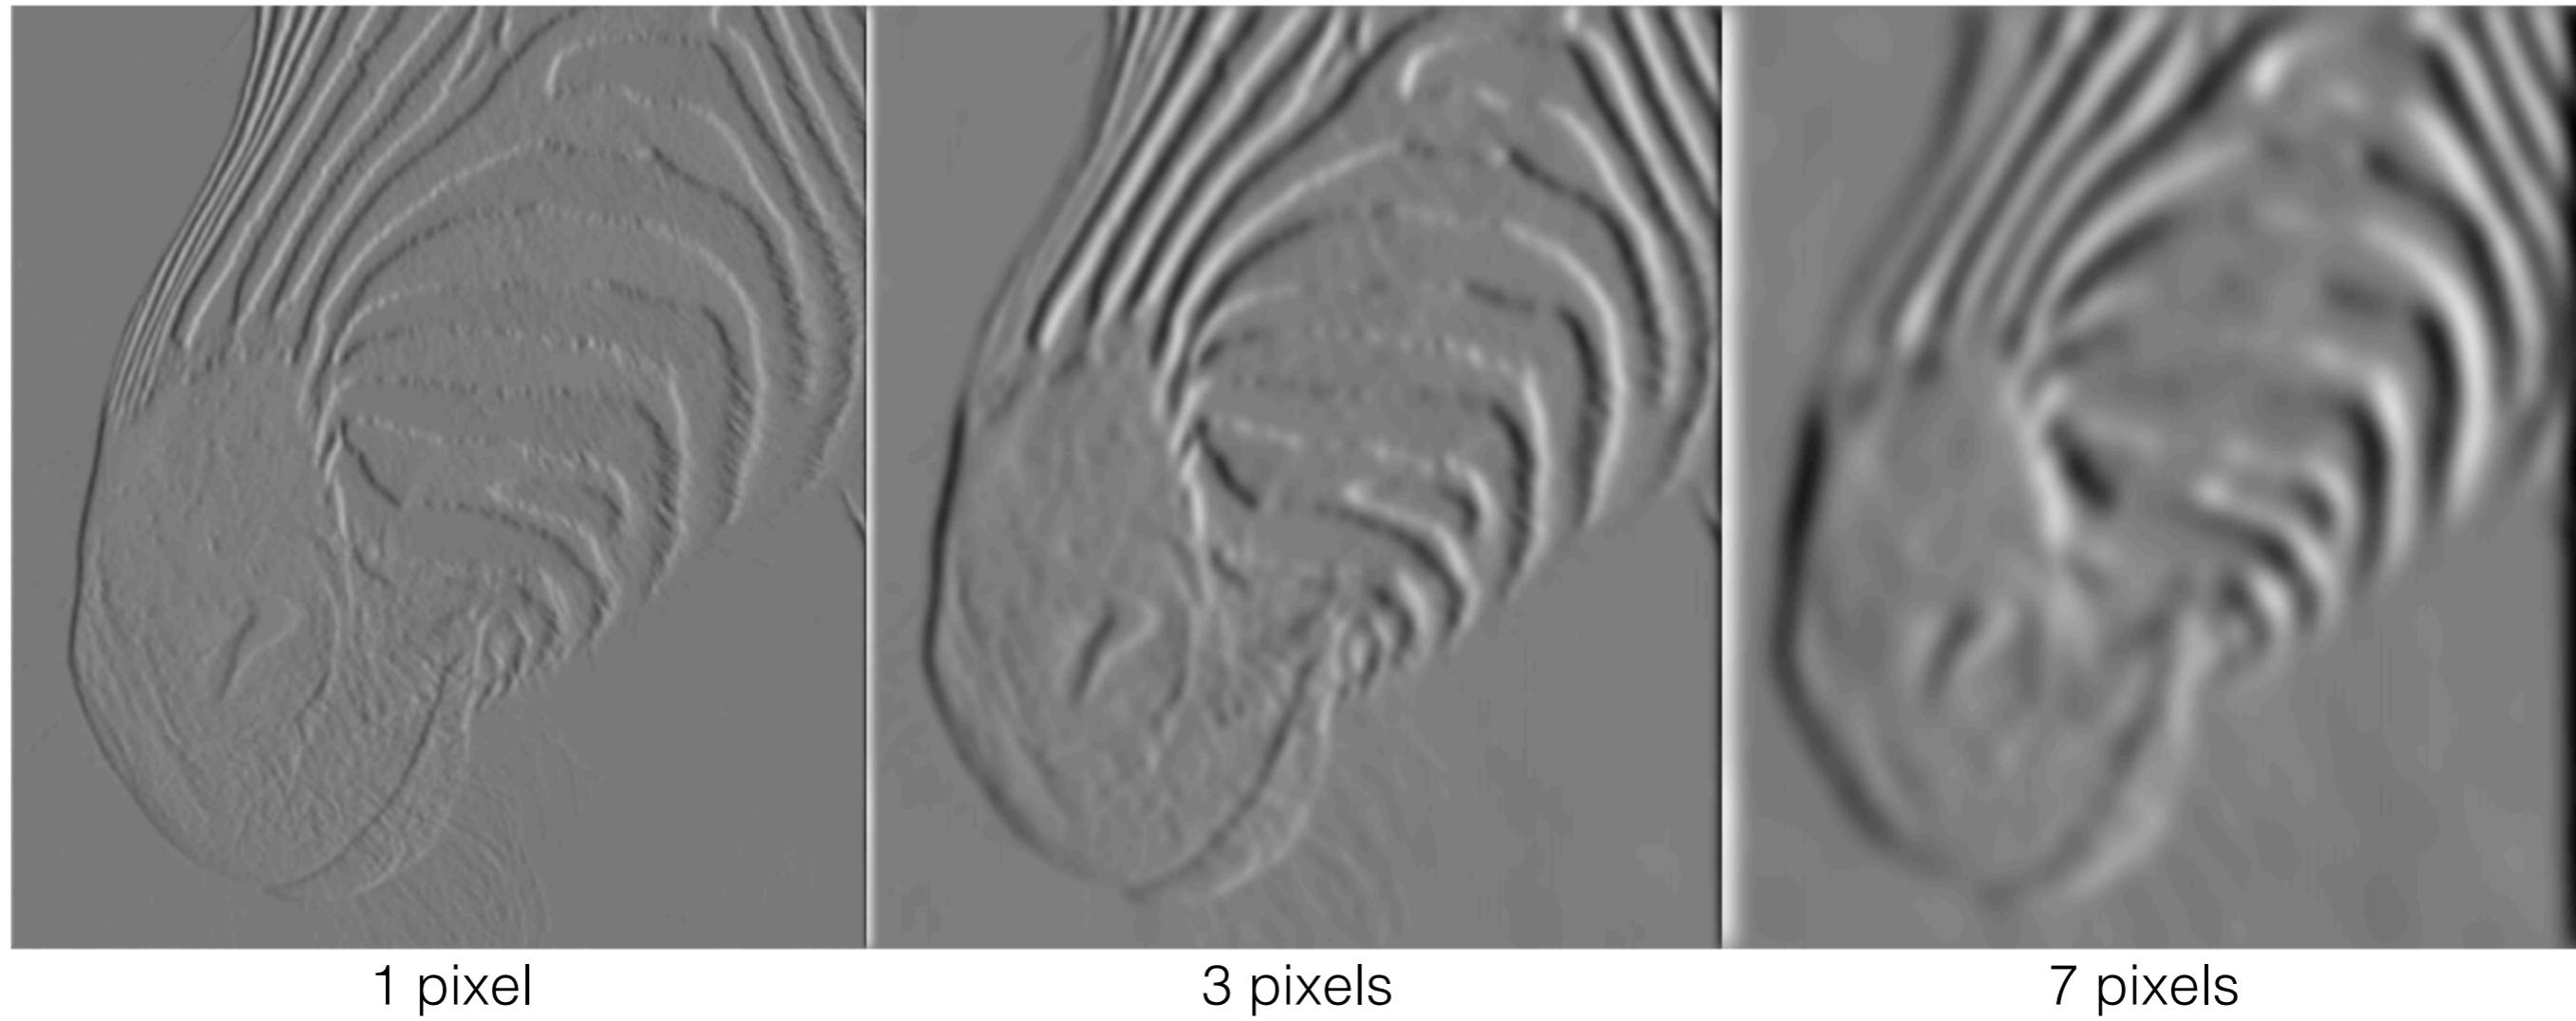


Note: separability of the Gaussian Filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

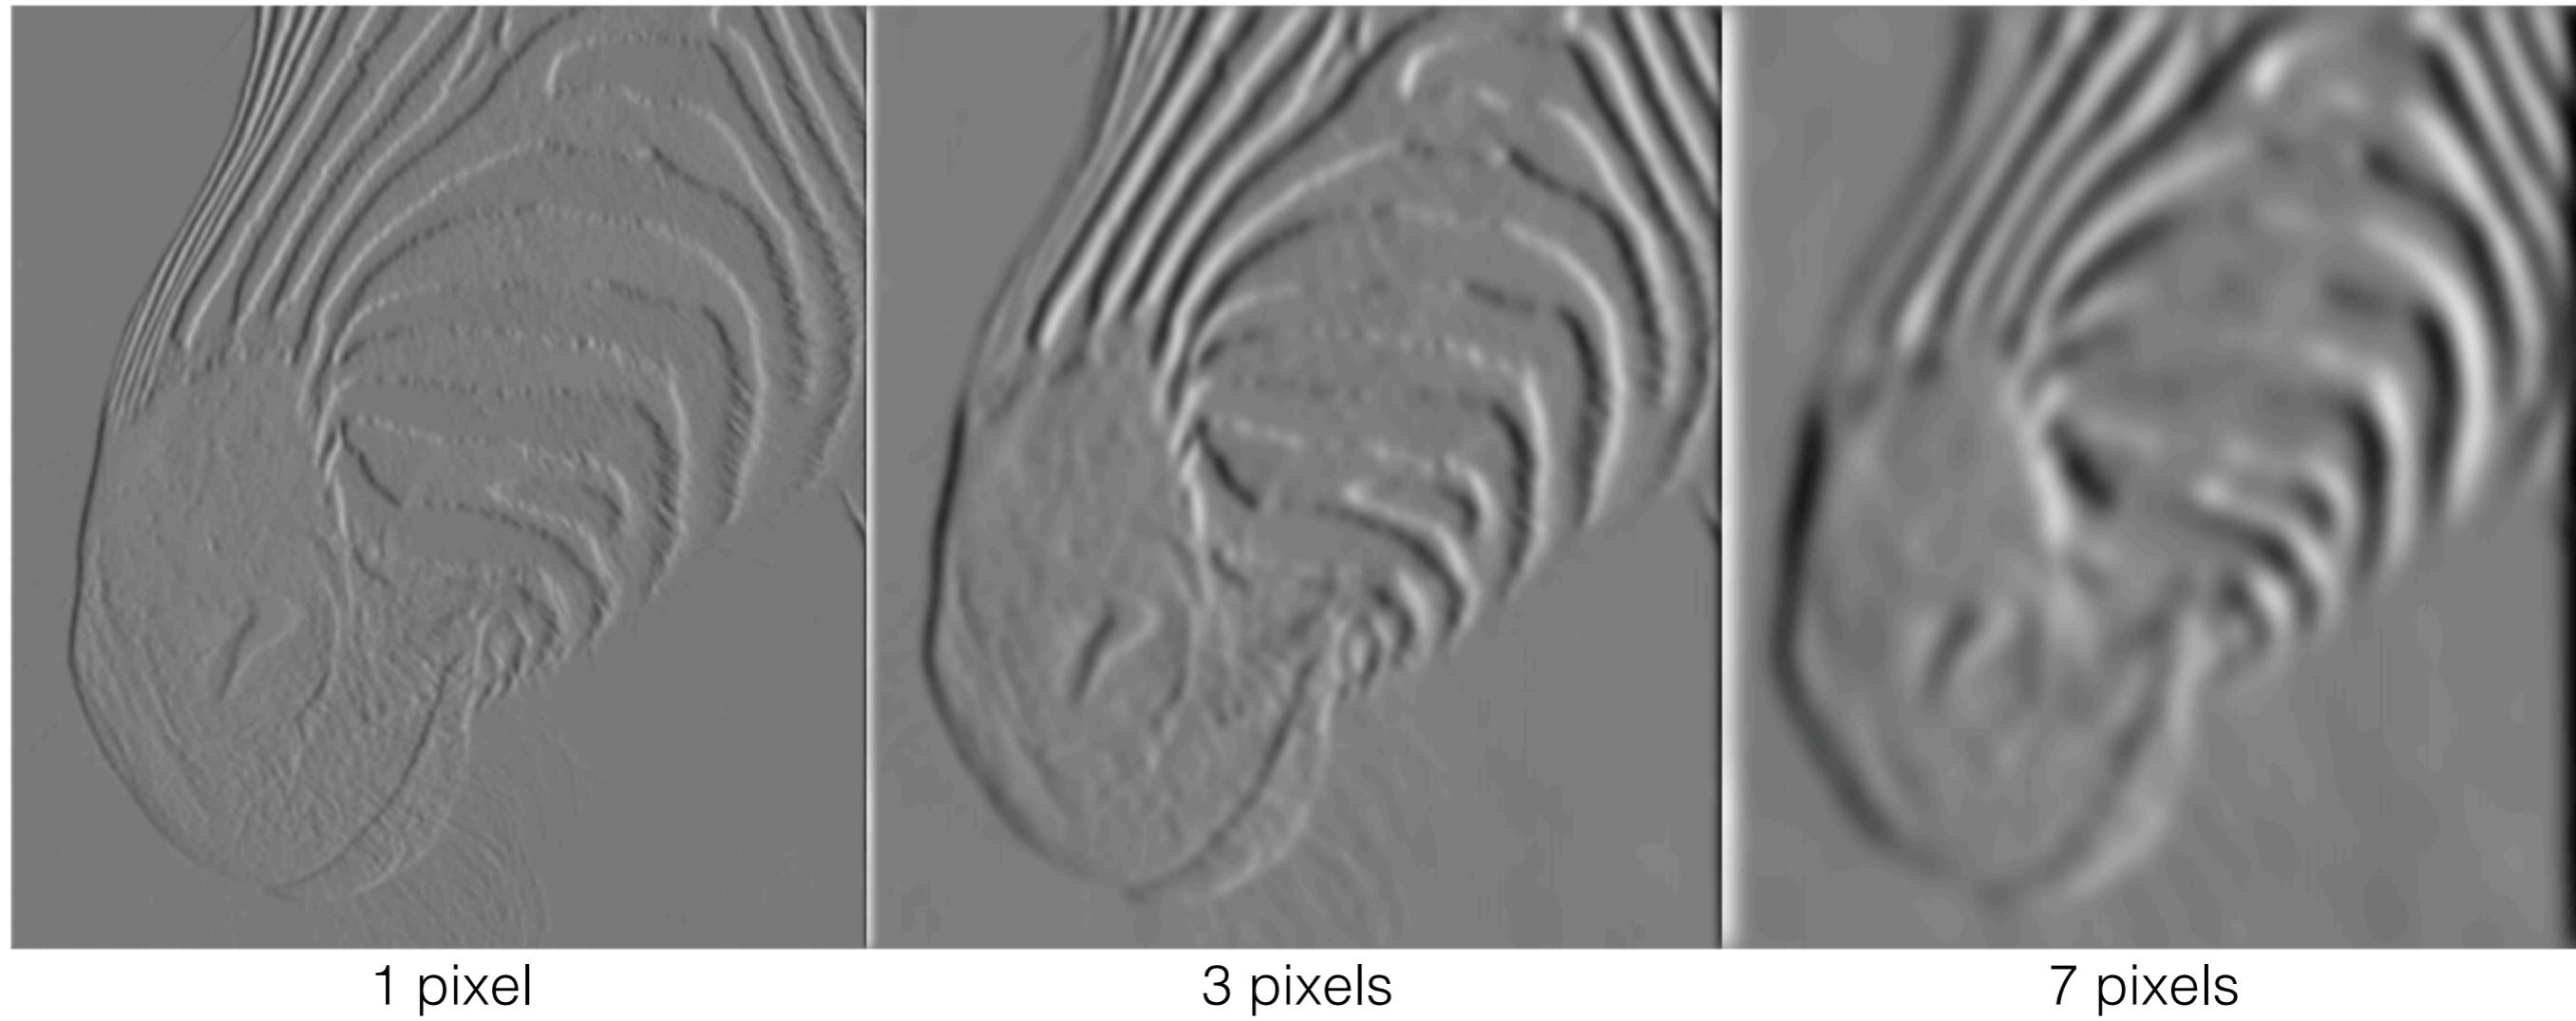
- ▶ The 2D Gaussian can be expressed as the product of 2 functions (one a function of x and the other a function of y)
- ▶ In this case the two functions are the (identical) 1D Gaussian

Scale of Gaussian derivative filter



- ▶ Note: σ is the *scale (width/spread)* of the Gaussian kernel
- ▶ Smoothed derivative remove noise, but blurs edges

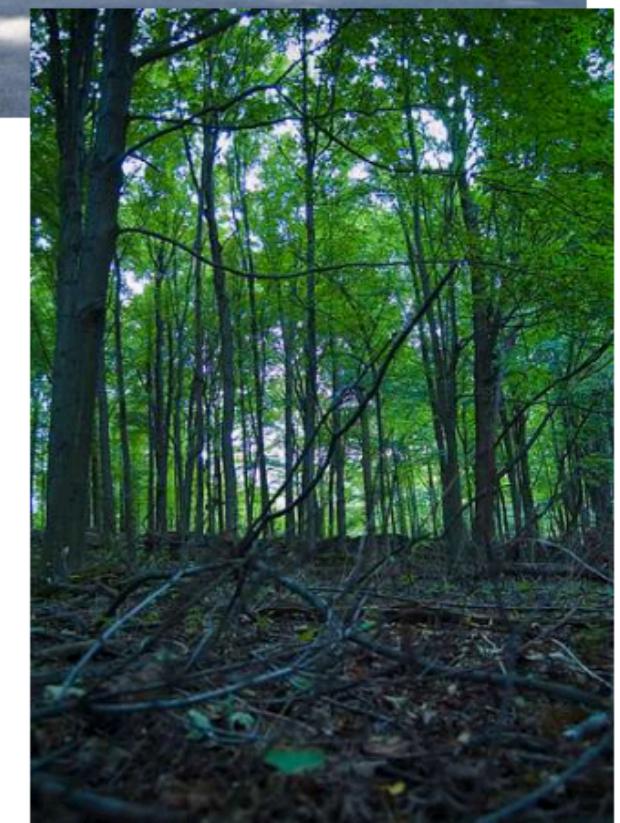
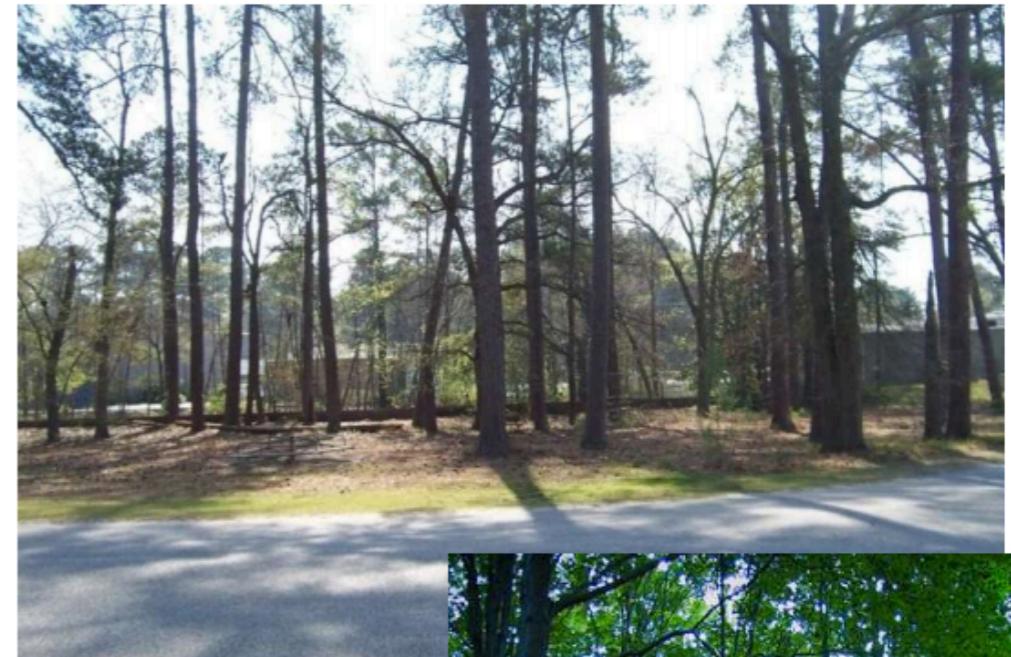
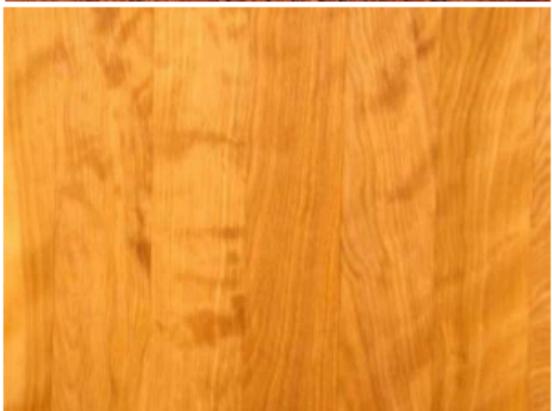
Scale of Gaussian derivative filter



- The apparent structures differ depending on scale σ
 - Larger values: larger scale edges detected
 - Smaller values: finer features detected

So, what scale to choose?

- It depends what we are looking for...



Edge detection: Sobel filter

- A derivative filter + some smoothing



Input image



*Left Sobel
(Horizontal Derivatives)*

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

Filter returns large response on vertical or horizontal lines?

A: Responds to vertical lines

Is the output always positive?

A: Output can be positive or negative

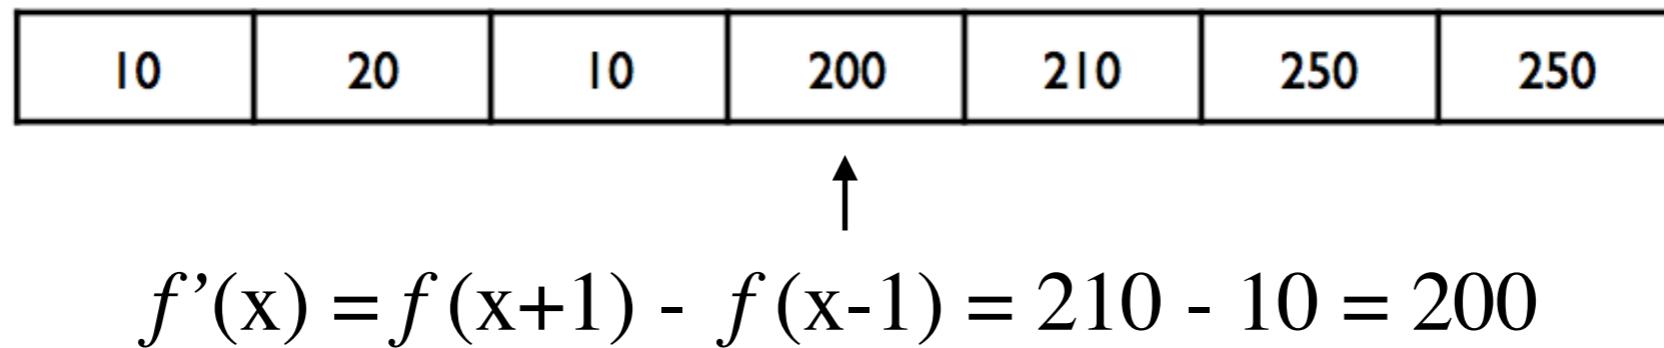
Edge detection: Sobel filter

- Where does this filter come?
 - Remember: (central) discrete derivative in 1D

$$\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$$

$[1 \ 0 \ -1]$
1D derivative filter

- Example:



Edge detection: Sobel filter

- Decomposing the Sobel Filter

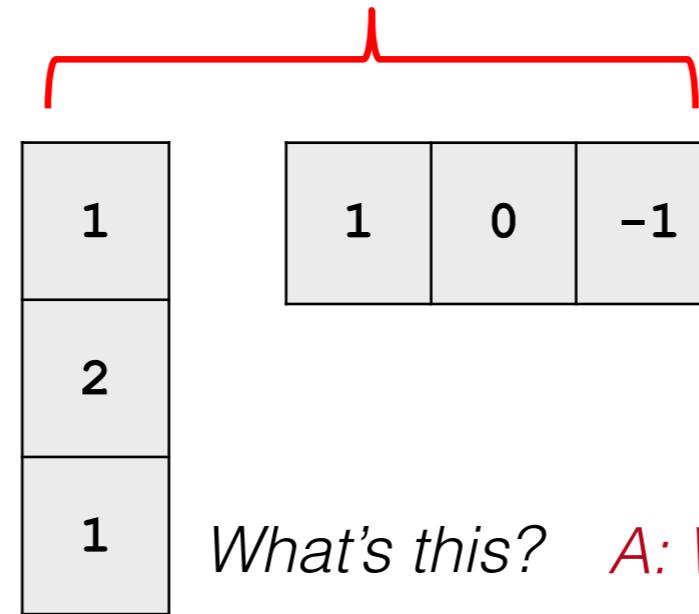


*

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

Left Sobel

Input image



1D x-derivative filter

What's this? A: Weighted average and scaling

Edge detection: derivative filters

- Common derivative filters:

Sobel

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

| | | |
|----|----|----|
| 1 | 2 | 1 |
| 0 | 0 | 0 |
| -1 | -2 | -1 |

Prewitt

| | | |
|---|---|----|
| 1 | 0 | -1 |
| 1 | 0 | -1 |
| 1 | 0 | -1 |

| | | |
|----|----|----|
| 1 | 1 | 1 |
| 0 | 0 | 0 |
| -1 | -1 | -1 |

Scharr

| | | |
|----|---|-----|
| 3 | 0 | -3 |
| 10 | 0 | -10 |
| 3 | 0 | -3 |

| | | |
|----|-----|----|
| 3 | 10 | 3 |
| 0 | 0 | 0 |
| -3 | -10 | -3 |

Roberts

| | |
|----|---|
| 0 | 1 |
| -1 | 0 |

| | |
|---|----|
| 1 | 0 |
| 0 | -1 |

Building a “robust” edge detector



Original (input) image



“Expected” Output

Building a “robust” edge detector

- The Canny Edge detector
 - ▶ Filter image with Derivative of Gaussian
 - ▶ Find magnitude and orientation of gradient
 - ▶ ...

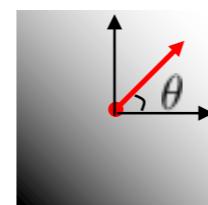
J. Canny, “ A Computational Approach To Edge Detection”, TPAMI 1986

Image gradient

$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

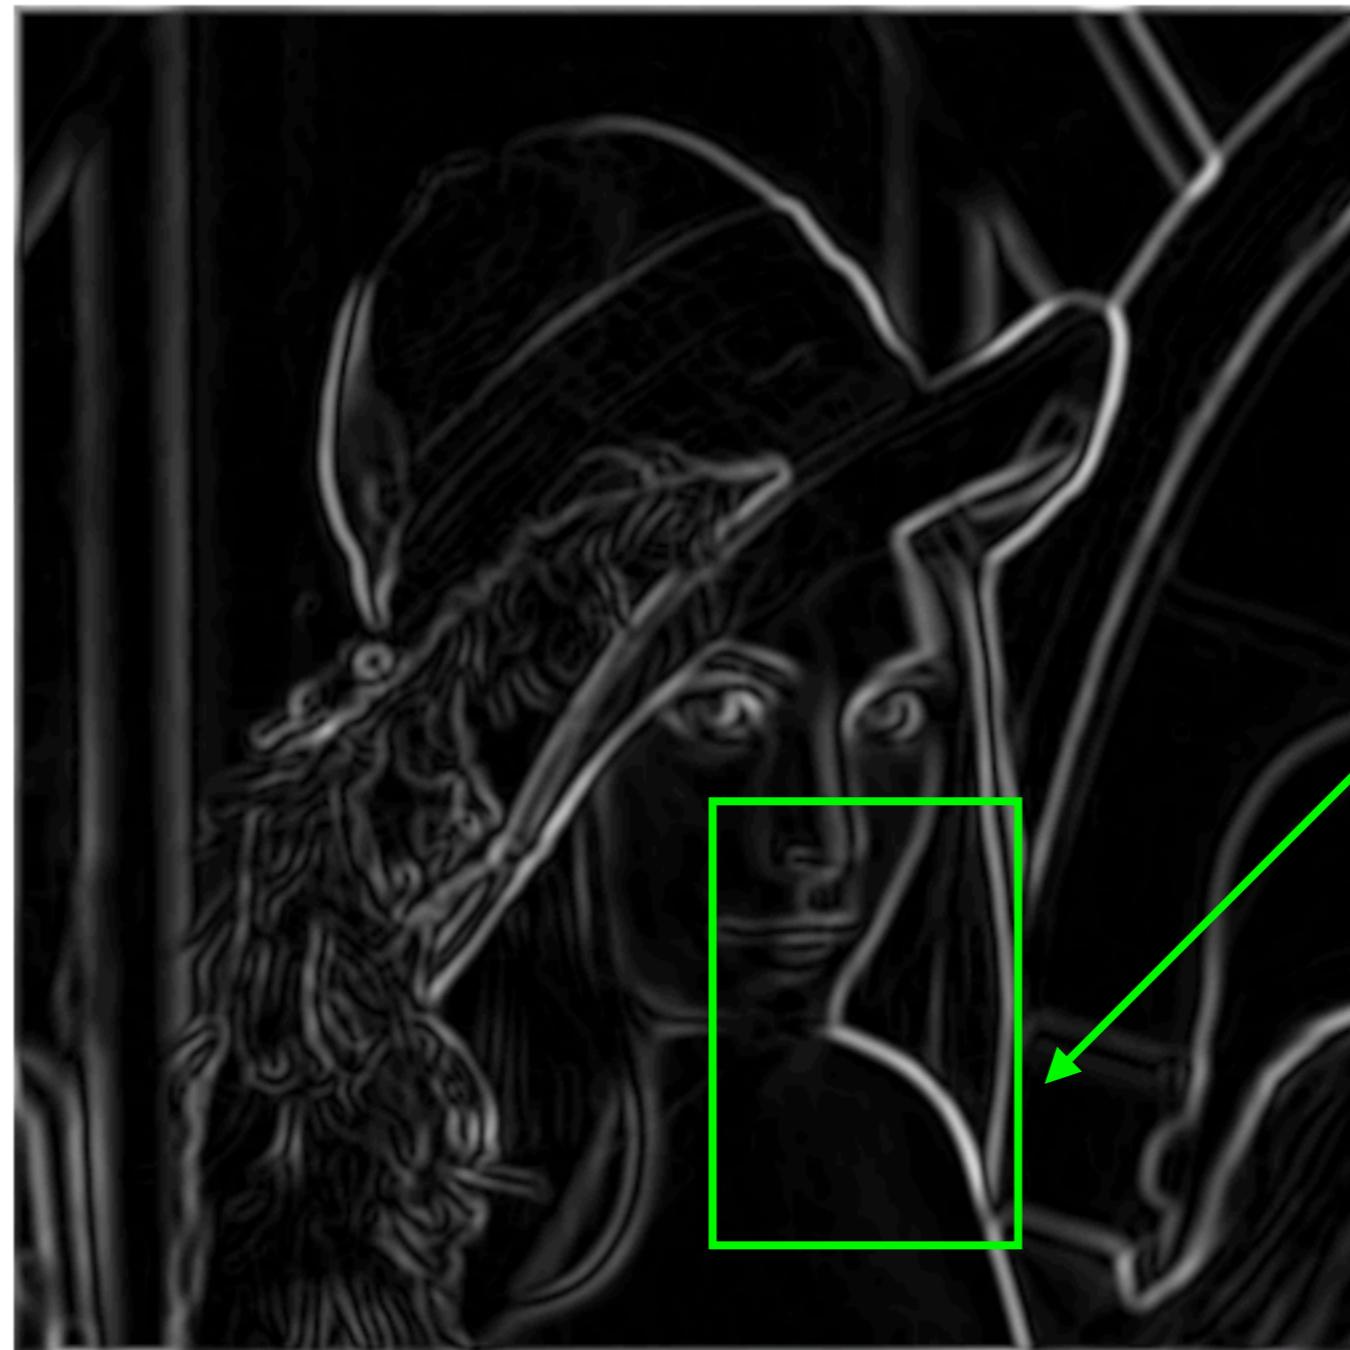
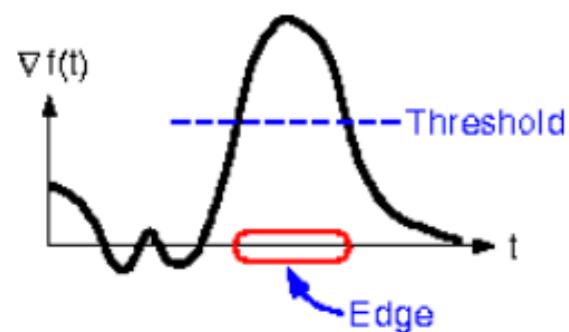
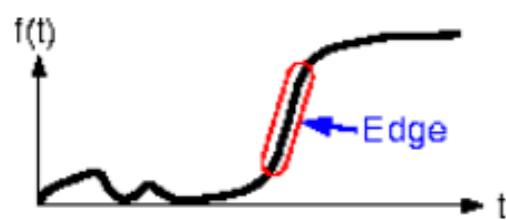
Building an edge detector: Canny



Gradient magnitude

Building an edge detector: Canny

Thresholding

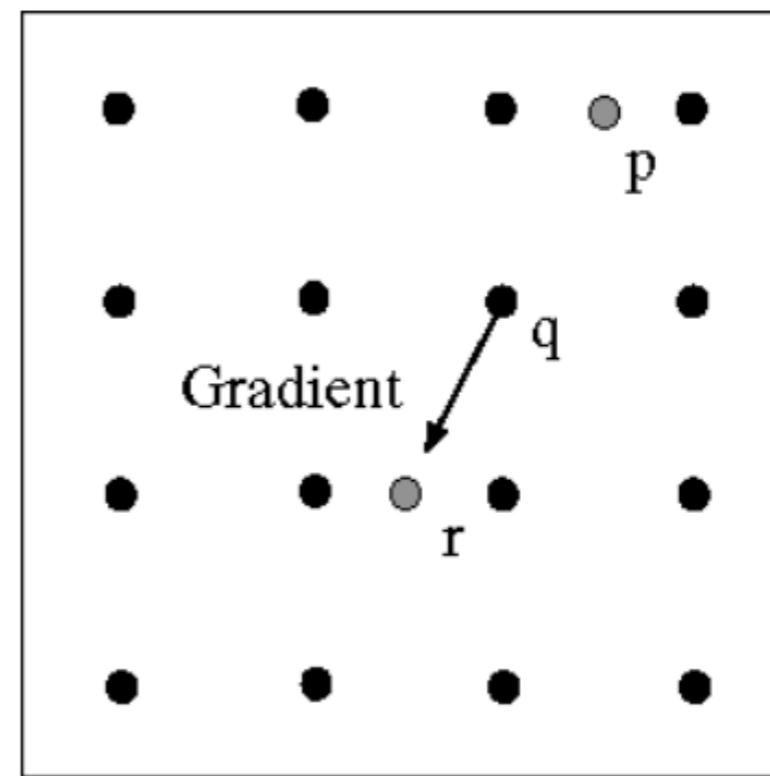
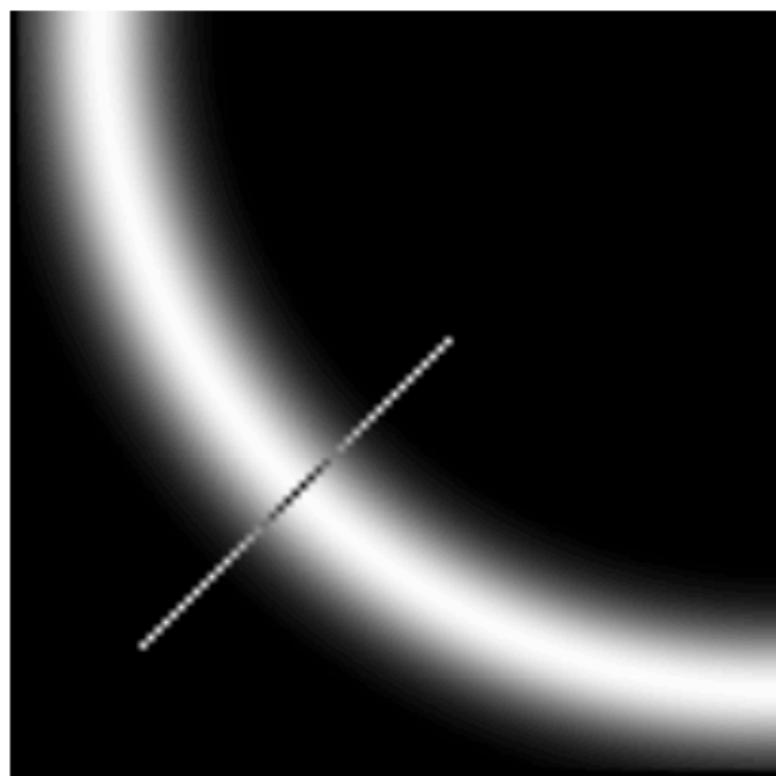


Gradient magnitude

*How to turn these
thick regions of
the gradient into
curves?*

Building an edge detector: Canny

- Non-maximum suppression
 - ▶ Check if pixel is local max along gradient direction
 - ▶ Select single maximum across width of the edge (requires checking interpolated pixels p and r)



Building an edge detector: Canny



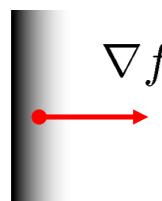
*Problem: pixels
along this edge
didn't survive the
thresholding*

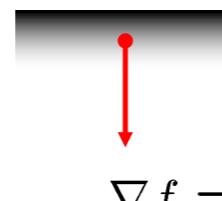
Non-maximum suppression: thinning

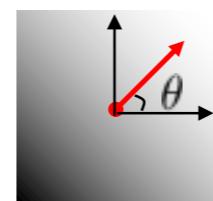
Building an edge detector: Canny

- The Canny Edge detector
 - ▶ Filter image with derivative of Gaussian
 - ▶ Find magnitude and orientation of gradient
 - ▶ Non-maximum suppression
 - ▶ Thin wide “ridges” down to single pixel width

Image gradient

$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


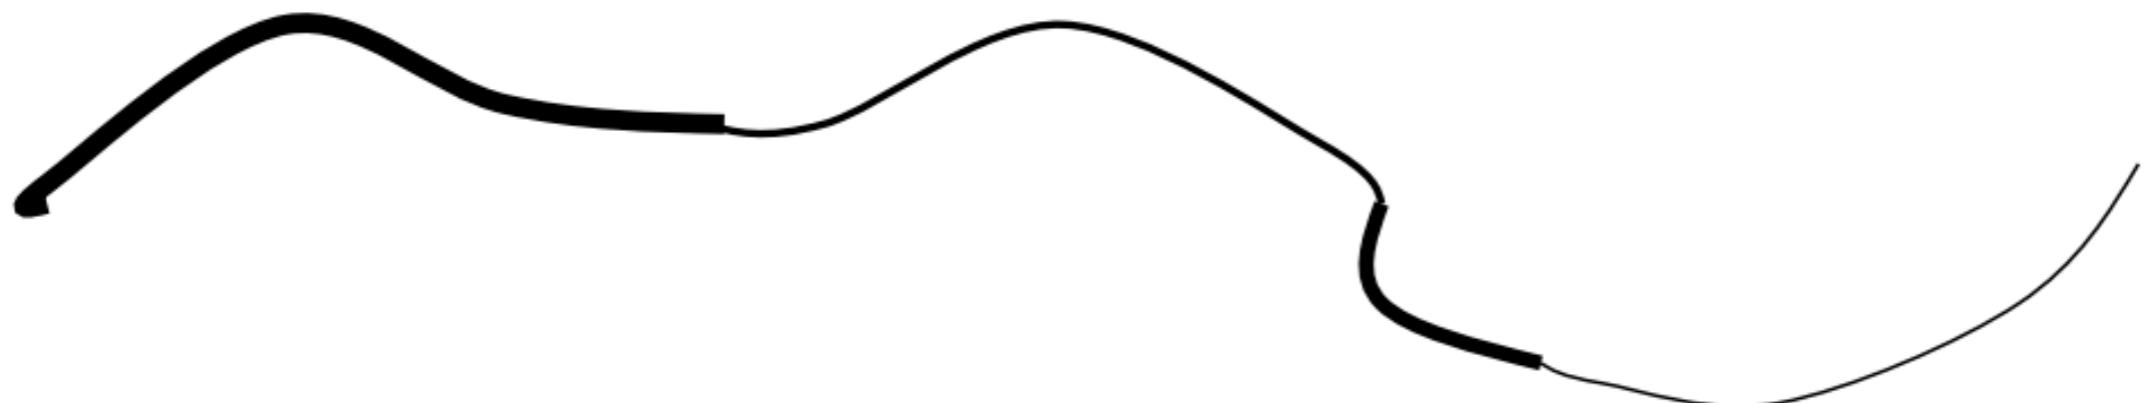
$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$


$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Building an edge detector: Canny

- Hysteresis thresholding: use a high threshold to start edge curves, and a low threshold to continue them



Building an edge detector: Canny

- Hysteresis thresholding



Original image



high threshold (strong edges)



low threshold (weak edges)



Hysteresis threshold

Building an edge detector: Canny

- The Canny Edge detector
 - Filter image with derivative of Gaussian
 - Find magnitude and orientation of gradient
 - Non-maximum suppression
 - Thin wide “ridges” down to single pixel width
 - Linking and thresholding (hysteresis)
 - Two thresholds: use the high threshold to start edge curves and the low threshold to continue them

Matlab: `edge(image, 'canny');`

J. Canny, “A Computational Approach To Edge Detection”, TPAMI 1986

Recall: a *(rough)* computer vision timeline

1966: Marvin Minsky assigns computer vision as an undergrad summer project

1970s: interpretation of synthetic worlds and carefully selected images

1980s: shift towards geometry and increased mathematical rigor

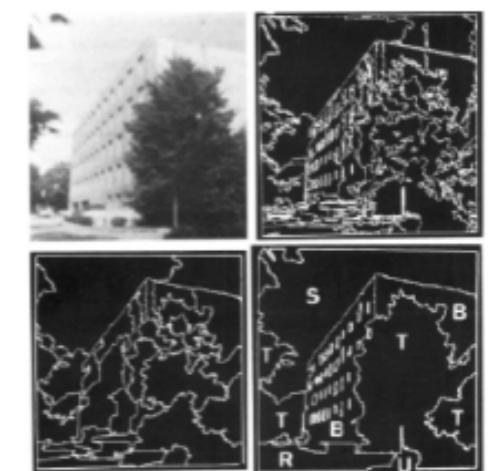
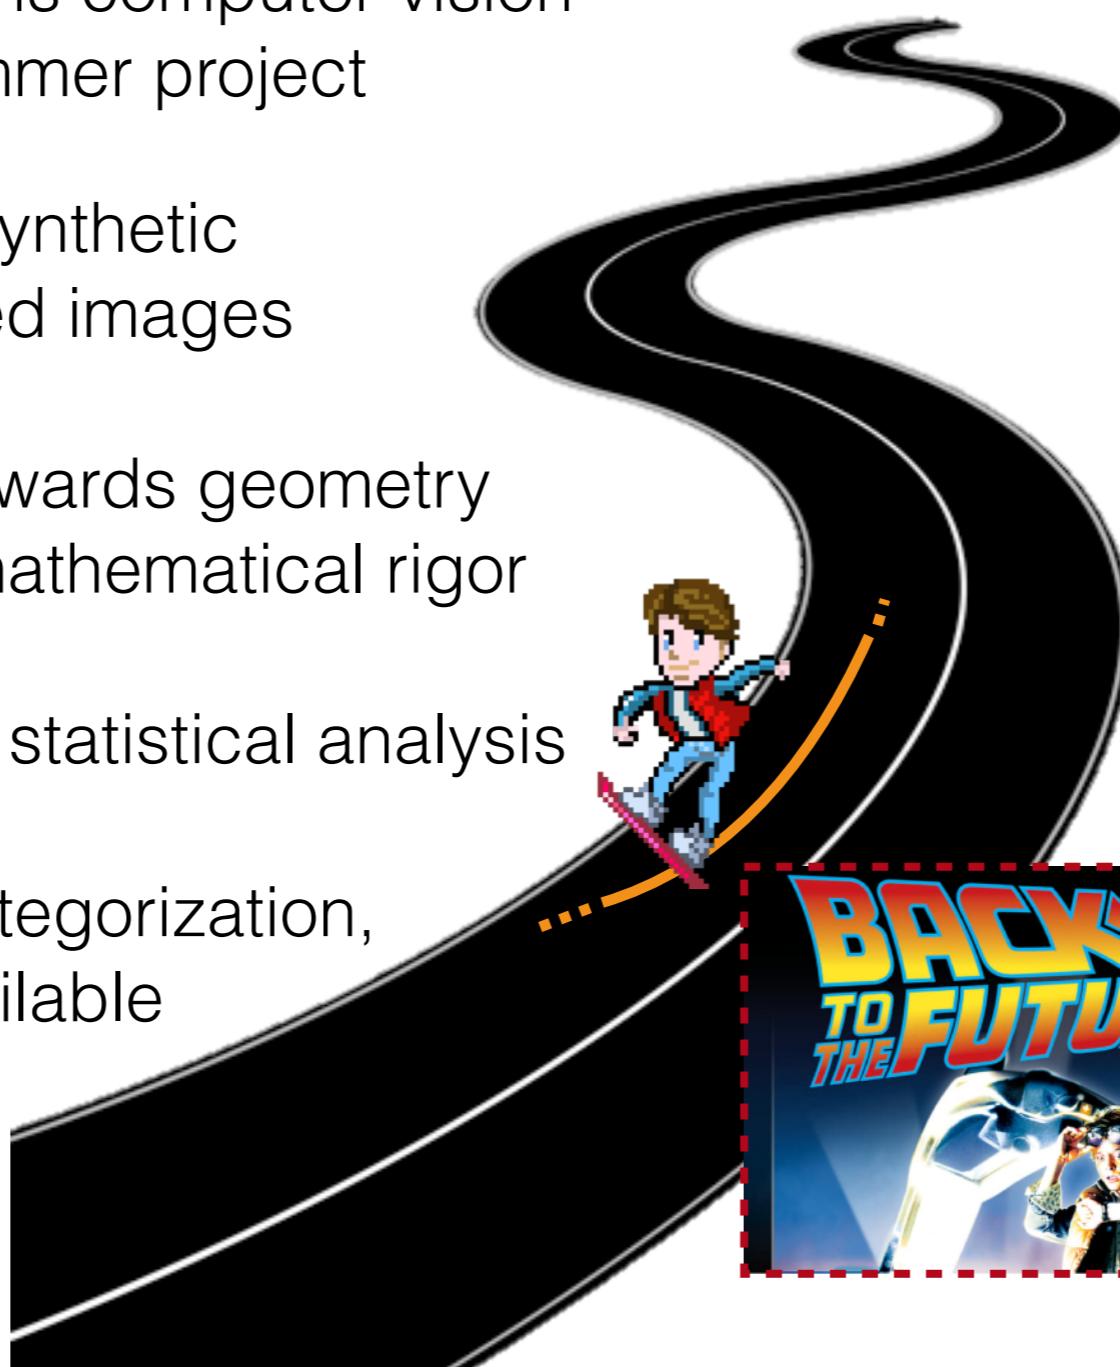
1990s: face recognition, statistical analysis

2000s: object recognition, categorization, annotated datasets available

2010s: large-scale visual recognition, visual intelligence

2020s: ???

(big/“foundation” models, ...)



Turk and Pentland '91

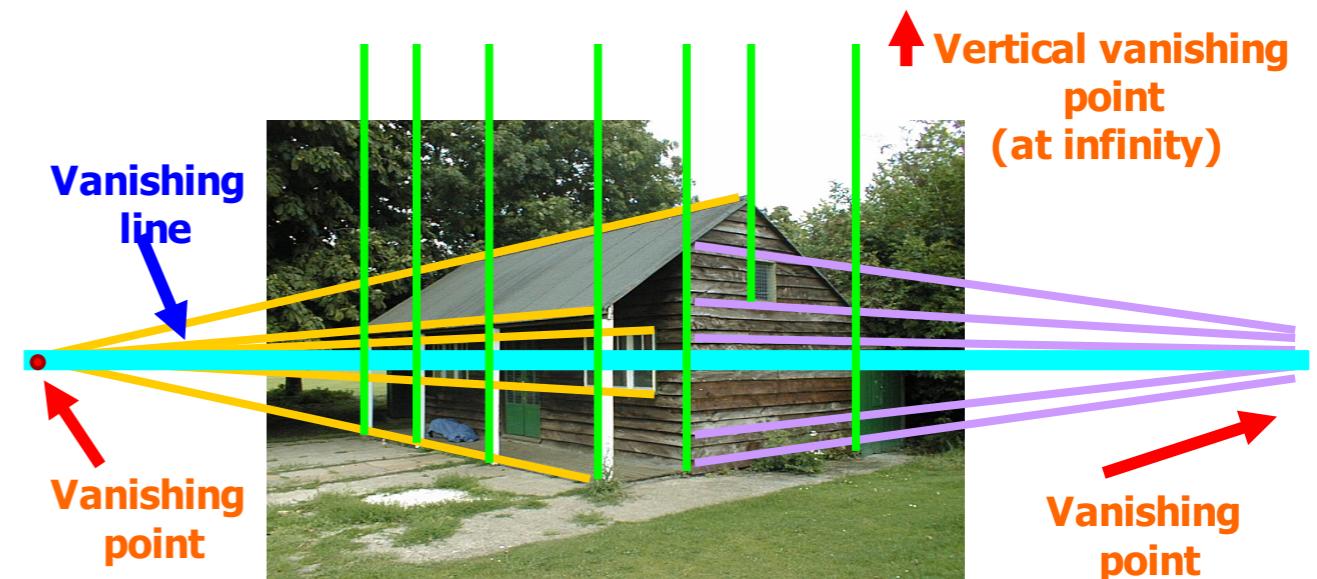
Recall: why do we care about edges?

- Extract information, recognize objects

Edges are great, but they are not very robust to a number of transformations...

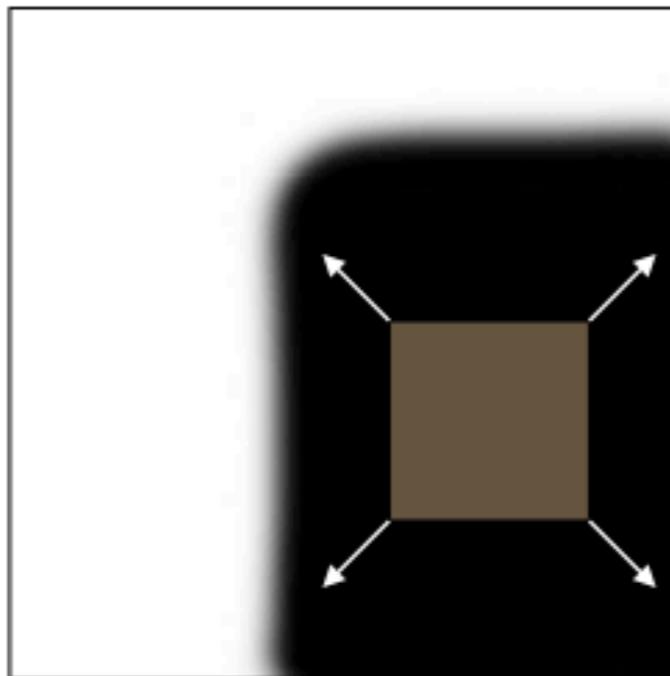


- Recover geometry and viewpoint

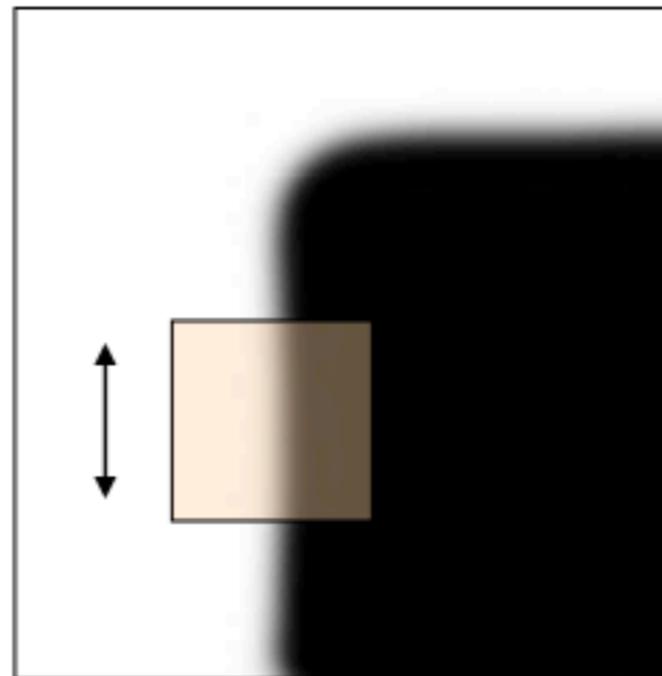


Edges, corners and blobs (keypoints)

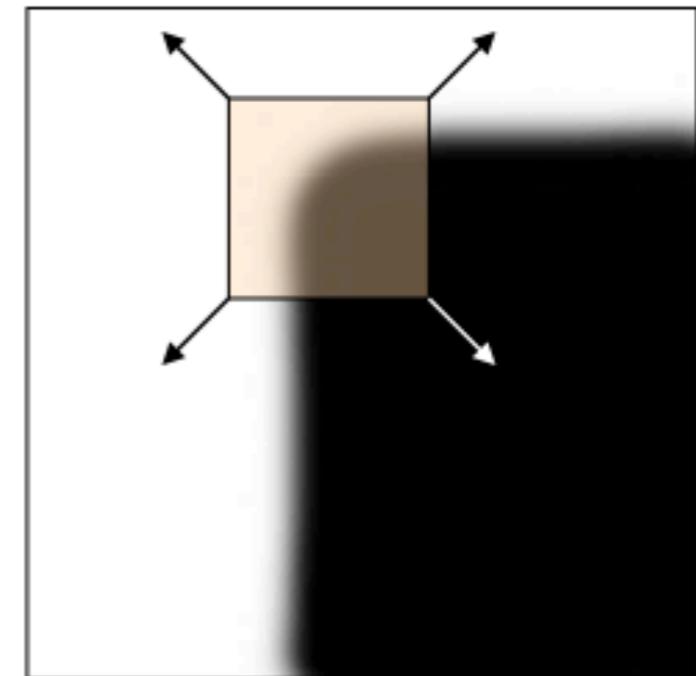
- Low-level image features: a basic classification



“flat” region:
no change in all
directions



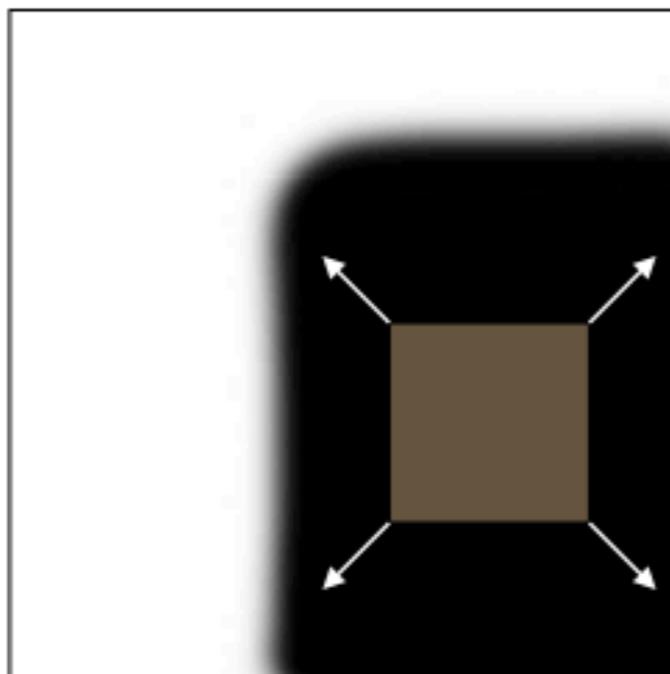
“edge”:
no change along
the edge direction



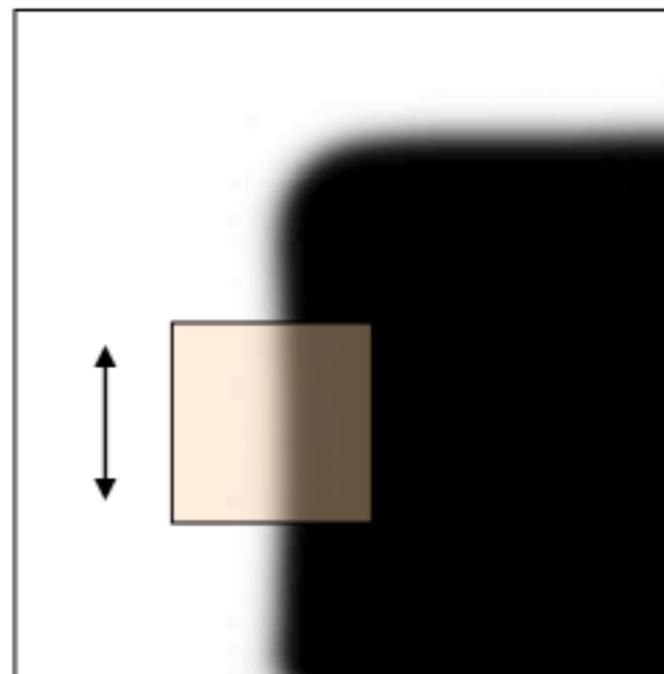
“corner”:
significant change
in all directions

Corner detection: intuition

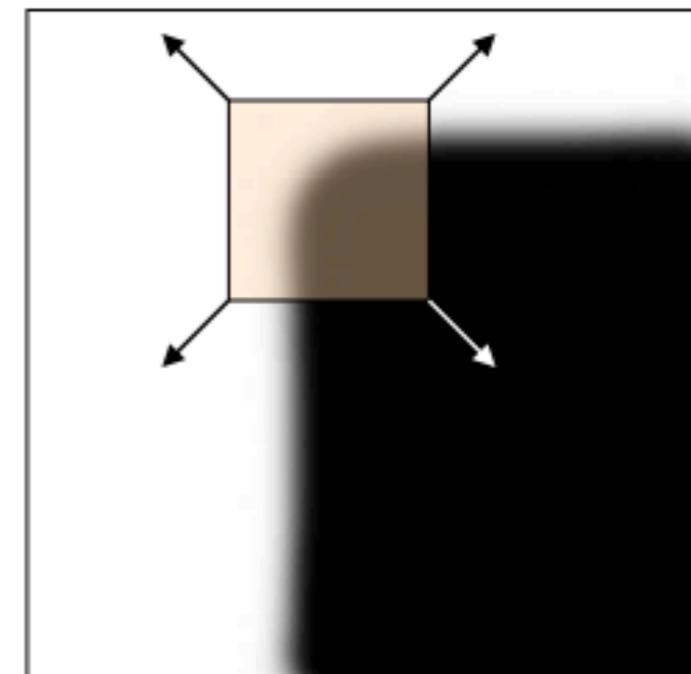
- We should easily recognize the point by looking through a small window (locally)
 - ▶ Shifting a window in any direction should give a large change in intensity



“flat” region:
no change in all directions



“edge”:
no change along the edge direction

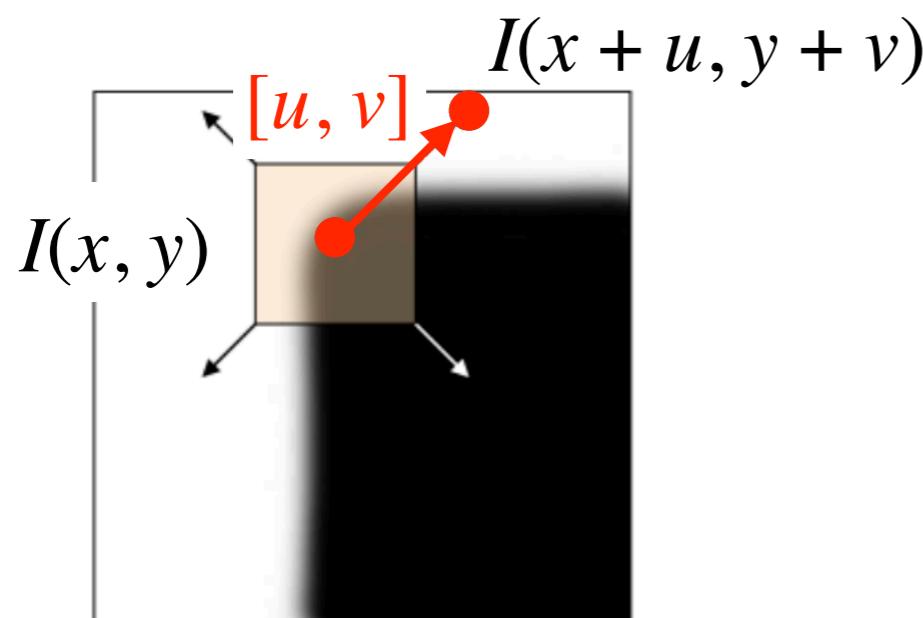


“corner”:
significant change in all directions

Corner detection: basic idea

- Given an image I , captures change in appearance of window W for the shift $[u, v]$:

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$



“corner”:
significant change
in all directions

It measures the intensity change at the center pixel (x, y) when we shift by $[u, v]$

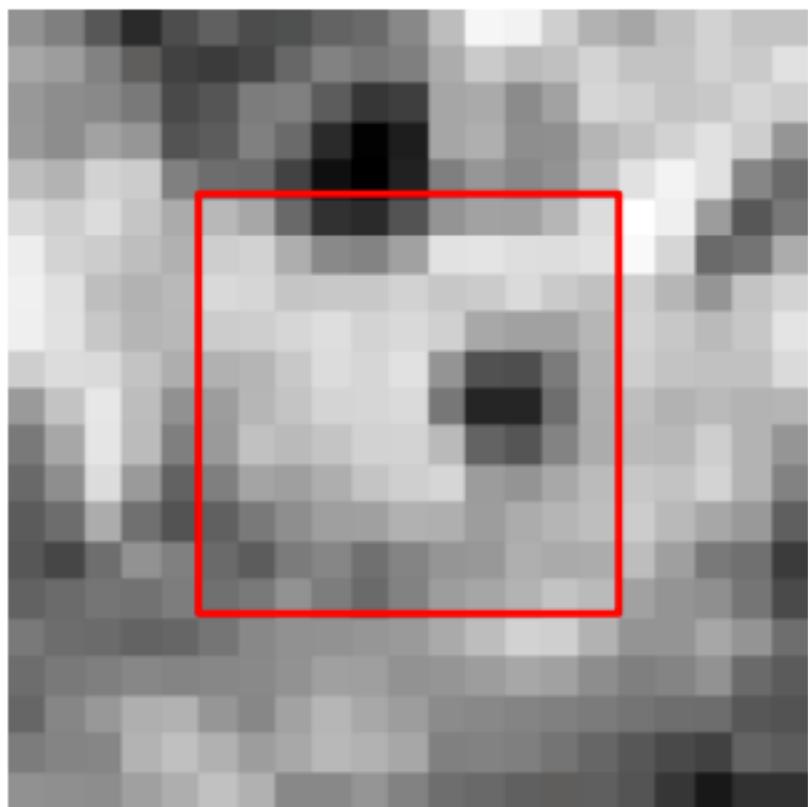
That's for a single point, $E(u, v)$ accumulates over the patch (small window) around the center pixel

Corner detection: basic idea

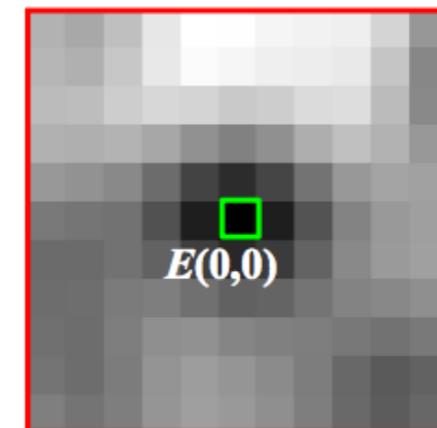
- Given an image I , captures change in appearance of window W for the shift $[u, v]$:

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

$I(x, y)$



$E(u, v)$



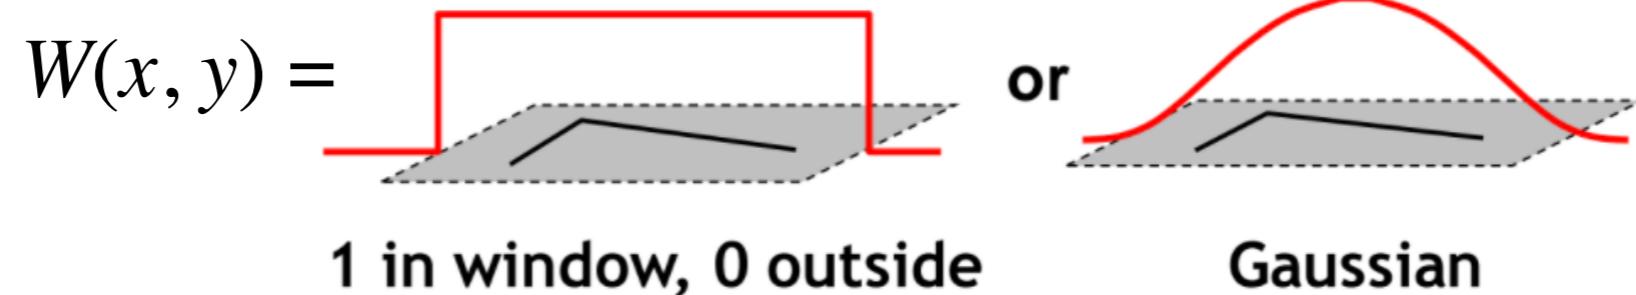
Corner detection: basic idea

- Given an image I , captures change in appearance of window W for the shift $[u, v]$:

$$E(u, v) = \sum_{(x,y)} W(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window function **Intensity change**

**Window
function**



Corners and keypoints

- Corners (e.g. Harris corner detector) are a good example of keypoints / interest points



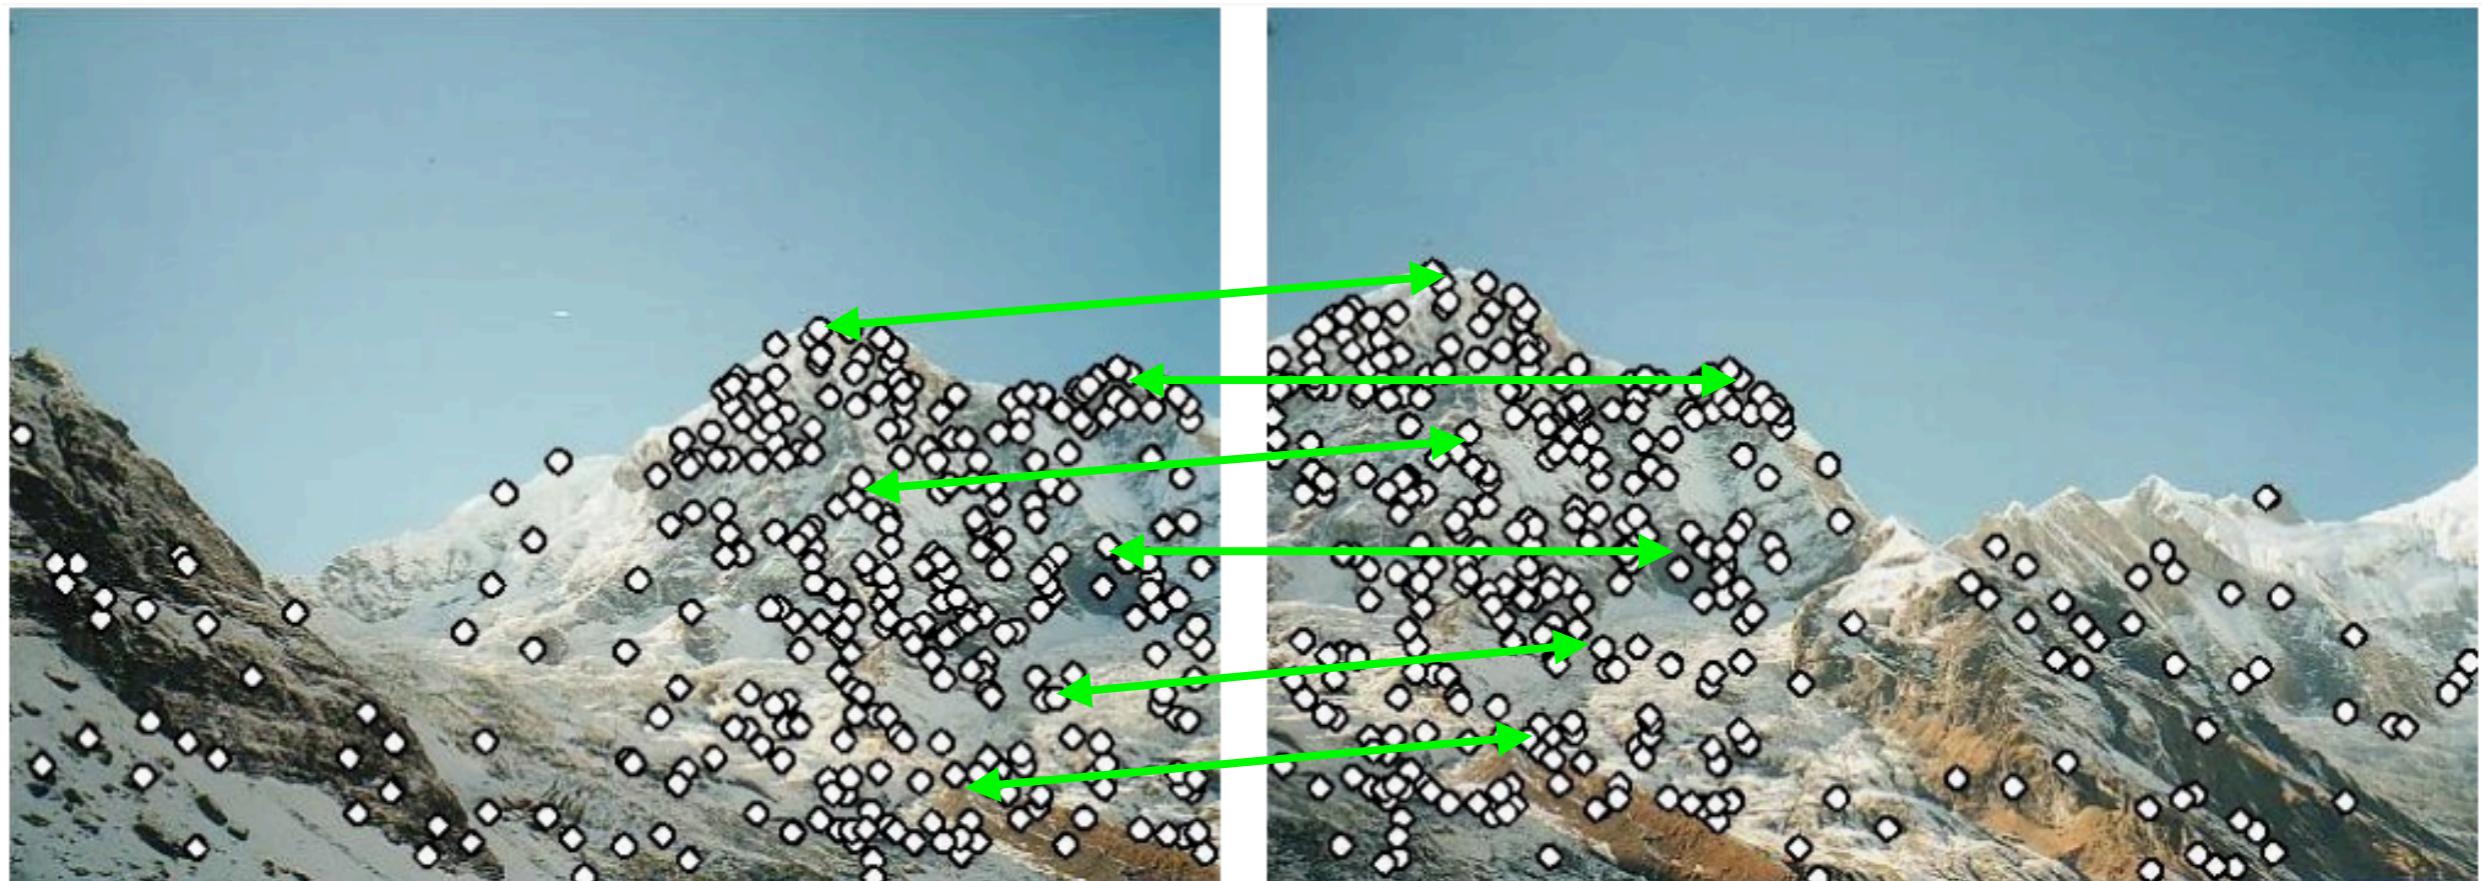
Why extract keypoints from an image?

- Motivation: a key example is panorama stitching
 - ▶ i.e. we have two images, how to combine them?



Why extract keypoints from an image?

- Motivation: a key example is panorama stitching
 - ▶ i.e. we have two images, how to combine them?

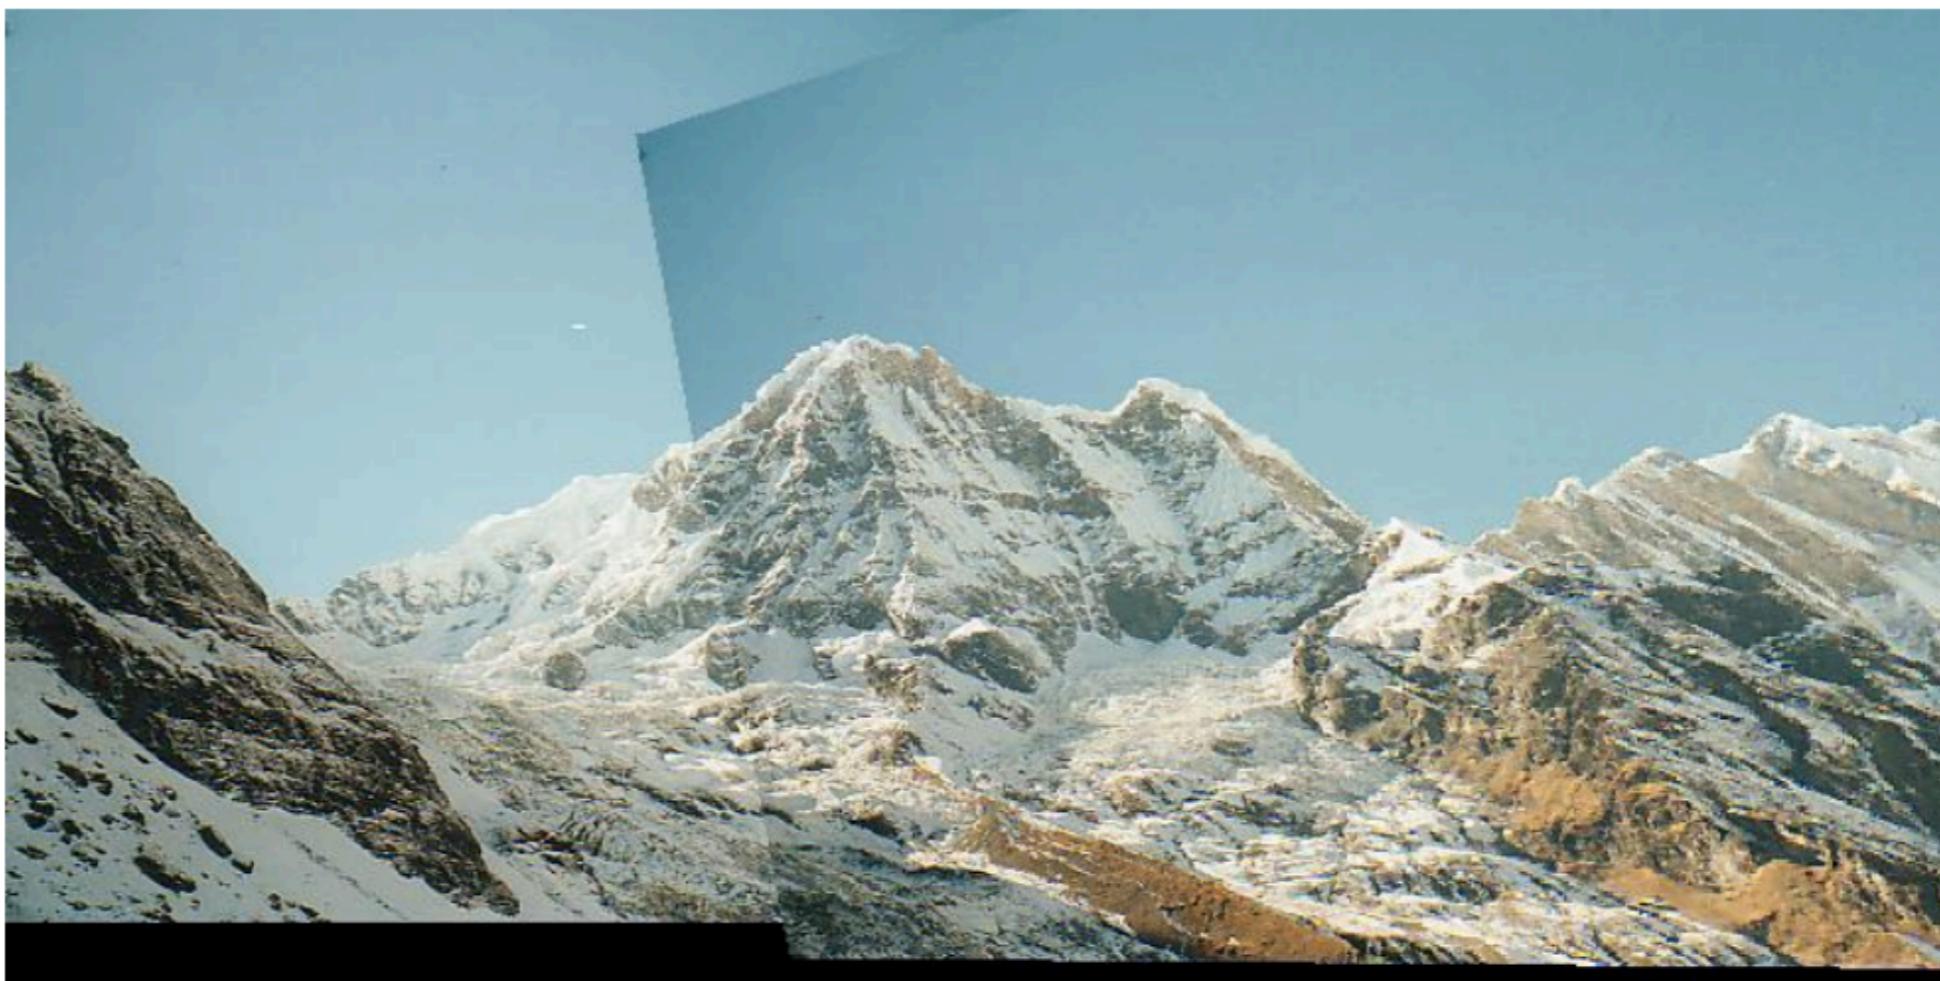


Step 1: extract keypoints

Step 2: match keypoint features

Why extract keypoints from an image?

- Motivation: a key example is panorama stitching
 - ▶ i.e. we have two images, how to combine them?



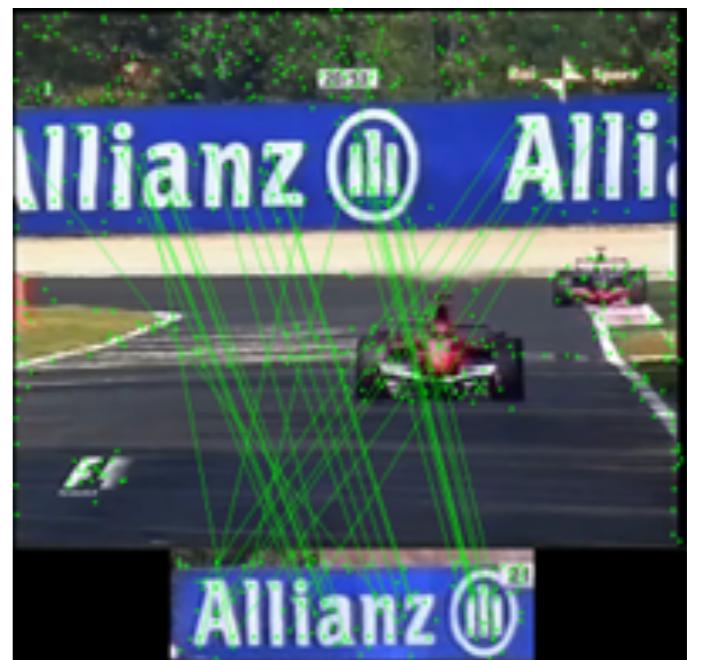
Step 1: extract keypoints

Step 2: match keypoint features

Step 3: align images

Keypoints: applications

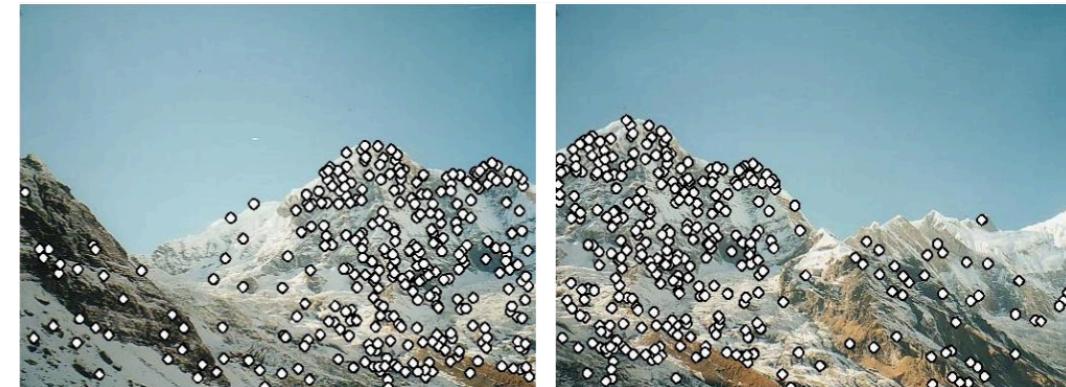
- Keypoints are used for:
 - ▶ Image alignment
 - ▶ 3D reconstruction
 - ▶ Motion tracking
 - ▶ Robot navigation
 - ▶ Indexing and database retrieval
 - ▶ Object recognition



Characteristics of Good Keypoints

- Repeatability

- Can be found despite geometric and photometric transformations



- Salience

- Each keypoint is distinctive

- Compactness and efficiency

- Many fewer keypoints than image pixels

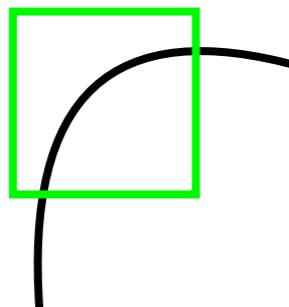
- Locality

- Occupies small area of the image: therefore more robust to clutter and occlusion

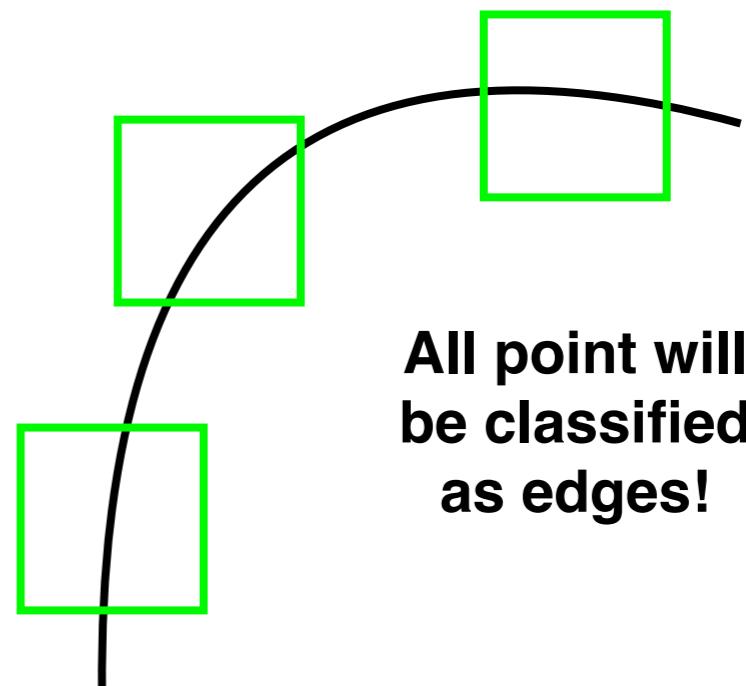
An example: Harris corner detector

- Properties (i.e. geometric invariances):

- Translation invariance? Yes
- Rotation invariance? Yes
- Scale invariance? No

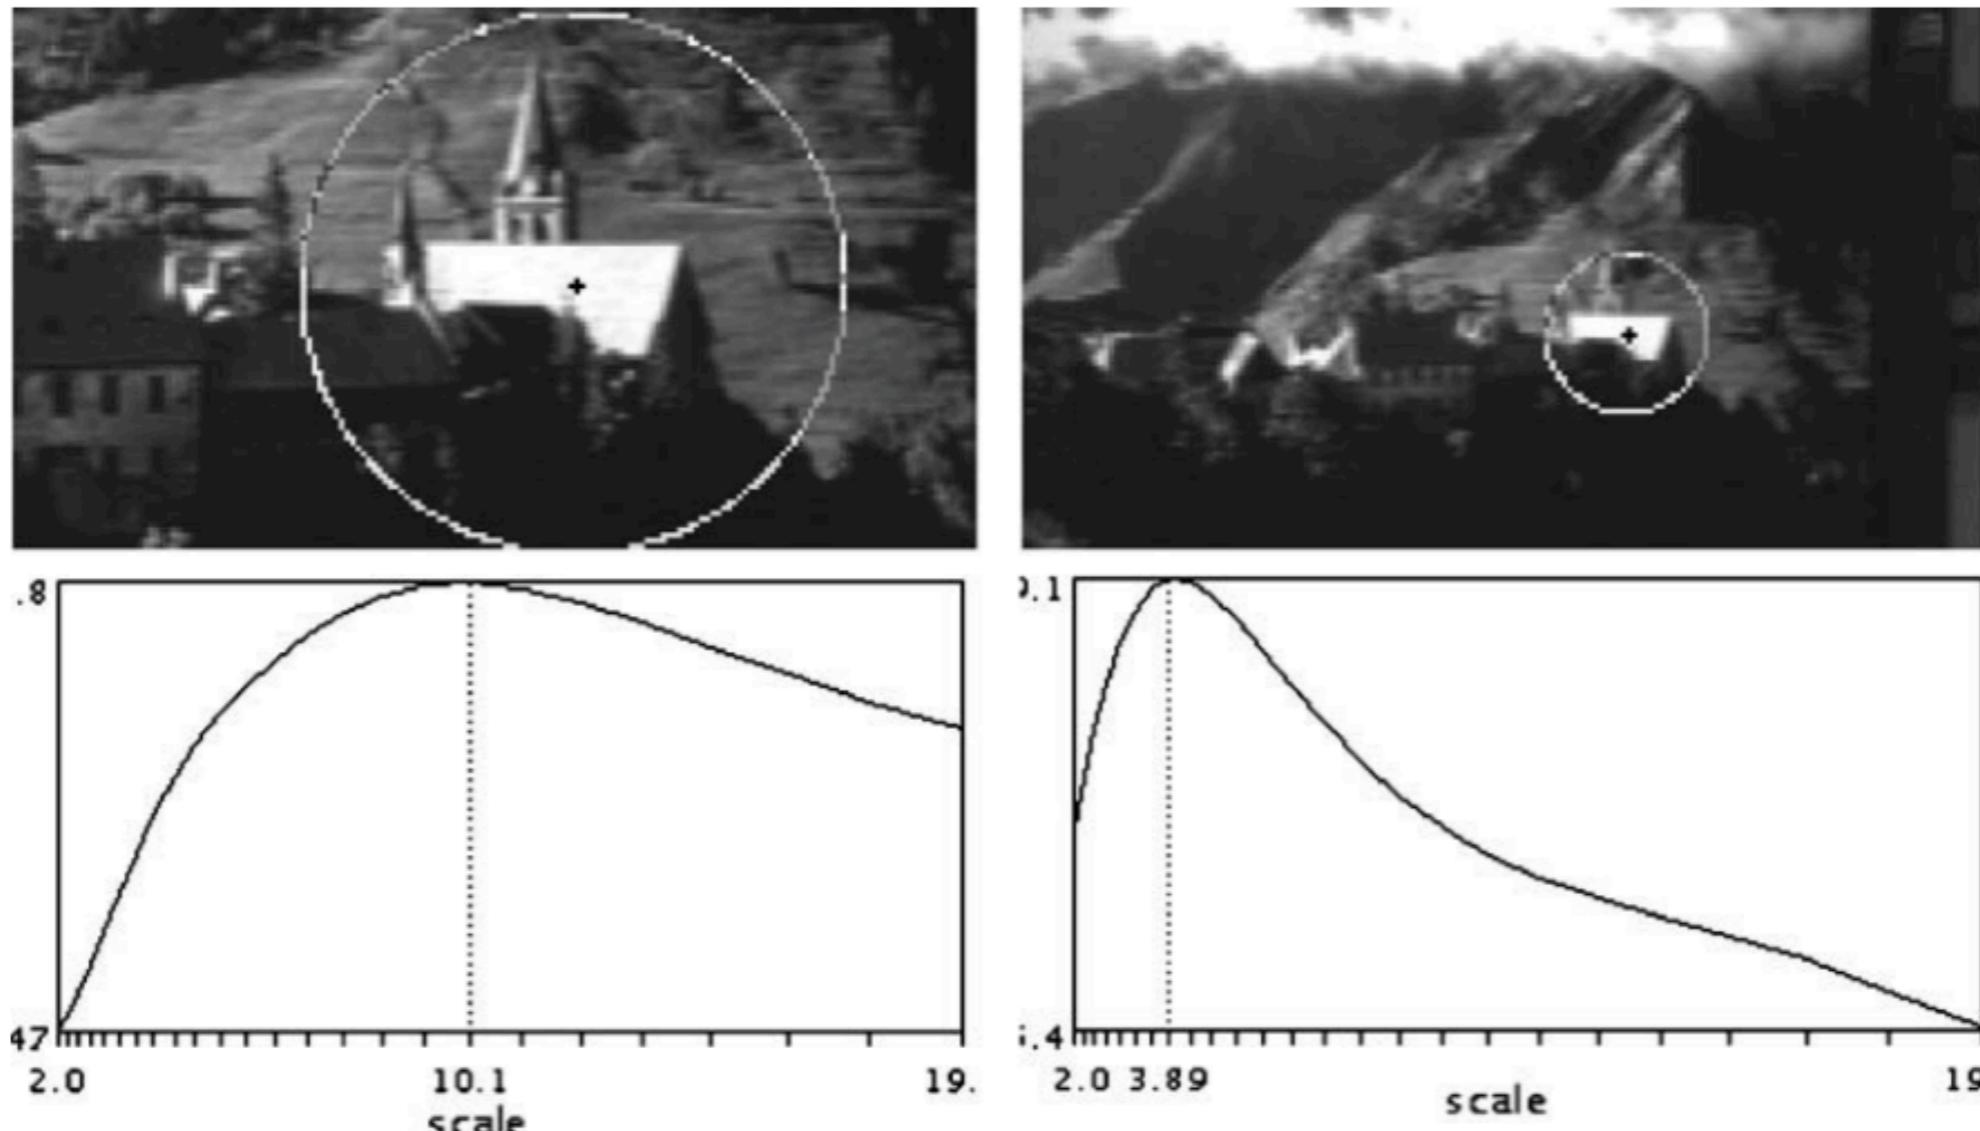


Corner



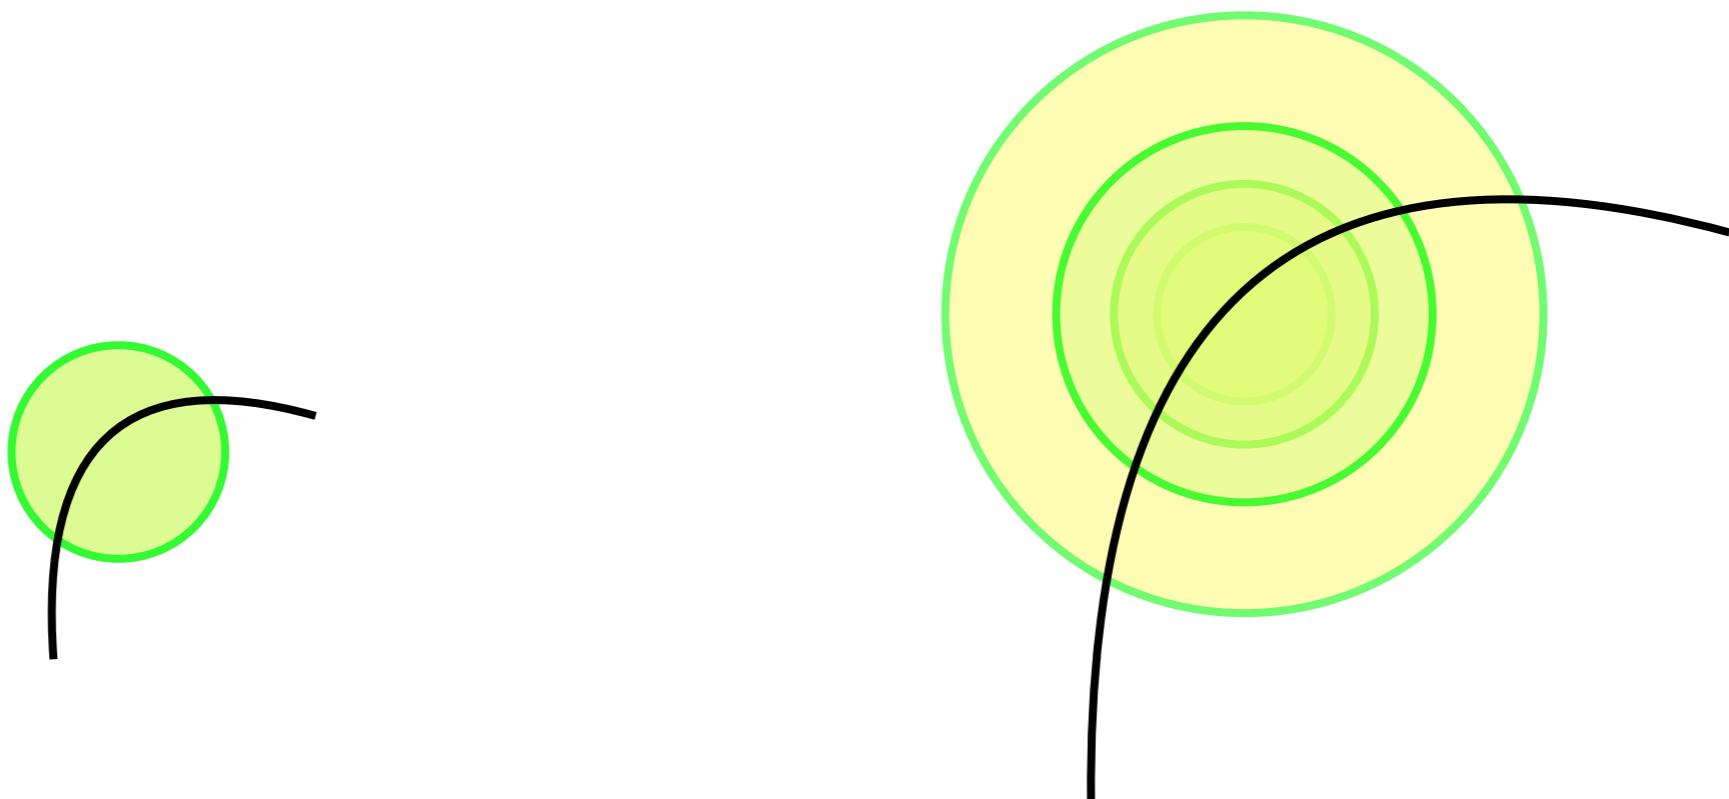
Feature detection with scale selection

- We want to extract features with characteristic scale that is invariant (*covariant*) with the image transformation



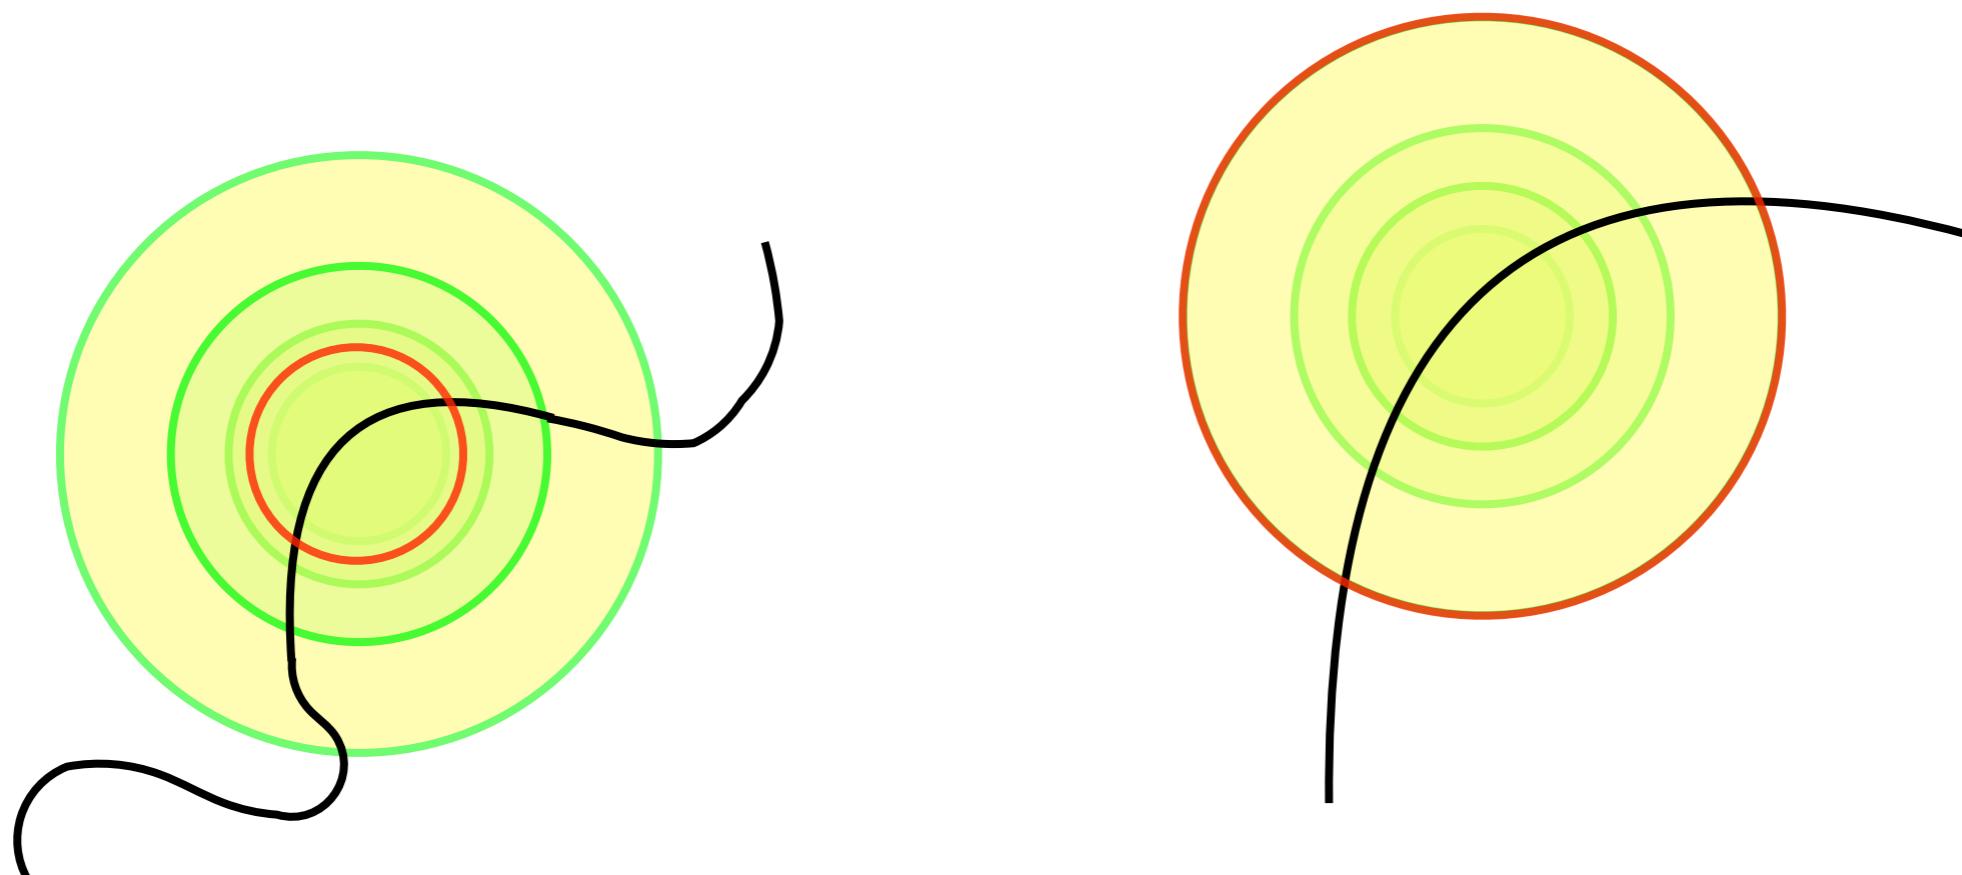
Scale invariant detection

- Consider regions (e.g. circles) of different sizes around a point
- How to find regions of corresponding sizes that will look the same in both images?



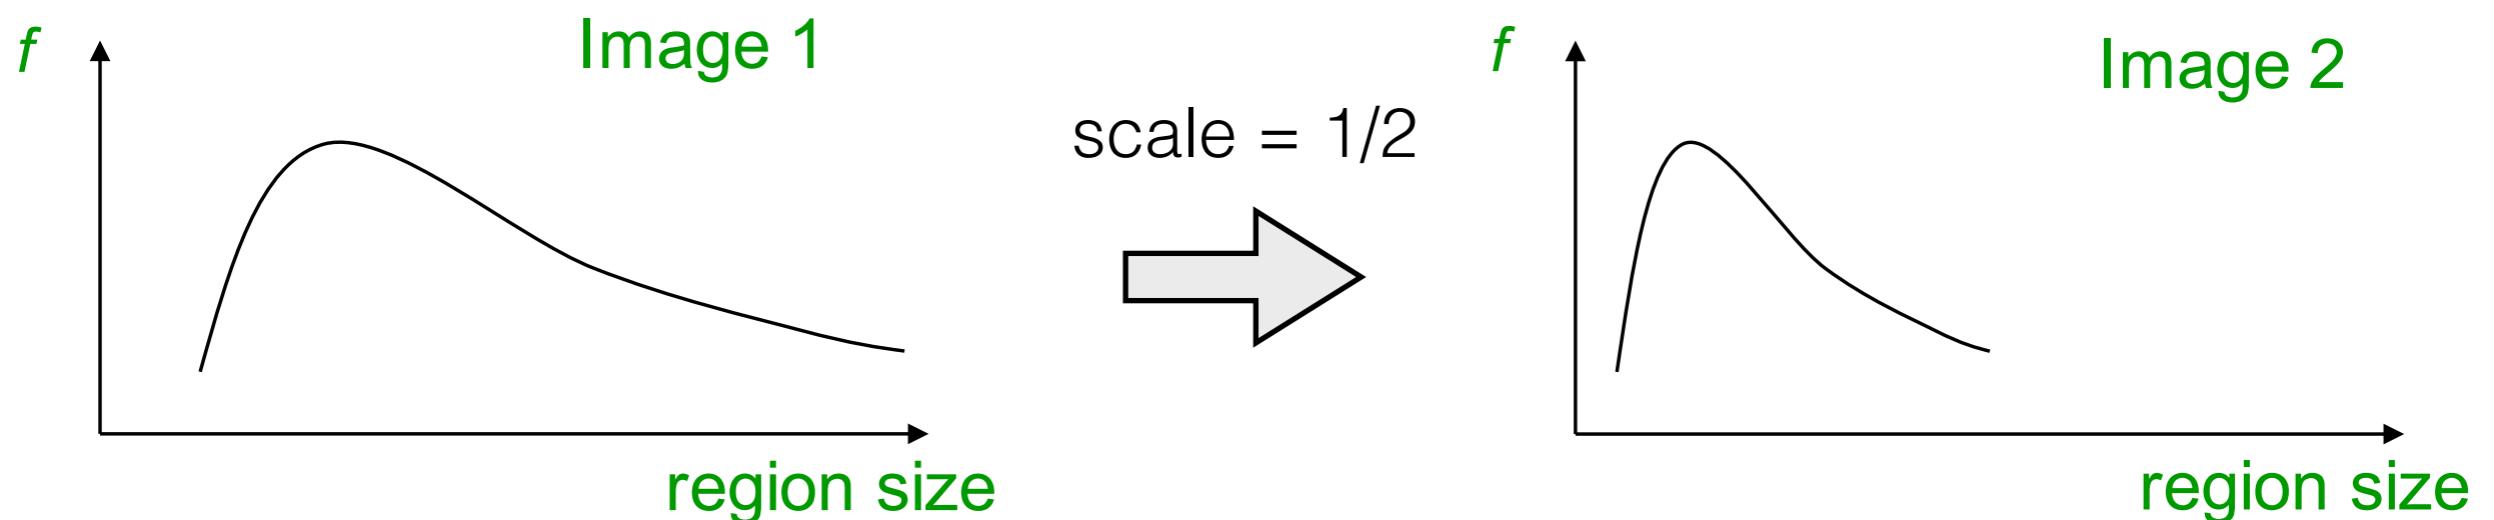
Scale invariant detection

- The problem: how do we choose corresponding circles **independently** in each image?



Scale invariant detection

- Solution: design a function on the circular region which is scale invariant *
 - e.g. average intensity (same even for different sizes)
- For a point in one image, we can consider it as a function of region size (circle radius)

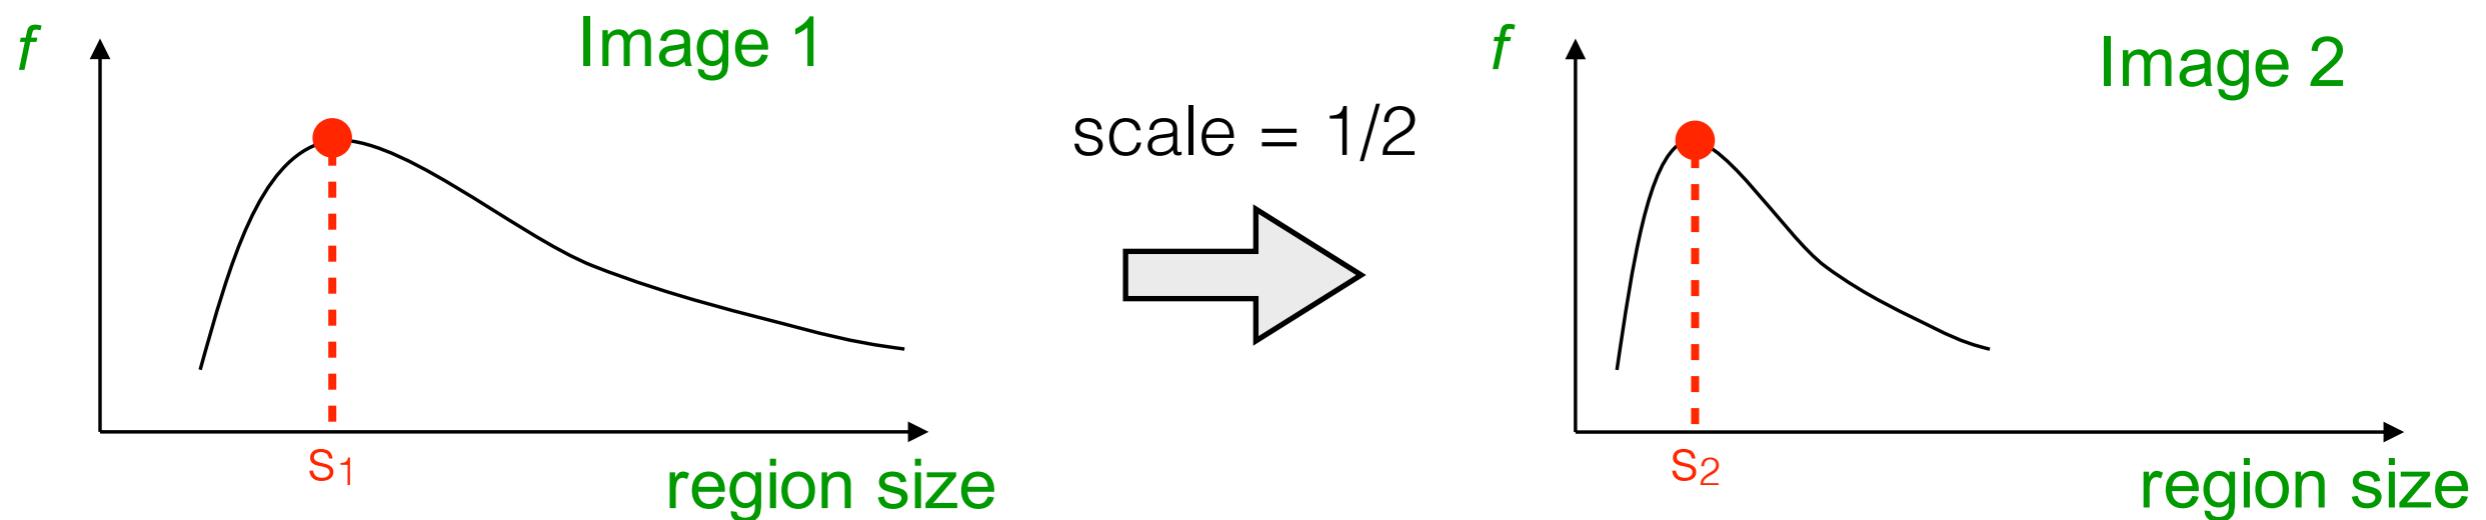


* i.e. the same for corresponding regions, even if they are at different scales

Scale invariant detection

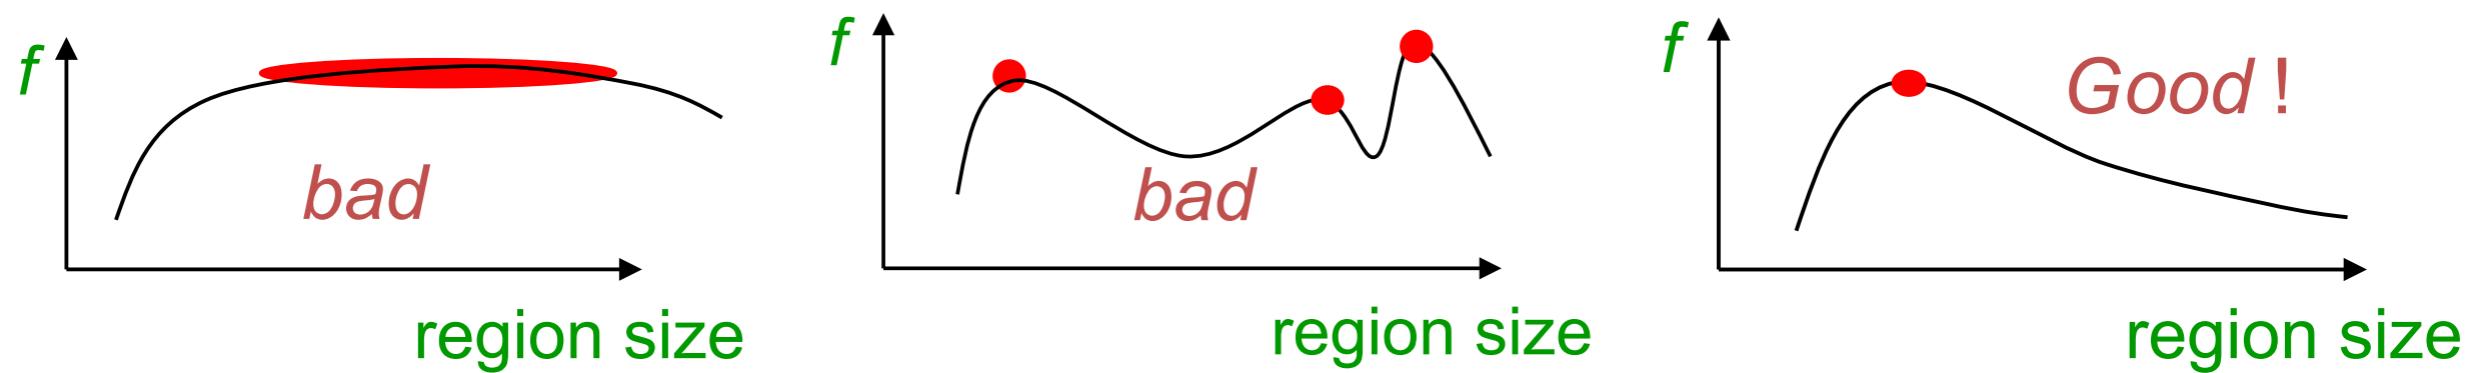
- Common approach: take a local max of this function
- Observation: region size for which the maximum is achieved should be *invariant (co-variant)* with scale

Important: this scale invariant region size is found in each image **independently!**



Scale invariant detection

- A good function for scale detection has one stable sharp peak



- So a good function would be one which responds to contrast (sharp local intensity change)

Contact

- **Office:** Torre Archimede, room 6CD3
- **Office hours** (ricevimento): Friday 9:00-11:00

✉ lamberto.ballan@unipd.it
⬆ <http://www.lambertoballan.net>
⬆ <http://vimp.math.unipd.it>
{@} twitter.com/lambertoballan